

Stochastic Gradient Descent

Contents

1. Overview	2
2. The problem with batch gradient descent	2
3. How stochastic gradient descent works	3
4. Faster progress per epoch	7
5. A noisy path	8
5.1 The zigzag	8
5.2 Advantages of the noise	8
5.3 Disadvantage: jitter at the bottom	9
6. Learning schedules	10
7. SGD in scikit-learn	11
8. When to use which	13
9. Summary	15
10. Sources	15
11. Key terms	15

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** [Gradient descent](#).
- **Leads to:** [Perceptron loss](#); [Batch size in Keras](#).
- **Compare with:** [Batch gradient descent](#); [Mini-batch gradient descent](#); [SGD with momentum](#).

1. Overview

Key point: Stochastic gradient descent updates the coefficients after every single random observation. Each step is cheap and noisy, so it gets close to the answer in far fewer epochs than batch gradient descent, but it jitters around the minimum.

A **feature** (G-772) is an input variable (one column of the data table), an **observation** (G-1374) is one record (one row), and the **target** (G-1949) is the output we predict.

Batch gradient descent (G-264), from [batch gradient descent with many features](#), uses every observation of the training data for each update. **Stochastic gradient descent (SGD)** (G-1892) uses just one observation, picked at random, for each update. SGD and its variants are probably the most used **optimisation algorithms** (G-1398) in machine learning, and nearly all of deep learning is trained with it (Goodfellow §8.3.1, §5.9).

This Note explains:

- the problem with batch gradient descent (section 2);
- how SGD works (section 3) and why it progresses faster per epoch (section 4);
- its noisy path (section 5) and how a learning schedule calms it (section 6);
- how to use it in scikit-learn (section 7), and when to choose which (section 8).

2. The problem with batch gradient descent

Key point: Every batch update touches every observation and every feature. With big data, every step is a huge amount of work, and the whole dataset must sit in memory.

To update once, batch gradient descent computes the derivative for every coefficient, and each derivative sums over all n observations. With n observations and m features, one epoch costs about $n \times m$ multiplications, and the whole run costs about $n \times m \times \text{epochs}$.

With numbers: a dataset of 100,000 observations and 100 features, trained for 1,000 epochs, needs about

$$100,000 \times 100 \times 1,000 = 10^{10}$$

multiplications, ten billion, just for the derivatives.

Figure 1 compares the cost of one update. Each bar is one method; the left chart uses a log scale, so each gridline is ten times the one below (a bar four gridlines shorter is 10,000 times smaller). Batch gradient descent needs 10 million multiplications for every single step. SGD needs only 100, the features of one observation, so it can take 100,000 steps in the time batch takes one. Both read the whole table once per **epoch** (G-696); SGD spends that work on many small steps instead of one big one.




Figure 1: The cost of one update for 100,000 observations and 100 features. Left: multiplications per update, 10,000,000 for batch and 100 for SGD (log scale). Right: updates per epoch, 1 for batch and 100,000 for SGD.

So each step is slow and needs the whole dataset in memory: the two disadvantages listed in [advantages and disadvantages](#).

3. How stochastic gradient descent works

Key point: In each epoch, repeat n times: pick one random observation, compute the derivatives from that observation alone, and update.

The word **stochastic** (G-1893) means random. The algorithm:

1. Start with any coefficients, for example $\beta_0 = 0$ and every $\beta_j = 1$.
2. For each epoch, repeat n times:
 - pick one observation i at random;
 - compute its prediction $\hat{y}_i$ and error $y_i - \hat{y}_i$;
 - update every coefficient using the derivatives from that one observation.
3. Stop after the chosen number of epochs.

The derivatives are the batch ones with the sum over observations removed: only observation i remains. Where they come from, for one observation: its loss is the squared error $(y_i - \hat{y}_i)^2$.

By the chain rule, the derivative is 2 times the error times the derivative of the error with respect to the coefficient:

$$\frac{\partial(y_i - \hat{y}_i)^2}{\partial\beta_j} = 2(y_i - \hat{y}_i) \times \left(-\frac{\partial\hat{y}_i}{\partial\beta_j}\right)$$

The prediction is

$$\hat{y}_i = \beta_0 + \beta_1 x_{i1} + \dots$$

so the derivative of $\hat{y}_i$ with respect to β_j is x_{ij} (and 1 for β_0), which gives the two results below:

$$\frac{\partial L}{\partial\beta_0} = -2(y_i - \hat{y}_i)$$

$$\frac{\partial L}{\partial\beta_j} = -2(y_i - \hat{y}_i) x_{ij}$$

With numbers: if observation i has error $y_i - \hat{y}_i = 5$ and $x_{ij} = 0.04$, then the derivative is

$$\frac{\partial L}{\partial\beta_j} = -2 \times 5 \times 0.04 = -0.4$$

and with **learning rate** (G-1068; the step size) 0.01 the coefficient rises by

$$0.01 \times 0.4 = 0.004$$

So an epoch of SGD makes n small updates instead of one big one.

This Note runs SGD on two datasets from the earlier Notes:

- the **diabetes data** of [batch gradient descent in code](#): 10 features and 353 training patients, so an epoch is 353 updates;
- the **100-point example** of [gradient descent](#): 100 points with one feature, so only a slope m and an intercept b to learn, which lets us draw the paths.

Why one observation is enough. Real data repeats itself: many observations look alike and say almost the same thing about the coefficients. Take 12 points that sit in three tight clusters. The gradient from one point of a cluster is close to the gradient from its neighbours, so summing all 12 before every step mostly repeats work. One random point already points roughly downhill, and the next random point corrects it. SGD gains the most when the data has such redundancy (StatQuest, “Stochastic Gradient Descent, Clearly Explained!!!”).

The next figure draws paths on a contour map (see [reading a contour map](#)), so first the surface it comes from. Every pair (m, b) is one line, and its loss is the mean squared error of

the 100-point example: for the start line $m = -127.8$, $b = 150$ the loss is about 41,800, and at the best line $m = 27.8$, $b = -2.3$ it is about 283. The loss is a smooth bowl over the (m, b) floor; [the shape of the loss function](#) shows the same kind of bowl tilting from the side view to the top view. Figure 2 puts the bowl beside its map: the map is the bowl seen from above, and each line joins points at the same height. Lines close together mean a steep slope; the centre ring, with the black cross, is the lowest point. The grey square is the start used in the races below.




Figure 2: The loss surface of the 100-point example (left) and the same surface seen from above, a contour map (right). Grey square: the start $m = -127.8$, $b = 150$. Black cross: the best line.

Figure 3 makes the race fair: every frame, each method reads the same 10 observations of the 100-point example (one feature, so only the slope m and intercept b to learn). Watch the orange SGD path: 10 small jumps per frame, so it reaches the bottom of the bowl within half an epoch, then keeps zigzagging there. Batch gradient descent (red) has to read all 100 observations before its first step.




Figure 3: Batch, mini-batch and stochastic gradient descent reading the same data, 10 observations per frame, learning rate 0.05. Left: the paths on the loss contours. Right: the loss against the observations read.

Python: Stochastic gradient descent from scratch.

```
class MySGDRegressor:
    def __init__(self, learning_rate=0.01, epochs=100, seed=0):
        self.lr, self.epochs = learning_rate, epochs
        self.rng = np.random.default_rng(seed)

    def fit(self, X, y):
        n = X.shape[0]
        self.intercept_, self.coef_ = 0.0, np.ones(X.shape[1])
        for _ in range(self.epochs):
            for _ in range(n):
                i = self.rng.integers(n)  # one random row
                err = y[i] - (X[i] @ self.coef_ + self.intercept_)
                self.intercept_ -= self.lr * (-2 * err)
                self.coef_ -= self.lr * (-2 * err * X[i])
        return self
```

The fixed seed makes the random choice of observations the same on every run, so the results repeat.

Extra: Picking observations at random can choose some observations twice in an epoch and others not at all. An alternative is to shuffle the observations once per epoch and walk through them in that order, so each observation is used exactly once. scikit-learn's `SGDRegressor` does this by default: its `shuffle` parameter, `True` unless changed, shuffles the training data after each epoch (scikit-learn docs).

4. Faster progress per epoch

Key point: In 40 epochs, SGD reaches test R^2 0.42 on the diabetes data; batch gradient descent only reaches 0.19.

Figure 4 compares the two on the diabetes data, epoch by epoch. Across is the number of epochs (passes over the data); up is the test R^2 , the share of the variation in unseen data that the model explains (1 is perfect). A curve that rises sooner has learned sooner.




Figure 4: Test R^2 after each epoch: batch vs stochastic

After	Batch (learning rate 0.5)	Stochastic (learning rate 0.01)
10 epochs	0.05	0.27
40 epochs	0.19	0.42
OLS	0.44	0.44

SGD updates 353 times per epoch, batch only once. So per epoch, SGD makes far more progress, even though each of its steps is based on a single observation.

Extra: A common trap: for the *same number of epochs*, batch gradient descent is faster, not SGD. In 100 epochs batch makes 100 updates, SGD makes $100 \times n$. SGD wins on time only because it needs far fewer epochs, which shows on large data. Here the Python loop above is also slower per epoch than batch (under a second for 40 epochs against a few milliseconds), because a Python loop over observations is slow and batch uses one fast matrix product. scikit-learn's SGD runs the same loop in compiled code, so it is fast. The real advantage of SGD is elsewhere: one update costs the same however many observations the data has, so SGD is the main way to train linear models on very large

datasets (Goodfellow §5.9).

5. A noisy path

Key point: One observation gives a rough estimate of the whole gradient, so SGD's path zigzags. The noise makes SGD fast and able to escape local minima, but it never settles exactly at the minimum.

5.1 The zigzag

Key point: Each step points roughly, not exactly, downhill.

The gradient from one observation is a noisy estimate of the gradient from all observations. Figure 5 (left) shows the consequence on the 100-point example, on the contour map of the bowl of Figure 2 (right panel: a close-up of the same map near the cross): batch gradient descent (red) walks straight down the bowl, while SGD (orange) wanders but reaches the bottom area within one epoch.




Figure 5: Batch vs stochastic paths, and the effect of a learning schedule

5.2 Advantages of the noise

Key point: The noise speeds SGD up on large data and can shake it out of a local minimum.

- **Large data:** SGD needs only one observation at a time in memory, and it gets close to the answer in few passes.
- **Non-convex losses:** with a loss that has local minima (see [the shape of the loss function](#)), batch gradient descent settles in whichever dip it reaches first. SGD's random

jumps can carry it out of a sharp, shallow local minimum (Kleinberg et al. 2018), like shaking a tray so a marble hops out of a small dent.

Figure 6 shows the escape on a made-up loss with one coefficient w , because linear regression’s own loss has no **local minimum** (G-1110). The average loss has a narrow, shallow dip at $w = -1.49$ and a wide, deep minimum at $w = 1.92$. Each of the 100 observations sees the same curve tilted a little to the left or right, so a step from one observation is the true step plus some noise. Both methods start at $w = -3$ with learning rate 0.05. Watch the red batch marker: it rolls slowly into the shallow dip and stays there. The orange SGD marker is knocked across the dip by its noisy steps and settles in the deep minimum by epoch 7.




Figure 6: Batch gradient descent and SGD on a made-up loss with a shallow local minimum and a deep minimum, both from $w = -3$ with learning rate 0.05. Faint orange dots: SGD’s 100 steps inside the current epoch. Batch ends in the shallow dip; SGD ends in the deep minimum.

The escape needs enough noise. Over 200 random runs, SGD ended in the deep minimum every time; with the noise halved, only 6 of 200 runs escaped (script `local_minimum.py`).

5.3 Disadvantage: jitter at the bottom

Key point: With a constant learning rate, SGD keeps jumping around the minimum instead of stopping at it.

Near the minimum, every new observation still pushes the coefficients in its own direction. So SGD keeps bouncing around the best values (Figure 5, right, orange) and never stops exactly there. Running it twice with different random observations also gives slightly different answers.

6. Learning schedules

Key point: Making the learning rate shrink over time lets SGD move fast at first and settle down near the end.

The fix for the jitter is a **learning schedule** (G-1070): a learning rate that decreases as training goes on. Big steps early on cover distance quickly; small steps later let SGD settle close to the minimum.

A common schedule is

$$\eta_t = \frac{t_0}{t + t_1}$$

where t counts the updates done so far, and t_0 and t_1 are constants that we choose. The constant t_0 sets the size of the start: at $t = 0$ the rate is t_0/t_1 . The constant t_1 sets how long the rate stays near the start: at $t = t_1$ the rate has fallen to half of its start value, $t_0/(2t_1)$. With $t_0 = 5$ and $t_1 = 50$ the start is

$$\frac{5}{50} = 0.1$$

half of it (0.05) is reached at $t = 50$, and:

Update t	0	100	1,000	10,000
Learning rate	0.1	0.033	0.0048	0.0005

Figure 7 plots the same schedule over 10,000 updates; both axes are log scales, so equal distances mean equal ratios, not equal differences. Watch where the green curve crosses the constant rate 0.05: before update 50 the schedule takes bigger steps, and after it the steps keep shrinking.




Figure 7: The learning schedule $5 / (t + 50)$ against a constant rate of 0.05, both axes on a log scale

In Figure 5 (right), the scheduled run (green) stays much closer to the minimum than the constant-rate run (orange): over the last 200 updates its average distance from the best values is 2.1, against 4.7.

Extra: The idea resembles **simulated annealing** (G-1810), an optimisation method named after annealing metal: cooled slowly, a metal settles into a stable, low-energy state. Simulated annealing likewise lowers a “temperature” step by step (Kirkpatrick et al.).

7. SGD in scikit-learn

Key point: `SGDRegressor` (G-1783) trains linear regression with stochastic gradient descent; `learning_rate` and `eta0` set the schedule.

Python: scikit-learn’s `SGDRegressor`.

```
from sklearn.linear_model import SGDRegressor

reg = SGDRegressor(max_iter=100, learning_rate="constant",
                   eta0=0.01, random_state=0)
reg.fit(X_train, y_train)
reg.score(X_test, y_test)    #  $R^2$  0.43
```

The main parameters:

Parameter	Meaning
<code>max_iter</code> (G-115)	the maximum number of epochs
<code>eta0</code> (G-713)	the starting learning rate
<code>learning_rate</code>	the schedule: <code>"constant"</code> , <code>"invscaling"</code> (the default, $\eta_0/t^{0.25}$), <code>"optimal"</code> or <code>"adaptive"</code>
<code>tol</code>	stop early when the loss stops improving by at least this much
<code>random_state</code>	fixes the random order of observations
<code>loss</code>	the loss to minimise: <code>"squared_error"</code> gives linear regression; other losses (such as <code>"huber"</code>) turn the same class into other models
<code>penalty, alpha</code>	regularisation, from penalising large coefficients onwards

`SGDRegressor` is a general class: gradient descent only minimises whatever loss it is given, so one class trains several linear models.

On the diabetes data, `"constant"` with `eta0=0.01` reaches test R^2 0.43 in 97 epochs. The default `"invscaling"` schedule shrinks the rate as $\eta_0/t^{0.25}$, so it needs a larger start: with `eta0=0.2` it reaches 0.45 in 75 epochs, the best of these runs and above OLS's 0.44. OLS gives the lowest error on the training data, not on the test data: coefficients that stop a little short of OLS can score slightly higher on the 89 unseen test patients, so a gap of 0.01 does not mean SGD found a better fit. A shrinking schedule is like a car braking as it nears a parking spot: it must start fast enough to arrive at all.

Figure 8 puts these runs side by side. Each bar is a run of SGD with its starting rate written below it; its height is the test R^2 , so a taller bar is a better model. With the shrinking schedule, a start of 0.01 leaves the model far from the answer (0.16); raising the start to 0.05, 0.1 and 0.2 brings it to the constant-rate result (green) and then just past OLS (grey).




Figure 8: Test R^2 of SGDRegressor on the diabetes data after at most 100 epochs, with the shrinking “invscaling” schedule at four starting rates (bars), against the constant rate 0.01 (dashed green) and OLS (dotted grey).

Extra: With the same small start, `eta0=0.01`, the shrinking schedule stops at test R^2 0.16 after 100 epochs: the steps shrink before the coefficients get near the answer. The same schedule reaches 0.39 after 1,000 epochs and 0.45 after 5,000; switching the shrinking off (`power_t=0`) gives 0.43 in 100 epochs. Starting rates between: 0.05 gives 0.38, 0.1 gives 0.43 (notebook). Schedules need tuning like any other hyperparameter (Goodfellow §8.3.1).

Extra: scikit-learn may print a `ConvergenceWarning` when `max_iter` runs out before its stopping rule (`tol`) is met. The warning is not an error; its own message suggests increasing `max_iter`.

8. When to use which

Key point: Batch for small data and smooth convex losses; SGD for very large data, data that does not fit in memory, and non-convex losses.

	Batch	Stochastic
Observations per update	all	one
Path	smooth, exact gradient	noisy, zigzags
Progress per epoch	slow	fast
Memory	whole dataset	one observation
Final answer	settles at the minimum	jitters around it (use a schedule)

	Batch	Stochastic
Local minima	can get stuck	can escape

Extra: SGD also suits data that keeps arriving. When a new observation comes in, SGD takes one more step from the current coefficients; it does not start again from the beginning (StatQuest). Learning from data as it arrives is **online learning** (G-1391), and scikit-learn's `partial_fit` (G-1458) does it (see [scikit-learn's `partial\_fit`](#)).

[Mini-batch gradient descent](#), taught next, sits between the two and is what most practice uses.

9. Summary

- Batch gradient descent needs about $n \times m$ operations per update: too slow and too memory-hungry for big data.
- SGD updates after each random observation: n updates per epoch, each using one observation's derivatives.
- On the diabetes data, 40 epochs of SGD reach test R^2 0.42, against 0.19 for batch.
- The noise makes SGD fast and able to escape local minima, but it jitters at the bottom.
- A learning schedule (shrinking rate) lets SGD settle, if its starting rate is large enough; `SGDRegressor` offers several.

10. Sources

Built from

- CampusX, “Stochastic Gradient Descent”, YouTube, https://www.youtube.com/watch?v=V7KBAA_g
- StatQuest with Josh Starmer, “Stochastic Gradient Descent, Clearly Explained!!!”, YouTube, <https://www.youtube.com/watch?v=vMh0zPT0tLI>. The intuition of one random observation per step behind Figure 3 (redrawn with our own data), the redundancy argument of Section 3, and the link to online learning.

Other references

- **Goodfellow:** I. Goodfellow, Y. Bengio, A. Courville, *Deep Learning*, MIT Press, 2016 (deeplearningbook.org). §5.9 (SGD for large datasets), §8.3.1 (SGD and decaying learning rates).
- **scikit-learn docs:** scikit-learn documentation, `sklearn.linear_model.SGDRegressor` (version 1.9).
- **Kirkpatrick et al.:** S. Kirkpatrick, C. D. Gelatt, M. P. Vecchi, “Optimization by Simulated Annealing”, *Science* 220(4598), 671–680, 1983.
- **Kleinberg et al. 2018:** R. Kleinberg, Y. Li, Y. Yuan, “An Alternative View: When Does SGD Escape Local Minima?”, *Proceedings of the 35th International Conference on Machine Learning (ICML)*, PMLR 80, 2018.

11. Key terms

Term	Meaning
Feature	An input variable: one column of the data table
Observation	One record: one row of the data table
Target	The output we predict
Stochastic	Involving randomness
Stochastic gradient descent (SGD)	Gradient descent that updates the coefficients after each single random observation

Term	Meaning
Local minimum	A point lower than everything around it, but not the lowest overall
Online learning	Training step by step on data as it arrives
Learning schedule (G-1070)	A rule that changes the learning rate during training, usually shrinking it
SGDRegressor	scikit-learn's linear regression trained with stochastic gradient descent
eta0	The starting learning rate in SGDRegressor
max_iter (G-115)	The maximum number of epochs in SGDRegressor
Simulated annealing	An optimisation method that lowers a “temperature” step by step so the search settles into a low state; the idea behind shrinking learning rates