

Ridge Regression: The Maths and Code from Scratch

Contents

1. Overview	2
2. One feature	3
2.1 The loss	4
2.2 The intercept	5
2.3 The slope	5
2.4 Why the slope shrinks	7
3. Many features	8
3.1 The loss in matrix form	9
3.2 Expanding and differentiating	10
3.3 Solving for w	11
3.4 Not penalising the intercept	11
3.5 By hand on two points	11
4. Summary	16
5. Sources	16
6. Key terms	16

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Standardization ([Note ML-009](#)); Multiple linear regression ([Note ML-052](#)); Normal equation ([Note ML-053](#)); Gradient descent ([Note ML-056](#)); Bias-variance trade-off ([Note ML-061](#)); Regularisation ([Note ML-062](#)).
- **Leads to:** Elastic Net ([Note ML-068](#)).
- **Compare with:** Lasso regression ([Note ML-066](#)); L1 and L2 regularisation in neural networks ([Note DL-026](#)).

1. Overview

Key point: Ridge finds its coefficients the same way plain least squares does: write the loss, find where its slope is zero, solve. The penalty changes the answer in one place only. With one feature, it adds λ to the bottom of the slope fraction; with many features, it adds λ to the diagonal of a matrix. A bigger bottom means a smaller slope.

The previous Note explained what Ridge does: it adds λ times the squared coefficients to the loss, which keeps them small. Here λ (lambda) is the penalty strength, a number we choose, for example $\lambda = 10$. This Note derives the formulas that find the Ridge coefficients, the same way the OLS and normal equation Notes did for plain linear regression.

Here a **feature** (G-772) is an input variable (one column of the data table), an **observation** (G-1374) is one record (one row), and the **target** (G-1949) y is the value we predict. There are two cases:

- **One feature:** a formula for the slope m and the intercept b (Section 2).
- **Many features:** a matrix formula for all coefficients at once (Section 3).

Both are worked by hand on two points, then coded from scratch and checked against scikit-learn's Ridge.

Figure 1 shows what the formulas will produce on the 100-observation example of the previous Note. The ordinary least-squares line has slope 27.83. With $\lambda = 10$ the **Ridge regression** (G-1691) line tilts down to 24.95, and with $\lambda = 100$ to 12.93. All three lines pass through the same point, the point of means; Section 2 shows why.




Figure 1: Ridge lines on the 100-observation example: $\lambda = 0$ (ordinary least squares, slope 27.83), $\lambda = 10$ (24.95) and $\lambda = 100$ (12.93). Every line passes through the point of means.

2. One feature

Key point: The penalty pulls the slope towards 0 like a weight on one side of a seesaw. In the formula, that weight is λ added to the bottom of the OLS slope fraction. The intercept formula does not change.

Figure 2 shows the mechanism before any algebra. With the intercept at its best value for each slope, the loss is a curve in the slope m : the squared error (blue) plus the penalty λm^2 (orange). The penalty is a bowl centred on 0, so adding it pulls the bottom of the sum towards 0. As λ grows from 0 to 100, the bottom slides from 27.83 to 12.93, exactly the value of the formula we derive below.




Figure 2: The Ridge loss of the 100-observation example as a function of the slope m , as λ grows from 0 to 100: squared error (blue), penalty λm^2 (dashed orange) and their sum (green). The dotted line marks the OLS slope 27.83; the green minimum slides towards 0.

Picture a seesaw: the data pulls the slope up, and λ is a weight added to the other side. The heavier the weight, the less the slope rises.

2.1 The loss

Key point: The loss is the usual sum of squared errors plus λ times the slope squared.

With one feature, the prediction for observation i is the slope times its feature value plus the intercept:

$$\hat{y}_i = mx_i + b$$

Here m is the slope, as in the simple linear regression Notes (in the normal equation Note, m counted the features instead). The Ridge loss is

$$L = \sum_{i=1}^n (y_i - mx_i - b)^2 + \lambda m^2$$

The sum runs over the n observations. Only the slope m is penalised, not the intercept b .

We follow two points from the previous Note by hand, $(1, 2)$ and $(3, 5)$, with $\lambda = 1$. Their means are

$$\bar{x} = (1 + 3)/2 = 2 \quad \bar{y} = (2 + 5)/2 = 3.5$$

2.2 The intercept

Key point: The penalty does not contain b , so the derivative with respect to b is the same as in OLS, and so is the answer.

Differentiate L with respect to b and set it to zero. The term λm^2 does not contain b , so it disappears:

$$\frac{\partial L}{\partial b} = -2 \sum_{i=1}^n (y_i - mx_i - b) = 0$$

This derivative is exactly the OLS step, and it gives the same result:

$$b = \bar{y} - m\bar{x}$$

2.3 The slope

Key point: Differentiating with respect to m adds a $2\lambda m$ term, which ends up as $+\lambda$ in the denominator.

Differentiate L with respect to m . The squared-error part gives the OLS terms; the penalty λm^2 adds $2\lambda m$:

$$\frac{\partial L}{\partial m} = -2 \sum_{i=1}^n (y_i - mx_i - b) x_i + 2\lambda m = 0$$

Now solve for m , one step per line, with the two points alongside.

Step 1. Divide by -2 :

$$\sum_{i=1}^n (y_i - mx_i - b) x_i - \lambda m = 0$$

Step 2. Replace b by its formula from Section 2.2:

$$\sum_{i=1}^n (y_i - \bar{y} - m(x_i - \bar{x})) x_i - \lambda m = 0$$

Step 3. Split the sum into the part without m and the part with m :

$$\sum_{i=1}^n (y_i - \bar{y}) x_i - m \sum_{i=1}^n (x_i - \bar{x}) x_i - \lambda m = 0$$

Step 4. Move both m terms to the right and take m out as a common factor:

$$\sum_{i=1}^n (y_i - \bar{y}) x_i = m \left(\sum_{i=1}^n (x_i - \bar{x}) x_i + \lambda \right)$$

Step 5. Replace each lone x_i by the deviation from the mean. The sums do not change, because deviations from a mean add up to 0 (the same trick as in the OLS derivation):

$$\sum_{i=1}^n (y_i - \bar{y}) x_i = \sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x})$$

$$\sum_{i=1}^n (x_i - \bar{x}) x_i = \sum_{i=1}^n (x_i - \bar{x})^2$$

On the two points, the first pair of sums, one term per point:

$$(2 - 3.5) \times 1 + (5 - 3.5) \times 3 = -1.5 + 4.5 = 3$$

$$(2 - 3.5)(1 - 2) + (5 - 3.5)(3 - 2) = 1.5 + 1.5 = 3$$

and the second pair:

$$(1 - 2) \times 1 + (3 - 2) \times 3 = -1 + 3 = 2$$

$$(1 - 2)^2 + (3 - 2)^2 = 1 + 1 = 2$$

Step 6. Divide by the bracket:

$$m = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2 + \lambda}$$

The top of the fraction is the same as in OLS, and the bottom has λ added. With $\lambda = 0$ the formula is the OLS slope.

On the two points with $\lambda = 1$:

$$m = \frac{3}{2 + 1} = 1.0$$

$$b = 3.5 - 1.0 \times 2 = 1.5$$

Without the penalty ($\lambda = 0$) the slope would be the line through both points:

$$m = \frac{3}{2+0} = 1.5$$

The Ridge slope 1.0 is the one the λ sweep of the [previous Note](#) found at $\lambda = 1$.

Extra: Step 5 in symbols. The two sides of the first line differ by one sum, which is 0:

$$\sum_{i=1}^n (y_i - \bar{y}) \bar{x} = \bar{x} \sum_{i=1}^n (y_i - \bar{y}) = \bar{x} \times 0 = 0$$

The same holds for the second line with the deviations of x.

2.4 Why the slope shrinks

Key point: A bigger λ makes the denominator bigger, so the slope gets smaller. The slope approaches 0 but never reaches it.

On the 100-observation example from the previous Note, each sum has 100 terms, so the computer adds them up (the notebook, two lines of NumPy). Input: the 100 pairs (x, y). Output:

$$\sum_{i=1}^{100} (x_i - \bar{x})(y_i - \bar{y}) = 2416.7$$

$$\sum_{i=1}^{100} (x_i - \bar{x})^2 = 86.85$$

So

$$m = \frac{2416.7}{86.85 + \lambda}$$

λ	Slope m	Intercept b
0	$2416.7/86.85 = 27.83$	-2.295
10	$2416.7/96.85 = 24.95$	-2.127
100	$2416.7/186.85 = 12.93$	-1.425

These slopes are exactly the ones scikit-learn gave in the previous Note. The intercept is not penalised, yet it moves too, because its formula from Section 2.2 contains the slope m:

$$b = \bar{y} - m\bar{x}$$

Figure 3 shows the whole curve.

$m = 2416.7 / (86.85 + \lambda)$: the slope shrinks towards 0 but never reaches it




Figure 3: The Ridge slope against λ

However large λ gets, the fraction stays above 0. So Ridge makes coefficients small but never exactly zero. Lasso, two Notes later, behaves differently.

Python: Ridge with one feature, from scratch.

```
class MyRidge:
    def __init__(self, alpha=0.1):
        self.alpha = alpha                # the  $\lambda$  of the formula

    def fit(self, x, y):
        num = np.sum((x - x.mean()) * (y - y.mean()))
        den = np.sum((x - x.mean()) ** 2) + self.alpha
        self.m = num / den
        self.b = y.mean() - self.m * x.mean()
        return self

MyRidge(alpha=10).fit(x, y).m           # 24.955, same as Ridge(alpha=10)
```

3. Many features

Key point: With many features the same recipe runs on matrices: the penalty adds λ to every number on the diagonal of the matrix that the normal equation inverts, except the intercept's.

Figure 4 shows the whole result before the steps. Compared with the normal equation of [Note ML-053](#), only one piece is new: λ times a matrix I that has 1s on its diagonal, with a 0 in the top-left corner so that the intercept is not penalised. Adding it puts λ on the

diagonal, which makes the inverse smaller, just as λ on the bottom of the fraction made the slope smaller in Section 2.

$$w = (X^T X + \lambda I)^{-1} X^T y$$

compared with the normal equation, only $+\lambda I$ is new; it adds λ to the diagonal

$$I = \begin{bmatrix} \mathbf{0} & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & & & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{bmatrix}$$

top-left entry 0: the intercept w_0 is not penalised

Figure 4: The Ridge formula and its identity matrix

3.1 The loss in matrix form

Key point: The penalty, λ times the sum of squared coefficients, is a row of coefficients times the same column.

With many features, the predictions are one matrix product, as in the normal equation Note:

$$\hat{y} = Xw$$

X has a first column of 1s, and w holds the intercept and all coefficients. The vector w is the same coefficient vector that the normal equation Note calls β ; w is the usual letter in Ridge and in the next Note. With two coefficients, for example

$$w = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} 1.5 \\ 1.0 \end{bmatrix}$$

The sum of squared coefficients is the row of coefficients times the column:

$$w^T w = w_0^2 + w_1^2 = 1.5^2 + 1.0^2 = 3.25$$

So the Ridge loss is

$$L = (y - Xw)^T (y - Xw) + \lambda w^T w$$

3.2 Expanding and differentiating

Key point: The squared-error part gives the same derivative as before; the penalty adds $2\lambda w$.

The first part multiplies out exactly as in Section 4 of the normal equation Note:

$$L = y^T y - 2w^T X^T y + w^T X^T X w + \lambda w^T w$$

Section 5.1 of the normal equation Note checked three rules on small numbers. Recap, one per part:

$$\frac{\partial}{\partial w} (y^T y) = 0$$

$$\frac{\partial}{\partial w} (2w^T X^T y) = 2X^T y$$

$$\frac{\partial}{\partial w} (w^T X^T X w) = 2X^T X w$$

The penalty needs one more rule, the matrix version of “the derivative of w^2 is $2w$ ”. Check it with two coefficients. Written out,

$$w^T w = w_0^2 + w_1^2$$

The two partial derivatives:

$$\frac{\partial}{\partial w_0} (w_0^2 + w_1^2) = 2w_0$$

$$\frac{\partial}{\partial w_1} (w_0^2 + w_1^2) = 2w_1$$

Stacked, they are 2 times w . So, with λ in front,

$$\frac{\partial}{\partial w} (\lambda w^T w) = 2\lambda w$$

At $w = (1.5, 1.0)$ and $\lambda = 1$ that is $(3, 2)$. Adding the four derivatives:

$$\frac{\partial L}{\partial w} = -2X^T y + 2X^T X w + 2\lambda w = 0$$

3.3 Solving for w

Key point: Collect the w terms, factor out w, and multiply by the inverse.

Divide by 2 and move the transpose of X times y to the right:

$$X^T X w + \lambda w = X^T y$$

To factor out w, write λw as $\lambda I w$, where I is the **identity matrix** (G-915), the square matrix with 1s on the diagonal and 0s elsewhere. Multiplying by I changes nothing, and a matrix cannot be added to a number, so we need the matrix form:

$$I = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad I w = \begin{bmatrix} 1 \times 1.5 + 0 \times 1.0 \\ 0 \times 1.5 + 1 \times 1.0 \end{bmatrix} = \begin{bmatrix} 1.5 \\ 1.0 \end{bmatrix} = w$$

Now w can be taken out as a common factor:

$$(X^T X + \lambda I) w = X^T y$$

Multiply both sides by the inverse of the bracket:

$$w = (X^T X + \lambda I)^{-1} X^T y$$

Back to Figure 4: compared with the normal equation, only $+\lambda I$ is new. It adds λ to every number on the diagonal of the transpose of X times X.

3.4 Not penalising the intercept

Key point: The top-left entry of I is set to 0, so the intercept is left out of the penalty.

The first entry of w is the intercept. The intercept only shifts predictions up or down, so it should not be penalised (the previous Note). Setting the top-left entry of I to 0 removes it from the penalty, as in Figure 4. With two coefficients:

$$I_0 = \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix} \quad \lambda I_0 w = \lambda \begin{bmatrix} 0 \\ w_1 \end{bmatrix}$$

The penalty now touches only the slope w_1 , as in Section 2.

3.5 By hand on two points

Key point: The matrix formula, run on the two points of Section 2, gives the same intercept 1.5 and slope 1.0.

Take the two points (1, 2) and (3, 5) again, with $\lambda = 1$.

Step 1. Build X with its column of 1s, and y:

$$X = \begin{bmatrix} 1 & 1 \\ 1 & 3 \end{bmatrix} \quad y = \begin{bmatrix} 2 \\ 5 \end{bmatrix}$$

Step 2. The transpose of X times X, one entry per pair of columns:

$$X^T X = \begin{bmatrix} 1+1 & 1+3 \\ 1+3 & 1+9 \end{bmatrix} = \begin{bmatrix} 2 & 4 \\ 4 & 10 \end{bmatrix}$$

Step 3. The transpose of X times y:

$$X^T y = \begin{bmatrix} 2+5 \\ 1 \times 2 + 3 \times 5 \end{bmatrix} = \begin{bmatrix} 7 \\ 17 \end{bmatrix}$$

Step 4. Add λI_0 , which puts $\lambda = 1$ on the slope's diagonal entry only:

$$X^T X + \lambda I_0 = \begin{bmatrix} 2 & 4 \\ 4 & 10+1 \end{bmatrix} = \begin{bmatrix} 2 & 4 \\ 4 & 11 \end{bmatrix}$$

Step 5. Invert with the 2×2 recipe of the normal equation Note (swap the diagonal, flip the other signs, divide by the determinant). The determinant:

$$2 \times 11 - 4 \times 4 = 22 - 16 = 6$$

$$(X^T X + \lambda I_0)^{-1} = \frac{1}{6} \begin{bmatrix} 11 & -4 \\ -4 & 2 \end{bmatrix}$$

Step 6. Multiply by the column (7, 17):

$$w_0 = \frac{11 \times 7 - 4 \times 17}{6} = \frac{77 - 68}{6} = 1.5$$

$$w_1 = \frac{-4 \times 7 + 2 \times 17}{6} = \frac{-28 + 34}{6} = 1.0$$

Intercept 1.5 and slope 1.0: the same as the one-feature formula of Section 2.3.

Check that the slope of the loss is zero there. Half the derivative of Section 3.2, with I_0 in the penalty, is the transpose of X times X times w, plus $\lambda I_0 w$, minus the transpose of X times y:

$$X^T X w = \begin{bmatrix} 2 \times 1.5 + 4 \times 1.0 \\ 4 \times 1.5 + 10 \times 1.0 \end{bmatrix} = \begin{bmatrix} 7 \\ 16 \end{bmatrix}$$

$$\lambda I_0 w = \begin{bmatrix} 0 \\ 1.0 \end{bmatrix}$$

$$\begin{bmatrix} 7 + 0 - 7 \\ 16 + 1.0 - 17 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Both entries are 0, so (1.5, 1.0) is the bottom of the Ridge loss.

Python: Ridge with many features, from scratch.

```
class MyRidgeND:
    def __init__(self, alpha=0.1):
        self.alpha = alpha

    def fit(self, X, y):
        X = np.insert(X, 0, 1, axis=1)           # column of 1s
        I = np.identity(X.shape[1])
        I[0, 0] = 0                               # intercept not penalised
        w = np.linalg.inv(X.T @ X + self.alpha * I) @ X.T @ y
        self.intercept_, self.coef_ = w[0], w[1:]
        return self

    def predict(self, X):
        return X @ self.coef_ + self.intercept_
```

On the diabetes data (test size 0.2, random state 4) with alpha 0.1:

	Intercept	Test R ²
MyRidgeND	150.89	0.4693
Ridge(alpha=0.1, solver="cholesky")	150.89	0.4693

The coefficients agree to within 3×10^{-12} , rounding error.

The closed form has the same limit as the normal equation: inverting the matrix costs about the cube of its size in operations, and building $X^T X$ needs the whole dataset in memory. On large data the next Note trains Ridge with gradient descent instead.

Extra: scikit-learn reaches the same answer differently: it centres X and y (subtracts their means), solves for the coefficients without an intercept column, and then computes the intercept from the means. Centring is the matrix version of $b = \bar{y} - m\bar{x}$. The "cholesky" solver (the **Cholesky solver**, G-383) then solves the same equation $(X^T X + \lambda I)w = X^T y$ directly as a **closed-form solution** (G-398), which is why the two rows of the table

above agree (scikit-learn docs, `Ridge`).

Extra: A bonus of $+\lambda I$: when columns of X are strongly correlated (**multicollinearity**, G-1273), $X^T X$ can be impossible or unstable to invert. Adding a positive λ to the diagonal always makes it invertible, so Ridge has an answer even when plain OLS does not. In fact this was the main reason Ridge was first introduced (ESL §3.4.1; Hoerl and Kennard 1970).

A small case: two observations with the same feature value 2. The two columns of X are then copies of each other up to a factor, and the determinant is 0, so the 2×2 recipe would divide by 0:

$$X = \begin{bmatrix} 1 & 2 \\ 1 & 2 \end{bmatrix} \quad X^T X = \begin{bmatrix} 2 & 4 \\ 4 & 8 \end{bmatrix} \quad 2 \times 8 - 4 \times 4 = 0$$

Adding $\lambda = 1$ to the slope's diagonal entry:

$$X^T X + \lambda I_0 = \begin{bmatrix} 2 & 4 \\ 4 & 9 \end{bmatrix} \quad 2 \times 9 - 4 \times 4 = 2$$

Now the inverse exists. In general, for any non-zero vector v the bracket gives a positive number, so no v is sent to zero and the matrix has an inverse:

$$v^T (X^T X + \lambda I) v = \|Xv\|^2 + \lambda \|v\|^2 > 0$$

The same line covers data with more features than observations, where $X^T X$ never has an inverse: Ridge still returns one answer (ISL §6.2.1).

Figure 5 measures the effect on the diabetes training data. With $\lambda = 0$ the smallest eigenvalue of $X^T X$ is only 0.0073, and the **condition number** (G-441), the largest eigenvalue divided by the smallest, is 48,311: the inverse exists but small rounding errors are blown up. Adding λ lifts the smallest eigenvalue by λ , and at $\lambda = 0.1$ the condition number is already down to about 3,290.




Figure 5: The diabetes training data with its column of 1s. Left: the smallest eigenvalue of $X^T X + \lambda I$ rises with λ from 0.0073. Right: the condition number falls from 48,311 at $\lambda = 0$. Both axes use log scales.

4. Summary

	Linear regression (OLS)	Ridge
Loss	$\sum (y_i - \hat{y}_i)^2$	$\sum (y_i - \hat{y}_i)^2 + \lambda \sum w_j^2$
Slope, one feature	$\frac{\sum (x - \bar{x})(y - \bar{y})}{\sum (x - \bar{x})^2}$	$\frac{\sum (x - \bar{x})(y - \bar{y})}{\sum (x - \bar{x})^2 + \lambda}$
Intercept, one feature	$\bar{y} - m\bar{x}$	$\bar{y} - m\bar{x}$
Many features	$(X^T X)^{-1} X^T y$	$(X^T X + \lambda I)^{-1} X^T y$

- Ridge adds λ to the denominator of the slope (one feature) or to the diagonal of $X^T X$ (many features).
- A larger λ gives smaller coefficients, but never exactly 0.
- The intercept is not penalised: the top-left entry of I is 0.
- The from-scratch code matches scikit-learn's **Ridge** exactly.

5. Sources

Built from

- CampusX, “Ridge Regression Part 2 | Mathematical Formulation & Code from scratch | Regularized Linear Models”, YouTube, https://www.youtube.com/watch?v=oDIZBQjk_3A

Other references

- **ESL:** Hastie, T., Tibshirani, R. and Friedman, J. *The Elements of Statistical Learning*, 2nd ed. Springer, 2009. Section 3.4.1, p. 64 (which credits Hoerl and Kennard, 1970).
- **Hoerl and Kennard (1970):** Hoerl, A. E. and Kennard, R. W. “Ridge regression: biased estimation for nonorthogonal problems.” *Technometrics* 12(1), 55–67.
- **ISL:** James, G., Witten, D., Hastie, T. and Tibshirani, R. *An Introduction to Statistical Learning*, 2nd ed. Springer, 2021. Section 6.2.1 (Ridge when there are more features than observations).
- **scikit-learn docs:** `sklearn.linear_model.Ridge` (`fit_intercept`, `solver`=“cholesky”), scikit-learn 1.9 documentation; centring step in `sklearn/linear_model/_base.py` (`_preprocess_data`).

6. Key terms

Term	Meaning
Identity matrix	A square matrix with 1s on the diagonal and 0s elsewhere; multiplying by it changes nothing
Closed-form solution	A formula that gives the answer directly, without repeated steps
Cholesky solver	A scikit-learn Ridge solver that solves the closed-form equation directly

Term	Meaning
Ridge regression (G-1691)	Linear regression with λ times the sum of squared coefficients added to the loss
Determinant (G-598)	For a 2×2 matrix, the diagonal product minus the off-diagonal product; the inverse divides by it