

Ridge Regression with Gradient Descent

Contents

1. Overview	2
2. The gradient	4
2.1 The loss with a factor of $\frac{1}{2}$	4
2.2 Expanding and differentiating	4
2.3 The update rule, and one update on numbers	5
3. Seeing the paths	8
4. Code from scratch	10
4.1 Why the learning rate has a limit	10
5. Gradient descent stops early	14
5.1 Early stopping rescues an overfitting model	14
5.2 Why the slow directions stay small	15
6. Ridge with gradient descent in scikit-learn	16
7. Summary	19
8. Sources	19
9. Key terms	19

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Standardization ([Note ML-009](#)); Multiple linear regression ([Note ML-052](#)); Normal equation ([Note ML-053](#)); Gradient descent ([Note ML-056](#)); Bias-variance trade-off ([Note ML-061](#)); Regularisation ([Note ML-062](#)).
- **Leads to:** Elastic Net ([Note ML-068](#)).
- **Compare with:** Lasso regression ([Note ML-066](#)); L1 and L2 regularisation in neural networks ([Note DL-026](#)).

1. Overview

Key point: Ridge can also be trained step by step, like plain linear regression. Each step does two things: it first pulls every coefficient a little towards 0, then takes the ordinary downhill step.

The [Ridge formula](#) finds the Ridge coefficients in one jump. Like plain linear regression, **Ridge regression** (G-1691; linear regression with a penalty on large coefficients, see [the idea of Ridge](#)) can instead be trained with **gradient descent** (G-862; [the idea](#)): start somewhere, and repeatedly step downhill on the loss (the number that says how wrong the model is).

Here a **feature** (G-772) is an input variable (one column of the data table), an **observation** (G-1374) is one record (one row), and the **target** (G-1949) y is the value we predict. Gradient descent is useful when there are many features, because the formula needs the inverse of a large matrix (see [the cost of the inverse](#)).

The next figures use a [contour map](#) of the squared error, so first the surface it comes from. Take the one-feature example of [Ridge with one feature](#): every pair (slope m , intercept b) is one line, and its squared error (the sum of the squared gaps between the line and the 100 targets) is

$$E(m, b) = \sum_{i=1}^{100} (y_i - mx_i - b)^2$$

where $\sum_{i=1}^{100}$ means “add the term for observation 1, 2, up to 100”. For the start line $m = -5$, $b = 20$ this is about 162,800 and for the best line $m = 27.8$, $b = -2.3$ about 28,300. E is a smooth bowl over the (m, b) floor (the same bowl as in [both parameters together](#) and [how stochastic gradient descent works](#)); here its height is drawn as $\log E$. Figure 1 shows the bowl and the same bowl seen from above. Each line of the map joins points at the same height; lines close together mean a steep slope, and the centre ring is the lowest point. The orange route is the Ridge step of Figure 2 (start, then shrink, then new w).




Figure 1: The squared-error surface of the one-feature example (left) and the same surface seen from above, a contour map (right). Orange: the Ridge step of Figure 2, from the start $(-5, 20)$ through the shrink to the new point $(11.2, 9.8)$.

Figure 2 shows one Ridge step on the one-feature example (100 observations), starting from slope -5 and intercept 20 . Watch the two arrows: the red one shrinks the slope to half its size, towards 0 ; the blue one is the ordinary least-squares step downhill. The intercept is never shrunk. Section 2.3 computes this exact step with numbers.




Figure 2: One Ridge step in two parts, on the squared-error contours: shrink the slope by $1 - \eta\lambda = 0.5$ (red), then take the ordinary least-squares step (blue). The intercept is not shrunk

This Note:

- derives the Ridge gradient step by step and runs one update on numbers (Section 2);
- watches the paths for three penalty strengths (Section 3);
- codes it from scratch and explains the learning-rate limit (Section 4);
- shows that stopping early is a regulariser too (Section 5);
- shows the scikit-learn options (Section 6).

2. The gradient

Key point: The Ridge gradient is the plain linear regression gradient plus one extra piece, λ times the coefficients. That extra piece is what pulls them towards 0.

2.1 The loss with a factor of $\frac{1}{2}$

Key point: Multiplying the loss by $\frac{1}{2}$ does not move its minimum; it only removes a 2 from the derivative.

The symbols, with the values of Figure 2:

- $\mathbf{w}$ is the column of coefficients, intercept first. [The normal equation](#) called the same column β ; Ridge texts and the code use $\mathbf{w}$. In Figure 2 the start is

$$\mathbf{w}_{\text{old}} = \begin{bmatrix} b \\ m \end{bmatrix} = \begin{bmatrix} 20 \\ -5 \end{bmatrix}$$

- $\mathbf{X}$ is the data with a first column of 1s, one row per observation (100 rows here), and $\mathbf{y}$ the column of targets.
- λ is the penalty strength, here $\lambda = 100$.

The Ridge loss (see [the loss in matrix form](#)), multiplied by $\frac{1}{2}$:

$$L = \frac{1}{2}(\mathbf{X}\mathbf{w} - \mathbf{y})^\top(\mathbf{X}\mathbf{w} - \mathbf{y}) + \frac{1}{2}\lambda \mathbf{w}^\top \mathbf{w}$$

Halving every value of the loss leaves the lowest point at the same $\mathbf{w}$. The $\frac{1}{2}$ just makes the derivative tidier, as the 2 from differentiating a square cancels the $\frac{1}{2}$.

2.2 Expanding and differentiating

Key point: The same expansion as the previous Note, then the same derivative rules, one term at a time.

Expanding the product as in [expanding the error](#):

$$L = \frac{1}{2}\mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - \mathbf{w}^\top \mathbf{X}^\top \mathbf{y}$$

$$+\frac{1}{2}y^T y + \frac{1}{2}\lambda w^T w$$

Differentiate one term per line with respect to w . The rules were checked on small numbers in [three rules of matrix calculus](#) and [expanding and differentiating](#):

$$\begin{aligned}\frac{\partial}{\partial w} \left(\frac{1}{2} w^T X^T X w \right) &= \frac{1}{2} \times 2 X^T X w \\ &= X^T X w\end{aligned}$$

$$\frac{\partial}{\partial w} \left(-w^T X^T y \right) = -X^T y$$

$$\frac{\partial}{\partial w} \left(\frac{1}{2} y^T y \right) = 0$$

$$\begin{aligned}\frac{\partial}{\partial w} \left(\frac{1}{2} \lambda w^T w \right) &= \frac{1}{2} \lambda \times 2 w \\ &= \lambda w\end{aligned}$$

Adding the four lines:

$$\frac{\partial L}{\partial w} = X^T X w - X^T y + \lambda w$$

The first two terms are the linear regression gradient. The Ridge penalty adds only λw .

Extra: Setting this gradient to zero gives $(X^T X + \lambda I)w = X^T y$, the formula of [solving for \$w\$](#) . Gradient descent walks towards the same answer step by step instead of jumping to it.

2.3 The update rule, and one update on numbers

Key point: New w = old w minus the learning rate times the gradient.

As in [gradient descent with many features](#), all coefficients are updated together:

$$w_{\text{new}} = w_{\text{old}} - \eta \left(X^T X w_{\text{old}} - X^T y + \lambda w_{\text{old}} \right)$$

where η is the **learning rate** (G-1068; [the size of each step](#)); here $\eta = 0.005$. The first entry of w is the intercept, which is not penalised. So, as in [not penalising the intercept](#), the λw term uses $\lambda I_0 w$, where I_0 is the identity matrix with its top-left entry set to 0:

$$I_0 = \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix}$$

One update, line by line. The two matrix pieces add up 100 observations, so the computer builds them (the notebook). Input: the 100 pairs (x, y). Output:

$$X^T X = \begin{bmatrix} 100 & 5.841 \\ 5.841 & 87.186 \end{bmatrix}$$

$$X^T y = \begin{bmatrix} -66.938 \\ 2412.823 \end{bmatrix}$$

The 100 counts the observations, 5.841 is the sum of the feature values, 87.186 the sum of their squares, -66.938 the sum of the targets and 2412.823 the sum of feature times target. The rest is by hand, starting from $w = (20, -5)$.

Step 1. The transpose of X (X^T , X with rows and columns swapped) times X, times w, one row of the result per block:

$$\begin{aligned} 100 \times 20 + 5.841 \times (-5) \\ = 2000 - 29.204 \\ = 1970.796 \end{aligned}$$

$$\begin{aligned} 5.841 \times 20 + 87.186 \times (-5) \\ = 116.815 - 435.931 \\ = -319.116 \end{aligned}$$

Step 2. Subtract the transpose of X times y:

$$1970.796 - (-66.938) = 2037.734$$

$$-319.116 - 2412.823 = -2731.939$$

Step 3. Add the penalty $\lambda I_0 w$. It is 0 for the intercept, and for the slope:

$$100 \times (-5) = -500$$

So

$$\text{gradient} = \begin{bmatrix} 2037.734 + 0 \\ -2731.939 - 500 \end{bmatrix}$$

$$\text{gradient} = \begin{bmatrix} 2037.734 \\ -3231.939 \end{bmatrix}$$

Step 4. Multiply by the learning rate:

$$0.005 \times 2037.734 = 10.189$$

$$0.005 \times (-3231.939) = -16.160$$

Step 5. Subtract from the old w :

$$b_{\text{new}} = 20 - 10.189 = 9.811$$

$$m_{\text{new}} = -5 - (-16.160) = 11.160$$

The new point (slope 11.16, intercept 9.81) is the “new w ” of Figure 2, and the first step of the $\lambda = 100$ path in Section 3.

The same step in two parts. Group the λw term with the old w instead:

$$w_{\text{new}} = (1 - \eta\lambda) w_{\text{old}} - \eta (X^T X w_{\text{old}} - X^T y)$$

For the slope, the two parts are the two arrows of Figure 2. The red shrink:

$$1 - \eta\lambda = 1 - 0.005 \times 100 = 0.5$$

$$0.5 \times (-5) = -2.5$$

The blue least-squares step, from Step 2:

$$-0.005 \times (-2731.939) = +13.660$$

$$-2.5 + 13.660 = 11.160$$

The same 11.160 as Step 5. The intercept has no shrink, only the least-squares step.

So every step first shrinks the coefficients by the factor $1 - \eta\lambda$, then takes the ordinary linear regression step. The per-step shrinking is why the L2 penalty is also called **weight decay** (G-2108), the name used for neural networks (ESL §3.4.1).

3. Seeing the paths

Key point: Gradient descent finds the bottom of whichever loss it is given. With $\lambda = 100$, the bottom sits at a smaller slope.

Figure 3 runs gradient descent on the one-feature example of [Ridge with one feature](#) for three values of λ , from the same start point ($m = -5$, $b = 20$, learning rate 0.005). The contours belong to the plain squared error, so only $\lambda = 0$ heads for their centre. Watch where each path stops: the larger λ , the closer to $m = 0$, and the flatter its line on the data (right).




Figure 3: Gradient descent with $\lambda = 0, 100$ and 250 from the same start. Left: the paths on the squared-error contours; the dashed curve holds every Ridge answer as λ grows from 0. Right: each path's current line on the data

With $\lambda = 250$ the path zig-zags across the valley: the penalty makes the bowl so steep in the m direction that a step of 0.005 overshoots, and each overshoot is smaller than the last (Section 4.1 explains why, with numbers). Figure 4 draws the bowls themselves for $\lambda = 0$ (the surface of Figure 1) and $\lambda = 100$ (top row: surfaces; bottom row: the same surfaces seen from above, the contour maps) over 60 steps.




Figure 4: Gradient descent on the loss with $\lambda = 0$ and $\lambda = 100$

- $\lambda = 0$: the path ends at slope 27.8 and intercept -2.3 , the linear regression answer.
- $\lambda = 100$: the penalty moves the bottom of the bowl to slope 12.9 and intercept -1.4 .
These values are exactly the Ridge values from [the formula](#).

The penalty makes the bowl steeper in the m direction only. The steepness of a bowl along one direction is its **curvature** (G-521), the second derivative. Along the slope m it is the bottom-right entry of the transpose of X times X , plus λ :

$$\frac{\partial^2 L}{\partial m^2} = \sum_{i=1}^{100} x_i^2 + \lambda$$

$$\lambda = 0 : \quad 87.19 + 0 = 87.19$$

$$\lambda = 100 : \quad 87.19 + 100 = 187.19$$

Along the intercept b the curvature is the count of observations, 100, whatever λ is.

4. Code from scratch

Key point: The loop computes the gradient with matrix products and updates all coefficients at once.

Python: Ridge with batch gradient descent.

```
class MyRidgeGD:
    def __init__(self, epochs, learning_rate, alpha):
        self.epochs, self.lr, self.alpha = (epochs,
                                             learning_rate, alpha)

    def fit(self, X, y):
        X = np.insert(X, 0, 1, axis=1)           # column of 1s
        w = np.insert(np.ones(X.shape[1] - 1), 0, 0.0)
        I = np.identity(X.shape[1])
        I[0, 0] = 0                               # no penalty on intercept
        for _ in range(self.epochs):
            grad = X.T @ X @ w - X.T @ y + self.alpha * I @ w
            w = w - self.lr * grad
            self.intercept_, self.coef_ = w[0], w[1:]
        return self

    def predict(self, X):
        return X @ self.coef_ + self.intercept_
```

On the diabetes data (test size 0.2, random state 4) with alpha 0.001, learning rate 0.005 and 500 epochs, test R^2 is 0.474 and the intercept 150.87.

4.1 Why the learning rate has a limit

Key point: A step that is too big jumps across the valley and lands higher up the other side. Each jump is then bigger than the last, and the coefficients blow up. The steeper the bowl, the smaller the step must be.

This gradient adds up the errors of all observations instead of averaging them, so its size grows with the number of observations. The summed gradient is why the learning rate must be tiny here.

Figure 5 shows the idea along one direction of a bowl, with the curvature of the steepest direction of the diabetes data, 353. Both runs start at distance 1 from the bottom. Watch

where each step lands. With learning rate 0.005 (left) every step jumps across the bottom but lands closer. With 0.006 (right) every step jumps across and lands farther away.




Figure 5: One direction of the loss bowl, curvature 353. Left: learning rate 0.005, each step multiplies the distance from the bottom by -0.765 , so the jumps shrink. Right: learning rate 0.006, each step multiplies it by -1.118 , so the jumps grow

Why each step multiplies the distance. Along one direction, call s the distance from the bottom and h the curvature. Near the bottom the loss is a parabola, and its slope is the curvature times the distance:

$$L = \frac{1}{2} h s^2$$

$$\frac{dL}{ds} = h s$$

One gradient descent step, then the distance taken out as a common factor:

$$s_{\text{new}} = s - \eta h s$$

$$s_{\text{new}} = (1 - \eta h) s$$

On numbers, with $h = 353$. Learning rate 0.005:

$$1 - 0.005 \times 353$$

$$= 1 - 1.765 = -0.765$$

$$s : \quad 1 \rightarrow -0.765 \rightarrow 0.585 \rightarrow -0.448 \rightarrow 0.342$$

The minus sign is the jump across the bottom; the size 0.765 below 1 makes each jump shorter. Learning rate 0.006:

$$1 - 0.006 \times 353$$

$$= 1 - 2.118 = -1.118$$

$$s : \quad 1 \rightarrow -1.118 \rightarrow 1.250 \rightarrow -1.397 \rightarrow 1.562$$

Now the size 1.118 is above 1, so each jump is longer than the last.

The limit. The distance shrinks only when the factor lies between -1 and 1 :

$$-1 < 1 - \eta h < 1$$

The right-hand side holds for any positive η and h . The left-hand side, with ηh moved across and the 1 subtracted:

$$\eta h < 2$$

$$\eta < \frac{2}{h}$$

With $h = 353$:

$$\eta < \frac{2}{353} = 0.0057$$

So 0.005 is just below the limit and 0.006 just above it.

Many directions at once. A bowl over many coefficients has a different curvature in each direction. The matrix of all its second derivatives is the **Hessian matrix** (G-888). For the Ridge loss it is

$$H = X^T X + \lambda I$$

For the one-feature example of Section 3 with $\lambda = 100$ (and the intercept left out of the penalty), its diagonal holds the two curvatures of Section 3:

$$H = \begin{bmatrix} 100 & 5.84 \\ 5.84 & 187.19 \end{bmatrix}$$

The steepest and the flattest directions of the bowl are the eigenvectors of H , and the curvatures along them are its **eigenvalues** (G-665). Eigenvalues are often written λ ; here λ is already the penalty, so we write h . The computer finds them (the notebook). Input: the matrix H . Output for the three paths of Section 3, with the step factor $1 - 0.005 h$ for the steepest direction:

λ	Eigenvalues h of H	Factor for the largest h
0	84.9 and 102.3	0.49
100	99.6 and 187.6	0.06
250	99.9 and 337.3	-0.69

With $\lambda = 250$ the factor is negative, so the path jumps across the valley each step, but its size 0.69 is below 1, so the jumps shrink: the zig-zag of Figure 3.

Extra: The same result in matrix form. Write w^* for the exact answer, where the gradient is zero:

$$Hw^* = X^T y$$

The gradient at any w is then H times the gap from the answer:

$$Hw - X^T y = Hw - Hw^*$$

$$= H(w - w^*)$$

One update, with the exact answer subtracted from both sides:

$$w_{\text{new}} - w^*$$

$$= (w - w^*) - \eta H(w - w^*)$$

$$= (I - \eta H)(w - w^*)$$

Along each eigenvector of H with eigenvalue h , the matrix in front multiplies by $1 - \eta h$, as in the one-direction case. So the steps shrink only if η is below $2/h$ for the largest h .

On the diabetes training set (353 observations) the largest eigenvalue is 353, so the limit is $2/353 = 0.0057$. With 0.006 the steps overshoot and grow: after 100 epochs the largest coefficient is about 15,000.

Figure 6 runs the code above at both learning rates. For about 50 epochs the two runs look alike; then the run with 0.006 bends upwards, because the overshoot along the steepest direction is multiplied every step by

$$|1 - 0.006 \times 353| = 1.12$$

and finally dominates.




Figure 6: The largest coefficient of the from-scratch code per epoch on the diabetes data, alpha 0.001. Learning rate 0.005 (blue, below the limit $2/353$) settles; 0.006 (red, above it) grows to about 15,000 after 100 epochs. Log scale.

5. Gradient descent stops early

Key point: Stopping gradient descent early keeps the coefficients small, so **early stopping** (G-656) is a regulariser of its own, much like the ridge penalty.

Think of pouring water into an ice-cube tray: stop pouring early and every cube is only partly full. Gradient descent starts with small coefficients and fills them up step by step; stopping early leaves the slow ones small.

5.1 Early stopping rescues an overfitting model

Key point: On data where linear regression overfits badly, stopping early helps about as much as the ridge penalty.

We give the model 65 features (the 10 diabetes measurements plus their squares and pairwise products) and train on 200 patients, as in [early stopping](#). With so many features for so

few observations, plain least squares has high variance, the case where shrinkage helps (ISL §6.2.1). Averaged over 50 random splits, test R^2 is:

Model	Test R^2
Linear regression (no regularisation)	0.06
Ridge, alpha chosen by cross-validation on the training data	0.44
Gradient descent without a penalty, stopped early	0.40

Both regularisers rescue the badly overfitting linear regression, and by a similar amount. For a linear model, stopping early and adding the ridge penalty are close cousins (Goodfellow et al. §7.8).

Extra: In the 65-feature experiment both models use standardised features. The stopping epoch is chosen on a validation part of the training data, and ridge's alpha by cross-validation on the training data, so the test data is used only once. Early stopping is one of the most common regularisers in deep learning (Goodfellow et al. §7.8).

5.2 Why the slow directions stay small

Key point: Gradient descent creeps along flat directions of the loss bowl, so after a few hundred steps the coefficients in those directions are still small.

Back on the 10 diabetes features (353 training observations), the exact Ridge answer from [the formula](#) scores test R^2 0.463 (R^2 is the share of the variation in the targets that the model explains: 1 is perfect, 0 is no better than the mean). Figure 7 tracks gradient descent for up to a million epochs.




Figure 7: Test R^2 and distance from the exact answer, per epoch

Epochs	Test R^2	Largest gap from the exact coefficients
100	0.390	897
500	0.474	823
50,000	0.464	104
1,000,000	0.463	0.0

- **Left:** test R^2 rises, peaks around 500 epochs, then settles at the exact answer's 0.463.
- **Right:** the coefficients need about a million epochs to reach the exact values.

The slow direction comes from s_1 and s_2 , which are strongly correlated (see [many features: the diabetes data](#)). The bowl is almost flat along “raise s_2 , lower s_1 ”, so gradient descent creeps along it, and stopping early leaves the coefficients in that direction smaller (Goodfellow et al. §7.8). With 353 observations least squares barely overfits (see [the diabetes data](#)), so on this data the size of the test- R^2 change says little; the 65-feature experiment above is the real test.

Extra: The numbers behind this. Along an eigenvector of the Hessian H (Section 4.1) with eigenvalue h , the error is multiplied by $1 - \eta h$ per step. When ηh is tiny, the error needs about $1/(\eta h)$ steps to shrink to about a third. The smallest eigenvalue here is 0.0083, so, one line at a time:

$$\eta h = 0.005 \times 0.0083 = 0.0000415$$

$$1 - \eta h = 0.9999585 \text{ per step}$$

$$\frac{1}{\eta h} = \frac{1}{0.0000415} \approx 24,000 \text{ steps}$$

Its eigenvector is mostly s_1 (-0.71) and s_2 ($+0.55$). After 500 epochs the gap from the exact answer, as a vector, has length 1,153, and 1,146 of that lies along this one direction; and the coefficient vector is about half as long as the exact one (774 against 1,455).

6. Ridge with gradient descent in scikit-learn

Key point: Use `SGDRegressor` with `penalty="l2"`, or `Ridge` with an iterative solver.

Python: Two scikit-learn options.

```
from sklearn.linear_model import SGDRegressor, Ridge

sgd = SGDRegressor(penalty="l2", alpha=0.001, max_iter=500,
                   eta0=0.1, learning_rate="constant",
                   random_state=0)
sgd.fit(X_train, y_train)
sgd.score(X_test, y_test)           #  $R^2$  0.450
```

```
ridge = Ridge(alpha=0.001, max_iter=500, solver="sparse_cg")
ridge.fit(X_train, y_train)
ridge.score(X_test, y_test)          # R2 0.463
```

- **SGDRegressor** (G-1783; [SGD in scikit-learn](#)) adds the Ridge penalty with `penalty="l2"` (G-1477). Its `alpha` is the penalty strength.
- **Ridge** has a `solver` (G-1836) setting. `"cholesky"` and `"svd"` use the formula directly. `"sparse_cg"`, `"lsqr"`, `"sag"` and `"saga"` reach the answer step by step, and `max_iter` limits the steps.

Ridge has no learning-rate setting at all, even with an iterative solver (scikit-learn docs, **Ridge**). Here `sparse_cg` reaches the exact answer, 0.463.

Extra: **SGDRegressor** averages the loss over the observations, while **Ridge** adds it up. The **SGDRegressor** objective is (scikit-learn user guide §1.5.8)

$$\frac{1}{n} \sum \frac{1}{2} (y_i - \hat{y}_i)^2 + \alpha \cdot \frac{1}{2} \|w\|^2$$

Multiplying by 2n:

$$\sum (y_i - \hat{y}_i)^2 + n\alpha \|w\|^2$$

That is the **Ridge** loss with `alpha` equal to `n` times α . So the same `alpha` means a stronger penalty in **SGDRegressor**: its `alpha` equals **Ridge**'s `alpha` divided by the number of observations. With 353 training observations, `SGDRegressor(alpha=0.001)` matches **Ridge** with `alpha`

$$353 \times 0.001 = 0.353$$

Run long, its coefficients land within 3 of that model's, but up to 840 away from `Ridge(alpha=0.001)`.

Figure 8 shows the three sets of coefficients side by side. The orange bars of `SGDRegressor(alpha=0.001)` sit on the blue bars of `Ridge(alpha=0.353)`, not on the grey bars of `Ridge(alpha=0.001)`; the biggest gap is in `sl`.




Figure 8: Coefficients on the diabetes training data (353 observations): Ridge with alpha 0.001 (grey) and 0.353 (blue), and SGDRegressor with alpha 0.001 run for 50,000 epochs (orange). The orange bars match the blue ones.

7. Summary

Method	How it finds w	Test R^2 here
Formula (<code>Ridge</code> , <code>cholesky</code>)	one step, $(X^T X + \lambda I)^{-1} X^T y$	0.463
<code>Ridge</code> , <code>sparse_cg</code>	iterative solver	0.463
<code>MyRidgeGD</code> , 500 epochs	batch gradient descent, stopped early	0.474
<code>SGDRegressor</code> , <code>l2</code>	stochastic gradient descent	0.450

- The Ridge gradient is $X^T X w - X^T y + \lambda w$: the linear regression gradient plus λw .
- Each update shrinks the coefficients by $1 - \eta \lambda$ first, hence the name weight decay.
- The intercept is left out of the penalty.
- Stopping gradient descent early also regularises: with many features for few observations it rescues an overfitting fit about as well as ridge.
- In scikit-learn: `SGDRegressor(penalty="l2")` or `Ridge` with an iterative solver.

8. Sources

Built from

- CampusX, “Ridge Regression Part 3 | Gradient Descent | Regularized Linear Models”, YouTube, https://www.youtube.com/watch?v=Fci_wwMp8G8

Other references

- **ESL:** Hastie, T., Tibshirani, R. and Friedman, J. *The Elements of Statistical Learning*, 2nd ed. Springer, 2009. Section 3.4.1, p. 63.
- **Goodfellow et al.:** Goodfellow, I., Bengio, Y. and Courville, A. *Deep Learning*. MIT Press, 2016. Section 7.8, “Early Stopping”.
- **scikit-learn docs:** `sklearn.linear_model.Ridge`; user guide Section 1.5.8, “Mathematical formulation” (SGD), scikit-learn 1.9.
- **ISL:** James, G., Witten, D., Hastie, T. and Tibshirani, R. *An Introduction to Statistical Learning*, 2nd ed. Springer, 2021. Section 6.2.1, p. 241.

9. Key terms

Term	Meaning
Weight decay	Another name for the L2 penalty: each gradient step shrinks the coefficients by a fixed factor
Solver	The method a scikit-learn model uses to find its coefficients
penalty	<code>SGDRegressor</code> setting that adds a regularisation penalty, such as “l2” for <code>Ridge</code>

Term	Meaning
Curvature (G-521)	How steep the loss bowl is along one direction: its second derivative there
Hessian matrix (G-888)	The matrix of all second derivatives of the loss; for Ridge it is $X^T X + \lambda I$
Eigenvalue (G-665)	The curvature of the bowl along one of its main directions (an eigenvector of the Hessian)