

Logistic Regression: Maximum Likelihood and the Log Loss

Contents

1. Overview	2
2. Comparing two models	2
3. Maximum likelihood	3
3.1 The probability of the true class	3
3.2 Multiply	4
4. From products to sums: the log	5
4.1 The problem with products	5
4.2 Taking logs	5
4.3 Changing the sign	5
5. The cost of one point	6
5.1 Why not the squared error?	7
6. One formula for both classes	8
6.1 Choosing the right probability automatically	9
6.2 The loss function	9
7. Minimising the loss	10
7.1 The search in a picture	10
7.2 Gradient descent	11
8. Summary	14
9. Sources	14
10. Key terms	14

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Classification problems ([Note ML-003](#)); Model-based learning ([Note ML-006](#)); Gradient descent ([Note ML-056](#)); Polynomial features ([Note ML-060](#)); Perceptron trick ([Note ML-069](#)); Sigmoid function ([Note ML-071](#)).
- **Leads to:** Softmax regression ([Note ML-078](#)); Gradient boosting ([Note ML-114](#)); Cost-sensitive learning ([Note ML-127](#)); ANN for classification ([Note DL-011](#)).
- **Compare with:** Support vector machines ([Note ML-086](#)); Hinge loss and soft margin ([Note ML-088](#)); Entropy, information gain and Gini ([Note ML-091](#)).

1. Overview

Key point: To find the best decision boundary, we need a loss function: one number that says how good a boundary is. For logistic regression it comes from maximum likelihood, and it is called binary cross entropy or log loss.

The sigmoid perceptron of the previous Note improved the **decision boundary** (G-555): the line where the model's probability is exactly 0.5, with one class predicted on each side. But we still had no way to say which decision boundary is **best**, or when to stop nudging. Both perceptron versions pick random points and nudge the decision boundary, without a number that measures how good the line is.

Machine learning normally works differently, as in linear regression:

1. Define a **loss function** (G-706): a formula that measures how wrong a model is.
2. Find the coefficients where the loss is smallest, with a formula or with gradient descent.

This Note builds the loss function for **logistic regression** (G-1120). The next Notes minimise it with gradient descent.

2. Comparing two models

Key point: When one model misclassifies points and the other does not, the better one is obvious. When both are perfect, or both make mistakes, we need a number.

Imagine four points, two green and two red, and two candidate decision boundaries. If model 1 misclassifies points and model 2 separates them perfectly, model 2 is clearly better.

But when two models both classify everything correctly, or both make one mistake, it is not obvious which is better. A loss function settles this: it gives each model a number, and the model with the better number wins.




Figure 1: The four points of the table in Section 3.1 under the two models. Each label is the probability the model gives to the point’s true colour; a cross marks a point on the wrong side of the decision boundary (black line). The background shades from red to green as $P(\text{green})$ grows.

In Figure 1, watch the two crosses: model 1’s decision boundary leaves points 2 and 3 on the wrong side, while model 2’s separates all four.

3. Maximum likelihood

Key point: For each point, take the probability the model gave to its true class. Multiply these over all points. The model with the larger product is better.

3.1 The probability of the true class

Key point: For a green point use $P(\text{green}) = \hat{y}$; for a red point use $P(\text{red}) = 1 - \hat{y}$.

Each point is one **observation** (G-1374) (one record, a row of the data table). Its position is given by its **features** (G-772) (the input values, such as the two coordinates), and its colour is the **target** (G-1949) (the output we predict). The sigmoid turns every point’s position into a probability: $\hat{y} = \sigma(w \cdot x)$ is the probability of green (the positive class), and $1 - \hat{y}$ is the probability of red.

For each point we take the probability the model gives to the colour the point **really** has. Suppose the two models give these values:

Point	True colour	Model 1: P(true colour)	Model 2: P(true colour)
1	green	0.7	0.7
2	red	0.4	0.6
3	green	0.4	0.7
4	red	0.8	0.6

A value above 0.5 means the point is on its correct side; below 0.5, on the wrong side. Model 1 misclassifies points 2 and 3.

3.2 Multiply

Key point: The product is the likelihood: how probable the model finds the data we actually see.

Treating the points as independent, the probability that the model produces exactly these colours is the product of the four values. This product is the **likelihood** (G-1086):

- Model 1: $0.7 \times 0.4 \times 0.4 \times 0.8 = 0.090$
- Model 2: $0.7 \times 0.6 \times 0.7 \times 0.6 = 0.176$

Model 2 has the higher likelihood, so it is the better model. **Maximum likelihood estimation** (G-1191) means choosing the coefficients that make this product as large as possible.




Figure 2: The probabilities and the costs of each point under the two models

In Figure 2, the left bars are the four factors of each product: model 1 has two factors of 0.4 where model 2 has 0.6 and 0.7, so model 1's product is the smaller one.

4. From products to sums: the log

Key point: Multiplying thousands of numbers below 1 gives a number too small for a computer. Taking logs turns the product into a sum, which is safe.

4.1 The problem with products

Key point: 10,000 probabilities of 0.7 multiply to about 10^{-1549} .

With 4 points the product is already only 0.09. A real dataset may have 10,000 observations. If each point got probability 0.7, the product would be $0.7^{10,000}$, about 10^{-1549} . Such a number is far below the smallest number a standard float can store (about 10^{-308}). The product **underflows** (G-2036): the computer stores the tiniest float it has (5×10^{-324}) or 0 instead of the true value, whatever the model, and can no longer tell two models apart.

An everyday picture: halve a sheet of paper again and again. After a few dozen cuts the piece is too small to see, and every model's product ends up as the same invisible scrap.

4.2 Taking logs

Key point: $\log(a \times b) = \log a + \log b$, so the log of the product is the sum of the logs.

The logarithm turns multiplication into addition:

$$\log(a \times b) = \log a + \log b$$

So instead of the likelihood, we compute its log, the **log-likelihood** (G-1113). Every log in this Note is the natural log (base e), which is what `np.log` computes; any other base gives different numbers but the same best model.

$$\log(0.7 \times 0.4 \times 0.4 \times 0.8) = \log 0.7 + \log 0.4 + \log 0.4 + \log 0.8 = -2.41$$

For 10,000 points of 0.7 the sum is simply $10,000 \times \log 0.7 = -3,567$, an ordinary number. Since log grows whenever its input grows, the model with the larger likelihood also has the larger log-likelihood: the comparison is unchanged.

4.3 Changing the sign

Key point: Logs of numbers below 1 are negative. Putting a minus sign in front gives a positive number to minimise: the cross entropy.

Every probability is between 0 and 1, and the log of such a number is negative. So the log-likelihood is always negative. Putting a minus sign in front makes it positive:

Model	Likelihood	Log-likelihood	Cross entropy
1	0.090	−2.41	2.41
2	0.176	−1.74	1.74

This negative log-likelihood is the **cross entropy** (G-505). Because of the sign change, the direction flips too: we **maximise** the likelihood but **minimise** the cross entropy. The better model, model 2, has the smaller cross entropy.




Figure 3: From probabilities to cross entropy

Figure 3 lays out the whole chain for model 1, with the number at each step.

5. The cost of one point

Key point: Each point costs $-\log p$. A confident correct prediction costs almost nothing; a confident wrong one costs a lot.

The cross entropy is a sum of one cost per point, $-\log p$, where p is the probability given to the true class. Figure 4 shows how this cost behaves.




Figure 4: The cost $-\log p$ of one point

- $p = 0.9$ (confident and correct): cost 0.11.
- $p = 0.7$: cost 0.36.
- $p = 0.4$ (on the wrong side): cost 0.92.
- $p = 0.1$ (confident and wrong): cost 2.30, and it grows without limit as p approaches 0.

So the loss punishes confident mistakes very heavily, and keeps rewarding the model a little for making correct points even more certain. The second effect matters for points close to the decision boundary: a correct point with $p = 0.6$ still costs 0.51, so lowering the loss means moving the decision boundary further away from it.

5.1 Why not the squared error?

Key point: For a badly wrong prediction, the log loss curve is far steeper than the squared error curve, so the model gets a much stronger push to correct itself.

Linear regression used the squared error, so a natural question is why logistic regression does not. For one point, the squared error would be $(1 - p)^2$: the gap between the probability given to the true class and 1, squared.

Figure 5 slides one point from a good prediction to a bad one and draws both costs. Watch the two black **tangent lines** (G-1945): the steepness of each tangent is the slope of that cost.




Figure 5: The cost of one point under log loss (red) and squared error (blue) as the probability p given to the true class falls from 0.95 to 0.03. The squared error never exceeds 1 and its slope never exceeds 2 in size. The log loss grows without limit, and at $p = 0.1$ its slope is already -10 against -1.8 . Idea after StatQuest, “Neural Networks Part 6: Cross Entropy”

p	Log loss: cost	Log loss: slope $-1/p$	Squared error: cost	Squared error: slope $-2(1 - p)$
0.9	0.11	-1.1	0.01	-0.2
0.5	0.69	-2.0	0.25	-1.0
0.1	2.30	-10.0	0.81	-1.8
0.03	3.51	-33.3	0.94	-1.9

Gradient descent (section 7) takes steps whose size depends on the slope of the loss. With the squared error, a terrible prediction ($p = 0.03$) gives almost the same slope as a mediocre one ($p = 0.1$). With log loss, the worse the prediction, the steeper the slope, so the worst mistakes are corrected with the largest steps.

Extra: On perfectly separable data, gradient descent on this loss slowly turns the decision boundary towards the one with the widest gap (Soudry et al. 2018).

6. One formula for both classes

Key point: $L = -(1/n) \sum [y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$. For each point, the y factors switch on the correct term.

6.1 Choosing the right probability automatically

Key point: $-[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$ equals $-\log \hat{y}$ for green points and $-\log(1 - \hat{y})$ for red points.

The cost of a point uses $\hat{y}$ if the point is green and $1 - \hat{y}$ if it is red. Writing it as $-\log \hat{y}$ for every point would be wrong for the red ones. One expression handles both cases:

$$\text{cost of point } i = -[y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

- **Green point** ($y_i = 1$): the second term is multiplied by 0 and vanishes, leaving $-\log \hat{y}_i$.
- **Red point** ($y_i = 0$): the first term vanishes, leaving $-\log(1 - \hat{y}_i)$.




Figure 6: The one formula draws two cost curves: green points pay $-\log \hat{y}$, red points pay $-\log(1 - \hat{y})$. Model 1's four points sit on their own curves; their costs add up to 2.41.

In Figure 6, watch where the curves rise: a green point is expensive when $\hat{y}$ is near 0, a red point when $\hat{y}$ is near 1, so each point is punished only for leaning towards the wrong colour.

With numbers, for model 1: point 2 is red with $\hat{y} = 0.6$, so its cost is $-\log(1 - 0.6) = -\log 0.4 = 0.92$. Point 3 is green with $\hat{y} = 0.4$, so its cost is $-\log 0.4 = 0.92$.

6.2 The loss function

Key point: Average the costs over all n points.

Summing over all points and dividing by n to get an average gives the loss function of logistic regression:

$$L = -\frac{1}{n} \sum_{i=1}^n [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)], \quad \hat{y}_i = \sigma(w \cdot x_i)$$

This loss is called **binary cross entropy** or **log loss** (G-303). For model 1 it is $2.41/4 = 0.603$.

Python: Log loss by hand and in scikit-learn.

```
from sklearn.metrics import log_loss

y = np.array([1, 0, 1, 0])          # green = 1, red = 0
y_hat = np.array([0.7, 0.6, 0.4, 0.2]) # model 1's P(green)

-np.mean(y * np.log(y_hat) + (1 - y) * np.log(1 - y_hat)) # 0.6031
log_loss(y, y_hat)                                         # 0.6031
```

7. Minimising the loss

Key point: There is no formula for the best w , so logistic regression is trained with gradient descent.

The task is now precise: find the coefficients w that make L as small as possible. For linear regression's squared error, setting the derivative to zero gave a formula (OLS). For log loss there is no such **closed-form solution** (G-398), a formula that gives the answer in one step, because w sits inside the sigmoid and the logs (Bishop §4.3.3).

7.1 The search in a picture

Key point: Try a curve, score it with the log loss, change the curve a little so that the score improves, and repeat until the score stops falling.

Without a formula, the best model has to be searched for. Figure 7 shows the search on one feature: the CGPA of 30 students and whether each was placed. A candidate model is an S-shaped sigmoid curve giving $P(\text{placed})$ for every CGPA.




Figure 7: Thirty students (green: placed, blue: not placed) and a candidate sigmoid curve. Each red bar is the gap between a student and the curve: 1 minus the probability the curve gives to the student’s true class. From curve to curve the bars shrink and the log loss falls from 1.190 to 0.138, where it stops falling. Idea after StatQuest, “Logistic Regression Details Pt 2: Maximum Likelihood”; the data and curves are our own

Step by step, in Figure 7:

1. **Curve 1** slopes the wrong way: it gives high-CGPA students a low probability of being placed. The red bars are long, and the log loss is 1.190.
2. **Curve 5** slopes the right way but is nearly flat. Every student gets a probability near 0.5, so the bars are still long: log loss 0.444.
3. **Curve 11** is steeper. Most bars are short: log loss 0.218.
4. **Curve 28** is the last one. The only long bars belong to the four students between CGPA 6.0 and 6.1, where placed and not-placed students overlap, and no curve can make all four short. The log loss is 0.138 and no longer falls.

The curve with the smallest log loss is the maximum likelihood model of section 3.

7.2 Gradient descent

Key point: Gradient descent decides how to change the curve at each try: it steps in the direction in which the loss falls fastest.

The method that chooses each next curve is **gradient descent** (G-862): compute the derivative of L with respect to w and step downhill repeatedly. Figure 8 runs it on the four points of section 2.




Figure 8: Gradient descent on the log loss of the four points, starting from model 1 (dotted line). Left: the decision boundary ($P(\text{green}) = 0.5$) turns; the labels give each point's probability for its true colour. Right: the log loss per step, with model 2's 0.434 dashed. Key frames: steps 0, 3 and 40.

In Figure 8, watch the loss pass model 2's value by step 3 and keep falling while the decision boundary turns to put all four points on their correct side. On these four perfectly separable points the loss keeps shrinking towards 0, as Section 5's Extra notes. The next Note works

out the derivative of the sigmoid, and the one after derives the gradient and codes logistic regression from scratch.

8. Summary

Step	Result for model 1
Probability of each point's true class	0.7, 0.4, 0.4, 0.8
Multiply: likelihood (maximise)	0.090
Take logs: log-likelihood	-2.41
Change sign: cross entropy (minimise)	2.41
Divide by n: log loss	0.603

- The perceptron versions had no measure of “best”; a loss function provides one.
- Maximum likelihood: choose the model that gives the observed classes the highest probability.
- Logs avoid tiny products; the minus sign gives a positive loss to minimise.
- Compared with the squared error, log loss has a much steeper slope for badly wrong predictions, so they are corrected with larger steps.
- Log loss: $L = -\frac{1}{n} \sum [y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$; it has no closed-form minimum.

9. Sources

Built from

- CampusX, “Logistic Regression Part 4 | Loss Function | Maximum Likelihood | Binary Cross Entropy”, YouTube, <https://www.youtube.com/watch?v=6bXOo0sxY5c>
- StatQuest with Josh Starmer, “Logistic Regression Details Pt 2: Maximum Likelihood”, YouTube, <https://www.youtube.com/watch?v=BfKanl1aSG0>
- StatQuest with Josh Starmer, “Neural Networks Part 6: Cross Entropy”, YouTube, <https://www.youtube.com/watch?v=6ArSys5qHAU>

Other references

- **Bishop:** Bishop, C. M. *Pattern Recognition and Machine Learning*. Springer, 2006. Section 4.3.3, p. 207.
- **Soudry et al. 2018:** Soudry, D., Hoffer, E., Nacson, M. S., Gunasekar, S. and Srebro, N. “The Implicit Bias of Gradient Descent on Separable Data.” *Journal of Machine Learning Research* 19(70), 1–57, 2018.

10. Key terms

Term	Meaning
Decision boundary	The line (or surface) where the model's probability is exactly 0.5; one class is predicted on each side
Loss function (G-706)	A formula that measures how wrong a model's predictions are

Term	Meaning
Likelihood	The product, over all points, of the probabilities the model gives to their true classes
Maximum likelihood estimation (MLE)	Choosing the parameters that make the likelihood as large as possible
Log-likelihood	The log of the likelihood: the sum of the log probabilities
Cross entropy	The negative log-likelihood; smaller is better
Binary cross entropy (log loss)	The average cross entropy for two classes, the loss function of logistic regression
Underflow	A number too close to 0 for the computer to store, which then becomes 0 or loses precision