

Logistic Regression: Gradient Descent and Code from Scratch

Contents

1. Overview	2
2. The idea on one weight	2
2.1 A model with one weight	2
2.2 Where the slope comes from: one student	3
3. Predictions in matrix form	4
3.1 One prediction	4
3.2 All predictions together	4
4. The loss in matrix form	5
5. The gradient	6
5.1 The first part	6
5.2 The second part	6
5.3 Putting them together	6
6. The update rule	7
7. The code	9
8. Checking against scikit-learn	10
9. Perfectly separable data	11
10. Summary	13
11. Sources	13
12. Key terms	13

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Classification problems ([Note ML-003](#)); Model-based learning ([Note ML-006](#)); Gradient descent ([Note ML-056](#)); Polynomial features ([Note ML-060](#)); Perceptron trick ([Note ML-069](#)); Sigmoid function ([Note ML-071](#)).
- **Leads to:** Softmax regression ([Note ML-078](#)).
- **Compare with:** Support vector machines ([Note ML-086](#)).

1. Overview

Key point: The gradient of the log loss is $-(1/m) X^T(y - \hat{y})$. Gradient descent repeats $w \leftarrow w + \eta (1/m) X^T(y - \hat{y})$, and reaches the same answer as scikit-learn.

The log loss Note gave logistic regression its loss function and noted that it has no closed-form minimum. This Note finishes the journey:

1. watch gradient descent work on a model with a single weight;
2. write the predictions and the loss in matrix form;
3. differentiate the loss with respect to the **weights** (G-407; the coefficients, the numbers the model learns);
4. code gradient descent and check it against scikit-learn.

2. The idea on one weight

Key point: Start with any weight. Measure the slope of the loss there. Step downhill by learning rate $\times$ slope. Far from the minimum the slope is large and the steps are big; close to it the slope is small and the steps are tiny.

Gradient descent (G-862) was built step by step in the [gradient descent Note](#). The recap in one line: the slope of the loss curve tells which way is downhill and how far away the bottom is, and each step moves the weight by the **learning rate** (G-1068) times that slope.

2.1 A model with one weight

To see the steps, take the smallest possible logistic regression: one **feature** (G-772) (an input variable), one weight w , no intercept. The data is 30 students; the feature is $x = \text{CGPA} - 6$ and the **target** (G-1949) (the output we predict) is $y = 1$ for placed and 0 for not placed. The model is $\hat{y} = \sigma(wx)$.

Figure 1 runs gradient descent on this model with learning rate 4, starting from $w = 0$. Watch the red ball on the left: its jumps shrink as the curve flattens. On the right, the S-curve of the current w steepens until it fits the students.




Figure 1: Gradient descent on one weight. Left: the log loss against w , with the ball at the current weight; each step is the learning rate (4) times the slope under the ball. Right: the curve $\sigma(wx)$ for the current w over the 30 students (green: placed, blue: not placed). The steps shrink from 1.640 to 0.020 over 60 steps while the log loss falls from 0.693 to 0.152. The picture of a ball taking big steps far from the minimum and small steps near it follows StatQuest, “Gradient Descent, Step-by-Step”; the data and model are our own

Step	w before	Log loss	Slope of the loss	Step = $-4 \times$ slope	w after
1	0	0.693	-0.410	1.640	1.640
2	1.640	0.300	-0.119	0.476	2.116
3	2.116	0.253	-0.082	0.328	2.445
4	2.445	0.229	-0.064	0.254	2.699

The slope is negative, so the loss falls as w grows, and every step moves w to the right. After 60 steps the slope is -0.005 and the step is down to 0.020.

2.2 Where the slope comes from: one student

The slope in the table is an average over the 30 students. One student shows what goes into it. Take a student with CGPA 6.8 who was placed: $x = 0.8$, $y = 1$. At $w = 0$:

- Predict.** $\hat{y} = \sigma(0 \times 0.8) = \sigma(0) = 0.5$.
- Error.** $y - \hat{y} = 1 - 0.5 = 0.5$. The prediction is too low.
- This student's slope.** $-(y - \hat{y})x = -0.5 \times 0.8 = -0.4$.
- This student's vote for the step.** $-4 \times (-0.4) = +1.6$: raise w , which raises $\hat{y}$ for this student.

Every student gives such a slope, $-(y_i - \hat{y}_i)x_i$. Their average at $w = 0$ is the -0.410 of the table. Sections 3 to 5 show where the expression $-(y - \hat{y})x$ comes from and write it for any number of features at once.

3. Predictions in matrix form

Key point: All m predictions at once: $\hat{y} = \sigma(Xw)$, where X has a first column of 1s.

3.1 One prediction

Key point: For one observation, $\hat{y} = \sigma(w_0 + w_1x_1 + \dots + w_nx_n)$.

Take a dataset with m **observations** (G-1374) (records, one row of the data table each) and n features (one column each). The target y is the class, 0 or 1. The model has $n + 1$ weights: w_0 (the **intercept** (G-960), the part of the prediction that does not depend on any feature) and w_1 to w_n . The prediction for observation i is

$$\hat{y}_i = \sigma(w_0 + w_1x_{i1} + w_2x_{i2} + \dots + w_nx_{in})$$

3.2 All predictions together


Key point: Stack the rows: the inner sums are the matrix product Xw .

Writing out all m predictions, each inner sum is one row of X (with a 1 added in front) times the weight vector w . So all of them together are a single matrix product:

$$\hat{y} = \sigma(Xw)$$

where σ is applied to each entry. Stacking the predictions as Xw is the same trick as in multiple linear regression.

one row per point;
first column all 1s



$$\frac{\partial L}{\partial w} = -\frac{1}{m} X^T (y - \hat{y})$$

X^T is $(n+1) \times m$, $(y - \hat{y})$ is $m \times 1$:
the gradient is $(n+1) \times 1$, one slope
per weight, the same shape as w

Figure 2: The shapes of X , w , $\hat{y}$ and the gradient

In Figure 2, read the top row left to right: the m rows of X times the one column w give one z per observation, and the sigmoid turns each z into a prediction.

4. The loss in matrix form

Key point: $L = -(1/m)[y^T \log \hat{y} + (1 - y)^T \log(1 - \hat{y})]$.

The log loss from the previous Notes is

$$L = -\frac{1}{m} \sum_{i=1}^m [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

The sum $\sum_i y_i \log \hat{y}_i$ multiplies matching entries of two vectors and adds them up, which is the **dot product** (G-634) $y^T \log \hat{y}$. So

$$L = -\frac{1}{m} [y^T \log \hat{y} + (1 - y)^T \log(1 - \hat{y})], \quad \hat{y} = \sigma(Xw)$$

5. The gradient

Key point: Differentiate each part with the chain rule; the sigmoid derivative cancels the logs, leaving $-(1/m) X^T(y - \hat{y})$.

5.1 The first part

Key point: The derivative of $y \log \hat{y}$ with respect to w is $y(1 - \hat{y}) x$.

For a single observation, differentiate $y \log \hat{y}$ step by step with the **chain rule** (G-371) (multiply the derivatives of the nested steps):

- the derivative of $\log \hat{y}$ with respect to $\hat{y}$ is $1/\hat{y}$;
- the derivative of $\hat{y} = \sigma(z)$ with respect to z is $\hat{y}(1 - \hat{y})$ (the sigmoid derivative Note);
- the derivative of $z = w \cdot x$ with respect to w is x .

Multiplying: $y \cdot \frac{1}{\hat{y}} \cdot \hat{y}(1 - \hat{y}) \cdot x = y(1 - \hat{y}) x$. The $\hat{y}$ cancels.

5.2 The second part

Key point: The derivative of $(1 - y) \log(1 - \hat{y})$ is $-(1 - y) \hat{y} x$.

In the same way, for $(1 - y) \log(1 - \hat{y})$:

- the derivative of $\log(1 - \hat{y})$ with respect to $\hat{y}$ is $-1/(1 - \hat{y})$;
- the rest of the chain is the same: $\hat{y}(1 - \hat{y}) x$.

Multiplying: $(1 - y) \cdot \frac{-1}{1 - \hat{y}} \cdot \hat{y}(1 - \hat{y}) \cdot x = -(1 - y)\hat{y} x$.

5.3 Putting them together

Key point: $y(1 - \hat{y}) - (1 - y)\hat{y}$ simplifies to $y - \hat{y}$.

Adding the two parts:

$$y(1 - \hat{y}) - (1 - y)\hat{y} = y - y\hat{y} - \hat{y} + y\hat{y} = y - \hat{y}$$




Figure 3: The cancellation along z . The log’s slope (orange, dashed) blows up where $\hat{y}$ nears 0 or 1, and the sigmoid’s slope (blue, dotted) shrinks there. Their product (green) is the plain $y - \hat{y}$.

In Figure 3, watch the green product stay between -1 and 1 even where the orange slope shoots off the chart: the two chain-rule factors cancel exactly.

So for one observation, the derivative of the bracket is $(y - \hat{y})x$. Summing over all observations, with the $-\frac{1}{m}$ in front, and writing the sum as a matrix product:

$$\frac{\partial L}{\partial w} = -\frac{1}{m} X^\top (y - \hat{y})$$

$X^\top$ has shape $(n + 1) \times m$ and $(y - \hat{y})$ has shape $m \times 1$, so the **gradient** (G-863) has shape $(n + 1) \times 1$: one slope per weight, the same shape as w (Figure 2).

Extra: This gradient has exactly the same form as the one for linear regression with mean squared error: $X^\top(\text{actual} - \text{predicted})$, scaled. The only difference is that the predictions now pass through the sigmoid (Bishop §4.3.2).

6. The update rule

Key point: $w_{\text{new}} = w_{\text{old}} + \eta (1/m) X^\top (y - \hat{y})$.

Gradient descent steps against the gradient, as in section 2. The step size η is the learning rate:

$$w_{\text{new}} = w_{\text{old}} - \eta \frac{\partial L}{\partial w} = w_{\text{old}} + \frac{\eta}{m} X^T (y - \hat{y})$$

Compare it with the sigmoid perceptron of the sigmoid Note: $w \leftarrow w + \eta(y - \hat{y})x$ for one random point. The new rule is the same idea, averaged over **all** points in each step. The perceptron was one-point **stochastic gradient descent** (G-1892) on this very loss, run for a fixed number of steps. **Batch gradient descent** (G-264) uses all points in every step and, run to the end, reaches the true minimum.

	Sigmoid perceptron	Batch gradient descent
Update	$w \leftarrow w + \eta (y_i - \hat{y}_i) x_i$	$w \leftarrow w + \eta \frac{1}{m} \sum_i (y_i - \hat{y}_i) x_i$
Points used per update	one, picked at random	all m , averaged

Each point pulls the weights with a strength equal to its error $y_i - \hat{y}_i$. Figure 4 draws that strength as the size of each point on the data of section 8.




Figure 4: The 100 points of section 8 during batch gradient descent; the area of each circle shows the size of the point’s error $y - \hat{y}$, and the orange line is the decision boundary. At the start many points have a large error and the boundary moves a lot per epoch. By epoch 5,000 only the points in the overlap still have a large error; their pulls balance out, and the boundary stops

1. **Epoch 0:** many circles are large (average error size 0.23). Their average pull is strong, so the boundary moves quickly.
2. **Epoch 20:** points far from the boundary on their correct side have shrunk to dots: a small error means a weak pull.

3. **Epoch 5,000:** the large circles left are the points in the overlap, on the wrong side or close to the boundary (average error size 0.09). The pulls from the two classes cancel in the average, so the update is close to zero.




Figure 5: Log loss on the 100 points of section 8 against the number of updates, both starting from $w = (1, 1, 1)$. Orange: the sigmoid perceptron’s rule, one random point per update (learning rate 0.1). Blue: the batch rule, all points per update (learning rate 0.5).

In Figure 5, watch the orange curve jitter above the dashed minimum (0.1383 after 5,000 updates) while the blue curve settles onto it (0.1366): each one-point step pulls towards a single point, the batch step towards all of them.

7. The code

Key point: Add a column of 1s, then repeat: predict with the sigmoid, and update with $X^T(y - \hat{y}) / m$.

Python: Logistic regression with batch gradient descent.

```
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def gd(X, y, lr=0.5, epochs=5000):
    X = np.insert(X, 0, 1, axis=1)           # column of 1s
    w = np.ones(X.shape[1])
    for _ in range(epochs):
        y_hat = sigmoid(X @ w)               # all predictions
        w = w + lr * (X.T @ (y - y_hat)) / X.shape[0]
```

```
return w[1:], w[0]
```

```
# coef_, intercept_
```

8. Checking against scikit-learn

Key point: After 5,000 epochs, the weights match scikit-learn's to three decimals.

One **epoch** (G-696) is one pass over all the training points; batch gradient descent makes one update per epoch.

The test data has 100 points with two features whose classes overlap a little (`make_classification, class_sep=1.5, random state 4`). scikit-learn's `LogisticRegression(penalty=None)` fits the same model without regularisation, so the two should agree.




Figure 6: Batch gradient descent on the 100 points. Left: the decision boundary ($p = 0.5$) turns into place; the dashed line is scikit-learn's. Right: the log loss per epoch (log scale) falls to scikit-learn's minimum.

In Figure 6, watch the orange line, the **decision boundary** (G-555) where $p = 0.5$, swing from its start at $w = (1, 1, 1)$ onto scikit-learn's dashed decision boundary while the loss curve flattens onto the dashed minimum.

	Intercept	w_1	w_2	Log loss
Start	1	1	1	
After 100 epochs	0.360	0.454	2.850	0.1484
After 1,000 epochs	-1.227	-0.094	3.542	0.1368
After 5,000 epochs	-1.568	-0.220	3.623	0.1366
scikit-learn	-1.569	-0.220	3.623	0.1366

The loss falls quickly at first and then levels off at scikit-learn's minimum (Figure 6, right). The decision boundary turns into place and ends on top of scikit-learn's (left). The model classifies 92% of the points correctly; the rest lie in the overlap, where no straight decision boundary can be perfect.

Python: The scikit-learn equivalent.

```
from sklearn.linear_model import LogisticRegression

lor = LogisticRegression(penalty=None)    # no regularisation
lor.fit(X, y)
lor.intercept_, lor.coef_                # -1.569, [-0.220, 3.623]
```

Extra: By default, `LogisticRegression` adds an L2 (Ridge) penalty with `C=1`, so its weights are slightly smaller than these. `penalty=None` switches it off to match our code; the hyperparameters Note covers the options.

9. Perfectly separable data

Key point: If a straight decision boundary separates the classes perfectly, the loss can always be made smaller by making the weights larger, so the weights never settle.

On the data of the perceptron Notes, where the classes are perfectly separated (`class_sep=20`), something different happens. This case is called **perfect separation** (G-1487):

Epochs	Weights (w_0, w_1, w_2)	Log loss
500	(3.63, 3.35, 0.11)	0.0052
5,000	(5.83, 4.84, 0.21)	0.00063
50,000	(8.20, 6.52, 0.37)	0.00007

The weights keep growing, and the loss keeps shrinking towards 0 without ever reaching it.




Figure 7: Batch gradient descent on perfectly separated points. Left: the pale band is where the model is unsure ($0.1 < p < 0.9$); the orange line is the decision boundary, $p = 0.5$. Right: the length of the weight vector per epoch (log scale).

In Figure 7, the decision boundary turns only slowly, yet the weights keep growing and the unsure band keeps narrowing: the model grows ever more confident instead of settling.

The reason: once every point is on its correct side, multiplying all the weights by 2 keeps the same decision boundary but makes every z twice as large, so every $\hat{y}$ moves closer to its correct 0 or 1. The loss always decreases, so there is no finite minimum (Bishop §4.3.2). The decision boundary itself still turns slowly: its slope $-w_1/w_2$ goes from -30 to -23 to -18 in the table. Gradient descent turns it towards the decision boundary with the widest gap, but only very slowly (Soudry et al. 2018).

In practice this is handled by stopping after a fixed number of epochs, or by adding **regularisation** (G-1659), a penalty on large weights in the loss, the standard fix (Bishop §4.3.2). scikit-learn uses a penalty by default.

10. Summary

- Each step moves the weights by the learning rate times the slope of the loss: big steps far from the minimum, tiny steps near it.
- Predictions: $\hat{y} = \sigma(Xw)$.
- Loss: $L = -\frac{1}{m}[y^\top \log \hat{y} + (1 - y)^\top \log(1 - \hat{y})]$.
- Gradient: $\frac{\partial L}{\partial w} = -\frac{1}{m}X^\top(y - \hat{y})$, thanks to $\sigma' = \sigma(1 - \sigma)$.
- Update: $w \leftarrow w + \frac{\eta}{m}X^\top(y - \hat{y})$; the code matches `LogisticRegression(penalty=None)`.
- On perfectly separable data the weights grow without limit; regularisation or early stopping keeps them finite.

11. Sources

Built from

- CampusX, “Logistic Regression Part 5 | Gradient Descent & Code From Scratch”, YouTube, <https://www.youtube.com/watch?v=ABrrSwMYWSg>
- StatQuest with Josh Starmer, “Gradient Descent, Step-by-Step”, YouTube, <https://www.youtube.com/watch?v=sDv4f4s2SB8>

Other references

- **Bishop:** Bishop, C. M. *Pattern Recognition and Machine Learning*. Springer, 2006. Section 4.3.2, pp. 205–207.
- **Soudry et al. 2018:** Soudry, D., Hoffer, E., Nacson, M. S., Gunasekar, S. and Srebro, N. “The Implicit Bias of Gradient Descent on Separable Data.” *Journal of Machine Learning Research* 19(70), 1–57, 2018.

12. Key terms

Term	Meaning
Decision boundary	The line where the model’s probability is exactly 0.5; one class is predicted on each side
Learning rate (η)	The step size η that scales each gradient descent update
Epoch	One pass over all the training points
Gradient (G-863)	The vector of slopes of the loss, one for each weight
Batch gradient descent	Gradient descent that uses all observations for every update
penalty=None	LogisticRegression setting that switches regularisation off

Term	Meaning
Perfect separation	When a straight decision boundary splits the training classes with no mistakes; unregularised weights then grow without limit