


Classification Metrics: Precision, Recall and F1 Score

Contents

1. Overview	2
2. Precision	3
2.1 A spam filter	3
2.2 The formula	4
3. Recall	5
3.1 A cancer detector	5
3.2 The formula	5
3.3 Choosing between them	6
4. F1 score	7
4.1 When neither mistake clearly matters more	7
4.2 Why the harmonic mean	8
5. In scikit-learn	9
6. More than two classes	10
6.1 Per class	10
6.2 Combining the classes	11
7. Summary	13
8. Sources	13
9. Key terms	13

Where this fits: Pipeline map step 9 (Evaluate). See the [Course map](#).



- **Builds on:** Imbalanced data ([Note ML-009](#)); Confusion matrix ([Note ML-075](#)).
- **Leads to:** ROC curve and AUC ([Note ML-077](#)).

1. Overview

Key point: Precision asks “of everything predicted positive, how much really is?”. Recall asks “of everything really positive, how much did we catch?”. F1 combines them into one number.

[Accuracy can be misleading](#), especially on imbalanced data (one class is far rarer than the other), and it hides which kind of mistake a model makes. This Note introduces three metrics built from the **confusion matrix** (G-449; the table of counts of right and wrong predictions, see [reading it](#)) that fix this: **precision** (G-1547), **recall** (G-1641) and the **F1 score** (G-743).

The confusion matrix of a two-class problem has four counts. The **positive** class is the one we look for (spam, disease); the other is **negative**:

	Predicted positive	Predicted negative
Actually positive	TP (true positive, G-2021): positive, and called positive	FN (false negative, G-747): positive, but called negative
Actually negative	FP (false positive, G-748): negative, but called positive	TN (true negative, G-2020): negative, and called negative

Figure 1 shades the part of this table that each metric reads: precision reads one column, recall reads one row.




Figure 1: Which part of the confusion matrix precision and recall use

2. Precision

Key point: Precision = $TP / (TP + FP)$. Use it when false positives are the expensive mistake.

2.1 A spam filter

Key point: Two spam filters with the same accuracy can differ a lot in false positives.

Two team members each build a spam filter, tested on 1,000 emails. Spam is the positive class.

	TP	FP	FN	TN	Accuracy
Model A	100	100	100	700	0.80
Model B	100	10	190	700	0.80

Both have accuracy 0.80, so accuracy cannot choose between them. They differ in the kind of mistake:

- **False positive:** a normal email sent to the spam folder.
- **False negative:** a spam email left in the inbox.

A missed job offer in the spam folder is far worse than an extra advert in the inbox. So here false positives are the expensive mistake, and we want the model with fewer of them: model B.

2.2 The formula

Key point: Of all the emails the model called spam, the fraction that really were spam.

$$\text{precision} = \frac{TP}{TP + FP}$$

In words: of everything the model **predicted** positive, how much really was positive? Precision uses the “predicted 1” column of the confusion matrix (Figure 1, left).

- Model A: half of what it calls spam is not spam.

$$100/(100 + 100) = 0.50$$

- Model B:

$$100/(100 + 10) = 0.91$$

Model B has the higher precision, matching the choice above. Figure 2 shows where the two numbers come from: precision reads only the outlined “predicted spam” column, and model B’s column holds far fewer false positives.




Figure 2: The confusion matrices of the two spam filters, both with accuracy 0.80. Precision reads the outlined column, everything predicted spam: 100 of 200 for model A (0.50), 100 of 110 for model B (0.91).

Extra: The precision formula has no TN in it. Suppose model A were tested on 100 times more normal email, with TN 70,000 instead of 700 and nothing else changed. Its precision would still be:

$$100/200 = 0.50$$

So when real positives are rare and negatives are plentiful, as with a rare disease, precision still reports honestly on the positive predictions (StatQuest, “ROC and AUC, Clearly Explained!”).

3. Recall

Key point: Recall = TP / (TP + FN). Use it when false negatives are the expensive mistake.

3.1 A cancer detector

Key point: In medical screening, missing a sick patient is the expensive mistake.

Two models screen 1,000 chest X-rays for cancer (cancer is the positive class):

	TP	FP	FN	TN	Accuracy
Model A	150	90	10	750	0.90
Model B	100	40	60	800	0.90

Again the accuracies are equal. Now:

- **False positive:** a healthy patient is told they may have cancer. Further tests will clear them.
- **False negative:** a patient with cancer is told they are fine, and goes untreated.

The false negative is far more dangerous, so we want the model with fewer of them: model A.

3.2 The formula

Key point: Of all the patients who really have cancer, the fraction the model found.

$$\text{recall} = \frac{TP}{TP + FN}$$

In words: of everything that **really** is positive, how much did the model catch? Recall uses the “actual 1” row of the confusion matrix (Figure 1, right).

- Model A:

$$150/(150 + 10) = 0.94$$

- Model B:

$$100/(100 + 60) = 0.63$$

Model A has the higher recall, matching the choice above. Figure 3 shows where the two numbers come from: recall reads only the outlined “actual cancer” row, and model B’s row holds six times as many missed patients.




Figure 3: The confusion matrices of the two cancer detectors, both with accuracy 0.90. Recall reads the outlined row, everyone who really has cancer: 150 of 160 for model A (0.94), 100 of 160 for model B (0.63).

Extra: Recall is also called **sensitivity** or the **true positive rate** (G-2022) (Fawcett 2006, §2). “Sensitivity” is the name used for medical diagnostic tests (Altman and Bland 1994). Its partner for the negatives is **specificity** (G-2209), $TN/(TN + FP)$: of everyone who really is negative, the fraction the model cleared. The two names return in [the true positive rate](#).

3.3 Choosing between them

Key point: Decide which mistake costs more. False positives worse: precision. False negatives worse: recall.

Problem	Worse mistake	Metric to watch
Spam filter	FP: losing a real email	precision
Cancer screening	FN: missing a patient	recall
Fraud alerts that block cards	FP: blocking honest customers	precision
Airport threat screening	FN: missing a threat	recall

Precision and recall usually pull against each other: making a model flag more cases catches more true positives (higher recall) but also more false alarms (lower precision). [How the](#)

[threshold moves the two rates](#) looks at the same trade-off (scikit-learn docs, “Precision-Recall” example).

Figure 4 shows the pull on the heart-disease test set of section 5 (61 patients). The logistic regression (a model that gives a probability for a yes-or-no target, see [exploring the data](#)) gives every patient a probability of disease, and a patient is flagged when that probability passes the threshold; scikit-learn uses 0.5. Watch the threshold slide: moving it left catches every patient (recall 1.00) but turns healthy people into false alarms (precision 0.56 at 0.05); moving it right clears the false alarms (precision 1.00 at 0.95) but misses most patients (recall 0.38).




Figure 4: The decision threshold slides across the heart-disease test set. Each dot is a patient at its predicted probability of disease, in the lane of its true class; the shaded side is flagged. Green: caught; orange: missed; red: false alarm; grey: correctly cleared. Right: precision, recall and F1 at that threshold

4. F1 score

Key point: F1 is one number built from precision (P) and recall (R), their harmonic mean. F1 is high only if both are high.

4.1 When neither mistake clearly matters more

Key point: Sometimes both errors matter equally, and one number is needed.

For a model that tells cats from dogs, calling a cat a dog is no worse than the reverse. Both precision and recall matter, and comparing two numbers for every model is awkward. The **F1 score** combines them.

4.2 Why the harmonic mean

Key point: The harmonic mean is pulled towards the smaller value, so a model cannot hide a bad precision behind a good recall.

$$F_1 = \frac{2 \times \text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

This formula is the **harmonic mean** (G-880) of the two, not the ordinary (arithmetic) average.

An everyday picture: a chain is only as strong as its weakest link. F1 behaves the same way: one weak score drags the whole number down, however good the other score is.

Precision	Recall	Arithmetic mean	F1
0.8	0.8	0.80	0.80
0.6	1.0	0.80	0.75
0.0	1.0	0.50	0.00

- When both are equal, the two means agree.
- When they differ, F1 stays closer to the smaller one. The second model has the same average as the first but a weak precision, and F1 penalises it.
- A model that flags everything as positive has recall 1 but may have precision near 0. The average would still give it 0.5; F1 gives 0.




Figure 5: F1 against the arithmetic mean

In Figure 5, the horizontal axis is the precision and each colour keeps the recall fixed (blue 1.0, green 0.8). The dashed lines are the arithmetic mean and the solid lines are F1. As

precision falls towards 0, the solid lines drop to 0, while the dashed lines stop at half the recall.

5. In scikit-learn

Key point: `precision_score`, `recall_score` and `f1_score`; by default they report the positive class (1).

Python: Precision, recall and F1.

```
from sklearn.metrics import precision_score, recall_score, f1_score
```

```
precision_score(y_test, y_pred)    # 0.800
recall_score(y_test, y_pred)      # 0.966
f1_score(y_test, y_pred)          # 0.875
```

On the heart-disease test set of [the confusion matrix](#) (61 patients):

Model	Precision	Recall	F1
Logistic regression (standardised)	0.800	0.966	0.875
Decision tree	0.788	0.897	0.839

A decision tree is a chain of yes-or-no questions on the features (see [a decision tree is nested if-else](#)). “Standardised” means each feature was rescaled to mean 0 and standard deviation 1 first.

Logistic regression misses fewer patients (higher recall), which matters most for a disease. Figure 6 draws the table: the two models are close on precision, and the gap in recall carries over into F1.




Figure 6: Precision, recall and F1 on the 61 heart-disease test patients: standardised logistic regression (blue) and a decision tree (orange).

Extra: In a two-class problem, scikit-learn reports the scores of class 1. With `average=None` it shows both classes: for logistic regression, precision is 0.962 for class 0 and 0.800 for class 1. The default comes from the setting `pos_label=1` (scikit-learn docs, `precision_score`).

6. More than two classes

Key point: Compute precision, recall and F1 for each class, then combine them with a macro average (plain mean) or a weighted average (by class size).

6.1 Per class

Key point: Each class in turn is treated as “positive” and all the others as “negative”.

A model sorts 108 animals into dog, cat and rabbit (Figure 7).

actual dog	25	10	5	total 40 recall 25/40 = 0.625
actual cat	1	30	3	total 34 recall 30/34 = 0.882
actual rabbit	3	11	20	total 34 recall 20/34 = 0.588
column total	total 29 precision 25/29 = 0.862	total 51 precision 30/51 = 0.588	total 28 precision 20/28 = 0.714	108 animals
	pred dog	pred cat	pred rabbit	row total

Figure 7: A three-class confusion matrix with per-class precision and recall

- **Precision of a class:** its diagonal count divided by its **column** total (everything predicted as that class). Many dogs and rabbits were called cats. For cat:

$$30/51 = 0.588$$

- **Recall of a class:** its diagonal count divided by its **row** total (everything really in that class). For cat:

$$30/34 = 0.882$$

Class	Precision	Recall	F1	Support
dog	0.862	0.625	0.725	40
cat	0.588	0.882	0.706	34
rabbit	0.714	0.588	0.645	34

The **support** (G-1924) is the number of items really in each class.

6.2 Combining the classes

Key point: Macro: the plain mean of the class scores. Weighted: each class counts in proportion to its support.

- **Macro average** (G-1143): every class counts equally. For precision, add the three class values, then divide by 3. With four decimals (25/29, 30/51 and 20/28), so that the rounding does not shift the last digit:

$$0.8621 + 0.5882 + 0.7143 = 2.1646$$

$$2.1646/3 = 0.7215$$

To three decimals, the macro precision is 0.722.

- **Weighted average** (G-2113): bigger classes count more. Multiply each class value by its support (its number of observations), add, then divide by the total 108:

Class	Precision	Support	Product
dog	0.862	40	34.48
cat	0.588	34	19.99
rabbit	0.714	34	24.28

Then the sum of the products, over the total:

$$34.48 + 19.99 + 24.28 = 78.75$$

$$78.75/108 = 0.729$$

When the classes are about the same size, the two agree closely. When they are very imbalanced, the choice matters: macro gives rare classes a full say, while weighted reflects performance on a typical item.

Python: Everything at once.

```
from sklearn.metrics import classification_report

precision_score(y_test, y_pred, average=None)      # per class
precision_score(y_test, y_pred, average="macro")   # 0.722
precision_score(y_test, y_pred, average="weighted") # 0.729
print(classification_report(y_test, y_pred))

classification_report (G-396) prints precision, recall, F1 and support for every class,
plus the accuracy and both averages.
```

On scikit-learn's handwritten digits (10 classes), logistic regression reaches a macro F1 of 0.971; its weakest class is the digit 5, with precision 0.875.

7. Summary

Metric	Formula	Question it answers	Use when
Precision	$TP/(TP + FP)$	of the predicted positives, how many are right?	false positives are costly
Recall	$TP/(TP + FN)$	of the real positives, how many were found?	false negatives are costly
F1	$2PR/(P + R)$	one number that needs both to be high	both mistakes matter

- Equal accuracy can hide very different mistakes; precision and recall separate them.
- F1 is the harmonic mean, so it stays near the weaker of the two.
- With several classes, compute per class and combine with a macro or weighted average.

8. Sources

Built from

- CampusX, “Precision, Recall and F1 Score | Classification Metrics Part 2”, YouTube, <https://www.youtube.com/watch?v=iK-kdhJ-7yI>
- StatQuest with Josh Starmer, “Machine Learning Fundamentals: Sensitivity and Specificity”, YouTube, <https://www.youtube.com/watch?v=vP06aMoz4v8>. Sensitivity and specificity as the medical names (Section 3.2).
- StatQuest with Josh Starmer, “ROC and AUC, Clearly Explained!”, YouTube, <https://www.youtube.com/watch?v=4jRBRDbJemM>. Precision does not use the true negatives (Section 2.2, Extra).

Other references

- **Altman and Bland 1994:** Altman, D. G. and Bland, J. M. “Diagnostic tests 1: sensitivity and specificity.” *BMJ* 308(6943), 1552, 1994.
- **Fawcett 2006:** Fawcett, T. “An Introduction to ROC Analysis.” *Pattern Recognition Letters* 27(8), 861–874, 2006. Section 2 and Figure 1.
- **scikit-learn docs:** `sklearn.metrics.precision_score` (pos_label); example “Precision-Recall”, scikit-learn 1.9.

9. Key terms

Term	Meaning
Precision	Of all items predicted positive, the fraction that really are positive

Term	Meaning
Recall (sensitivity)	Of all items that really are positive, the fraction the model found
Specificity (G-2209)	Of all items that really are negative, the fraction the model cleared: $TN / (TN + FP)$
F1 score	The harmonic mean of precision and recall
Harmonic mean	An average that stays close to the smaller of the values: $2ab/(a + b)$ for two values
Support	The number of items that really belong to a class
Macro average	The plain mean of a metric over all classes
Weighted average	The mean of a metric over classes, weighted by each class's support
classification_report	scikit-learn function that prints precision, recall, F1 and support for every class