

Softmax Regression (Multinomial Logistic Regression)

Contents

1. Overview	2
2. The softmax function	2
2.1 Raw scores, and why picking the largest is not enough	2
2.2 The formula	3
2.3 Its properties	4
2.4 Softmax has a slope	4
2.5 Two classes give the sigmoid	5
3. How the model predicts	5
4. How the model is trained	6
4.1 An intuition: one model per class	6
4.2 The real approach: one loss for all classes	6
5. In scikit-learn	7
6. Summary	9
7. Sources	9
8. Key terms	9

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** One-hot encoding ([Note 11](#)); Logistic regression ([Note 13](#)).
- **Leads to:** Categorical and sparse categorical cross-entropy ([Note 633](#)); ANN for classification ([Note 1011](#)); Unembedding, logits, temperature and sampling ([Note 1088](#)).

1. Overview

Key point: Softmax regression extends logistic regression to more than two classes. The model computes one score per class and turns the scores into probabilities that add up to 1, using the softmax function.

The logistic regression of the previous Notes handles **binary** classification: placed or not, spam or not. Many problems have more classes. A student might be placed, not placed, or opt out of placements altogether; an iris flower can be one of three species.

Softmax regression (G-1833), also called **multinomial logistic regression** (G-1276), handles any number of classes. Softmax is also the standard output layer of neural networks that classify, so it matters for deep learning too (Goodfellow et al. §6.2.2.3). With two classes it reduces exactly to ordinary logistic regression.

2. The softmax function

Key point: $\hat{y}_k = e^{(z_k)} / \sum_j e^{(z_j)}$. Each output is between 0 and 1, and the outputs add up to 1.

2.1 Raw scores, and why picking the largest is not enough

Key point: Picking the largest score gives the class, but it gives gradient descent nothing to learn from.

A model gives a flower one raw score per class. Take a flower with scores 2.81 for setosa, 1.84 for versicolor and -4.66 for virginica. Raw scores are hard to read: some are above 1, some are below 0, and nothing says how sure the model is.

The simplest fix is to pick the largest score. The function that does so is **argmax** (G-212): it gives 1 to the class with the largest score and 0 to all the others. Here argmax gives setosa 1, versicolor 0 and virginica 0. That output is easy to read, and it is how the final class is reported.

Argmax cannot be used to train the model, though. Training uses gradient descent, and gradient descent needs a slope: how much the output changes when a score changes a little.

1. Raise the setosa score from 2.81 to 2.91. Setosa is still the largest, so the argmax output stays 1, 0, 0.
2. Lower it to 2.71. The output is still 1, 0, 0.
3. The output did not move, so the slope is 0. A slope of 0 tells gradient descent to change nothing, and the weights never improve.

We need a function that still favours the largest score but changes smoothly. That function is softmax. The model trains with softmax and uses argmax only to report the final class (Figure 2 shows the two side by side).

2.2 The formula

Key point: Raise e to each score, then divide by the total.

Call the three scores z_1 , z_2 and z_3 . The **softmax function** (G-1830) turns them into probabilities in two steps: raise e to each score, then divide by the total.

$$\hat{y}_k = \frac{e^{z_k}}{e^{z_1} + e^{z_2} + e^{z_3}}$$

Here $\hat{y}_k$ is the softmax output for class k . In general, with K classes, the denominator sums e^{z_j} over all K classes.




Figure 1: Softmax step by step for one flower

With numbers (Figure 1), a flower with scores (2.81, 1.84, -4.66):

1. e^z : (16.69, 6.30, 0.01). Every value is now positive, and a larger score gives a much larger value.
2. Divide by the total, 23.00: (0.726, 0.274, 0.0004).

The flower is most likely setosa (73%), possibly versicolor (27%), and almost certainly not virginica.

2.3 Its properties

Key point: All outputs are in $(0, 1)$, they add up to 1, and the order of the scores is kept.

- Each output is between 0 and 1, because e^z is always positive and the denominator includes it.
- The outputs always add up to 1: the classes share out all the probability. A flower must be one of the three species.
- The order of the scores is kept: the class with the largest score always gets the largest probability. In Figure 1 the scores rank setosa, versicolor, virginica ($2.81 > 1.84 > -4.66$), and so do the probabilities ($0.726 > 0.274 > 0.0004$).

Because the outputs lie between 0 and 1 and add up to 1, we read them as probabilities. They are the model's own estimates: they come from the fitted weights, so a different fit gives somewhat different values.

2.4 Softmax has a slope

Key point: A small change in a score gives a small change in every probability, so gradient descent has a slope to follow.




Figure 2: The setosa score slides from -3 to 6 while the other two scores stay fixed. Left: the three softmax probabilities, which always add up to 1. Right: the setosa output against its score: argmax (dashed) is a step with slope 0; softmax (solid) is a smooth curve.

In Figure 2, watch the left panel first. As the setosa score rises, the setosa probability rises and the versicolor probability falls by the same amount. The classes compete for a fixed total of 1.

Now the right panel. The dashed argmax line is flat at 0, jumps to 1 where the setosa score passes the versicolor score (1.84), and is flat again. Its slope is 0 everywhere except at the jump. The solid softmax curve rises smoothly, so it has a slope at every point.

The slope has a simple formula. Write p_k for the softmax probability of class k :

- raising a class's own score raises its probability at the rate $p_k(1 - p_k)$;
- raising another class's score z_j lowers p_k at the rate $-p_k p_j$.

For our flower, $p_{\text{setosa}} = 0.726$, so its own slope is $0.726 \times (1 - 0.726) = 0.20$, and the slope with respect to the versicolor score is $-0.726 \times 0.274 = -0.20$. Neither is 0, so gradient descent can use them (Section 4.2).

2.5 Two classes give the sigmoid

Key point: With two classes, softmax is the sigmoid of the difference of the two scores.

With two classes, divide the top and bottom by e^{z_1} :

$$\frac{e^{z_1}}{e^{z_1} + e^{z_2}} = \frac{1}{1 + e^{-(z_1 - z_2)}} = \sigma(z_1 - z_2)$$

The result is the sigmoid of the difference of the two scores. So binary logistic regression is the special case of softmax regression with two classes.

3. How the model predicts

Key point: One weight vector per class gives one score per class; softmax turns the scores into probabilities; the largest probability wins.




Figure 3: From one flower to a predicted class

With K classes, the model has K weight vectors, one per class (Figure 3). To predict a flower x (with a 1 in front for the intercept):

1. compute a score for every class: $z_k = w^{(k)} \cdot x$;

2. apply softmax to get K probabilities;
3. predict the class with the largest probability.

On the iris data, each flower is one **observation** (G-1374; one record, a row of the data table). We use two **features** (G-772; input variables, one column each): sepal length and petal length. The **target** (G-1949; the output we predict) is the species. The trained model has three weight vectors:

Class	Intercept	Sepal length	Petal length
setosa	11.423	-0.210	-2.924
versicolor	1.605	0.347	-0.350
virginica	-13.028	-0.137	3.274

For a flower with sepal length 3.4 and petal length 2.7, the setosa score is $11.423 - 0.210 \times 3.4 - 2.924 \times 2.7 = 2.81$. The other two come out at 1.84 and -4.66. Softmax then gives the probabilities of Figure 1, and the prediction is setosa.

4. How the model is trained

Key point: One-hot encode the classes, then minimise the categorical cross entropy, the multi-class version of log loss, with gradient descent.

4.1 An intuition: one model per class

Key point: Turn the K-class problem into K yes-or-no problems, one per class.

A simple way to picture training is to **one-hot encode** (G-1379) the output (the one-hot encoding Note). A column with values 0, 1, 2 becomes three columns: “is it class 0?”, “is it class 1?”, “is it class 2?”. Each column is a binary problem, so one logistic regression could be trained per column, giving three weight vectors.

The one-model-per-class picture works, and scikit-learn offers it as **one-vs-rest** (G-1388). But the K models are trained separately, so nothing makes their probabilities add up to 1, and training K separate models gets slow when there are many classes.

4.2 The real approach: one loss for all classes

Key point: $L = -(1/m) \sum_i \sum_k y_{ik} \log \hat{y}_{ik}$. For each observation, only the log probability of its true class counts.

Softmax regression instead trains all K weight vectors together, by minimising one loss function, the **categorical cross entropy** (G-349):

$$L = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K y_{ik} \log \hat{y}_{ik}$$

Here y_{ik} is the one-hot value (1 if observation i is class k , else 0) and $\hat{y}_{ik}$ the softmax probability of class k . For each observation, every term is multiplied by 0 except the true class. So the loss is simply the average of $-\log(\text{probability given to the true class})$: exactly the log loss of the earlier Note, extended to K classes. With $K = 2$ it is the **binary cross entropy** (G-303).

With 2 features and 3 classes there are $3 \times 3 = 9$ weights. Gradient descent computes the derivative of L with respect to all nine and updates them together, just as in the gradient descent Note for logistic regression.

Figure 4 runs this on the training flowers, with the two features standardised. All nine weights start at 0, so every class gets probability $1/3$ and the loss is $\log 3 = 1.10$. As the weights move together, the three regions sort themselves out and the loss falls to 0.31 after 500 epochs, with 28 of the 30 test flowers right (0.933). scikit-learn's solver, run to convergence, reaches the 0.967 of Section 5.




Figure 4: Softmax regression trained by gradient descent on the iris training flowers (learning rate 0.1, standardised features). Left: the decision regions after 0 to 500 epochs. Right: the categorical cross entropy per epoch, falling from $\log 3 = 1.10$ to 0.31.

5. In scikit-learn

Key point: LogisticRegression uses softmax automatically when there are more than two classes. predict_proba returns the softmax probabilities.

Python: Softmax regression on iris.

```
from sklearn.linear_model import LogisticRegression
```

```
X = iris.data[:, [0, 2]] # sepal length, petal length
```

```

clf = LogisticRegression()           # multinomial for 3+ classes
clf.fit(X_train, y_train)
clf.score(X_test, y_test)           # 0.967
clf.predict_proba([[3.4, 2.7]])     # [0.726, 0.274, 0.0004]
clf.predict([[3.4, 2.7]])           # [0] (setosa)

```

Extra: Older versions of scikit-learn needed `LogisticRegression(multi_class="multinomial")`. Recent versions removed that setting: with the default solver, multi-class problems always use softmax. One-vs-rest is still available as `OneVsRestClassifier(LogisticRegression())` (scikit-learn docs, `LogisticRegression`).

On 30 test flowers, the model gets 29 right (accuracy 0.967); the one mistake is a versicolor predicted as virginica.




Figure 5: Decision regions of softmax regression on iris

Figure 5 shows the **decision regions** (G-557): each point of the plane is coloured by the class the model would predict. The boundaries between regions are straight lines, because each score is a linear function of the features, so softmax regression is still a linear classifier. The query flower (star) sits in the setosa region, close to the versicolor boundary, which matches its 73%/27% split.

6. Summary

- Softmax regression = logistic regression for $K \geq 2$ classes; also called multinomial logistic regression.
- Softmax: $\hat{y}_k = e^{z_k} / \sum_j e^{z_j}$; outputs are probabilities that add up to 1.
- Argmax picks the largest score but has slope 0, so the model trains with softmax and reports the class with argmax.
- One weight vector per class; predict the class with the largest probability.
- Training minimises the categorical cross entropy with gradient descent; with two classes, everything reduces to the sigmoid and the binary log loss.
- Decision boundaries are straight lines.

7. Sources

Built from

- CampusX, “Softmax Regression || Multinomial Logistic Regression || Logistic Regression Part 6”, YouTube, https://www.youtube.com/watch?v=Z8noL_0M4tw
- StatQuest with Josh Starmer, “Neural Networks Part 5: ArgMax and SoftMax”, YouTube, <https://www.youtube.com/watch?v=KpKog-L9veg> (argmax against softmax, the three properties, the slope of softmax)

Other references

- **Goodfellow et al.:** Goodfellow, I., Bengio, Y. and Courville, A. *Deep Learning*. MIT Press, 2016. Section 6.2.2.3, “Softmax Units for Multinoulli Output Distributions”.
- **scikit-learn docs:** `sklearn.linear_model.LogisticRegression`, scikit-learn 1.9.

8. Key terms

Term	Meaning
Softmax regression	Logistic regression extended to any number of classes using the softmax function
Multinomial logistic regression	Another name for softmax regression
Softmax function	Turns a list of scores into probabilities: $e^{z_k} / \sum_j e^{z_j}$
Argmax	Gives 1 to the class with the largest score and 0 to the others; used to report the class, not to train
Categorical cross entropy	The loss of softmax regression: the average of $-\log(\text{probability of the true class})$
One-vs-rest	Training one binary classifier per class, each separating that class from all others
Decision region	The part of the feature space in which a model predicts a given class