

Bagging vs Random Forest

Contents

1. Overview	2
2. Difference 1: the base model	3
3. Difference 2: tree-level against node-level feature sampling	4
3.1 Bagging: tree-level sampling	4
3.2 Random forest: node-level sampling	4
3.3 Why more randomness helps	4
4. Checking it in code	6
4.1 The bagged trees	6
4.2 The forest's trees	7
4.3 Does it pay off?	7
5. Summary	9
6. Sources	9
7. Key terms	9

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Overfitting ([Note ML-007](#)); Hyperparameter tuning ([Note ML-009](#)); Grid and random search ([Note ML-028](#)); Decision trees ([Note ML-091](#)); Feature importance ([Note ML-093](#)); Bagging ([Note ML-099](#)).
- **Leads to:** Balanced random forest ([Note ML-127](#)).
- **Compare with:** Bagging ([Note ML-101](#)); Dropout ([Note DL-024](#)).

1. Overview

Key point: Two differences: bagging can use any **base model** (G-260) while a random forest always uses decision trees; and bagging samples features once per tree while a random forest samples them again at every node.




Figure 1: Feature sampling with 2 of 5 features: bagging draws the pair once and the whole tree uses it; a random forest draws a new pair at every node

A **random forest** (G-1611) is built on **bagging** (G-251; the [random forest introduction Note](#)), but the two are not the same, even when bagging uses decision trees. Figure 1 shows the subtle difference: **where** the **features** (input variables, one column of the data table each) are sampled. Each **observation** (one record, one row of the table) has a value for every feature and a **target**, the class we predict.

This Note covers both differences and checks the second one in code. The Notebook ([notebook.ipynb](#)) runs every check.

2. Difference 1: the base model

Key point: Bagging is a general technique for any algorithm; a random forest is made of decision trees only.

In bagging, the base models can come from any algorithm, as long as they all use the same one: all decision trees, all KNN or all SVMs (the [bagging classifier Note](#), section 2.2). Decision trees, KNN and SVMs are the usual choices.




Figure 2: Difference 1: `BaggingClassifier` takes any one algorithm as its base model; `RandomForestClassifier` is always made of decision trees

Figure 2 sets the two classes side by side. Watch the left column: bagging offers a choice of base model, while the forest's column holds only trees.

In a random forest, the base model is always a decision tree. scikit-learn's classes show this directly:

- `BaggingClassifier` has an `estimator` parameter for choosing the base model; when it is left at `None`, it uses a decision tree.
- `RandomForestClassifier` has `n_estimators`, the number of trees, but **no** parameter for choosing the base model.

Extra: Older documentation and code call this parameter `base_estimator`. It was renamed `estimator` and the old name was removed in scikit-learn 1.4 (the [bagging classifier Note](#), section 3.2).

3. Difference 2: tree-level against node-level feature sampling

Key point: Bagging fixes each tree's features before the tree is grown; a random forest picks a fresh random set of features before every split.

So is a bagging ensemble of decision trees a random forest? **No.** The remaining difference is in how the features are sampled (feature sampling: the [random forest introduction Note](#), section 4).

Take a dataset with 5 features, and suppose each tree may use 2 of them (`max_features=2` (G-1185) in both classes).

3.1 Bagging: tree-level sampling

Key point: Two features are drawn once; every split in that tree uses one of those two.

Before tree 1 is grown, 2 of the 5 features are drawn at random, say col1 and col3. Tree 1 is then trained on those two features only, and every split in it is on col1 or col3 (Figure 1, left). Tree 2 gets its own pair, say col4 and col5; tree 3 col1 and col4; and so on.

Bagging therefore uses **tree-level feature sampling** (tree-level column sampling, G-2014): the features are decided once, before the tree is built, and no other feature is ever touched by that tree.

3.2 Random forest: node-level sampling

Key point: Two features are drawn at every node, so one tree can end up using all five.

In a random forest, the draw happens again **every time a node is about to split** (Figure 1, right). The forest therefore uses **node-level feature sampling** (node-level column sampling, G-1325):

- the root draws col2 and col3, and splits on the better of the two, col3;
- the next node draws again, col4 and col5, and splits on col4;
- the next draws col2 and col1, and splits on col1.

Building a forest tree this way is shown step by step in the [random forest introduction Note](#), section 4.1. Node-level sampling is exactly the `max_features` setting of a single decision tree (the [decision tree hyperparameters Note](#), section 4.6), applied in every tree of the forest. Node-level sampling adds more randomness than tree-level sampling.

3.3 Why more randomness helps

Key point: More randomness makes the trees more different from each other, and an ensemble of different models performs better.

From the ensemble Notes (the [voting ensemble Note](#)): an ensemble works best when its base models are each better than chance and as different from each other as possible. How alike

two models' predictions are is measured by the **correlation between base models** (G-488): 0 means unrelated, 1 means identical. The chain from sampling to accuracy:

1. Node-level sampling forces each split to choose among a random few features, so even trees grown on similar observations end up splitting on different features.
2. Trees that split differently make their mistakes on different observations, so their correlation is lower (ESL §15.2).
3. When the trees vote or are averaged, mistakes made by only some trees are outvoted, so a lower correlation leaves a smaller error.

The Extra below gives the formula, and Figure 3 plots it.

ESL Figure 15.1 shows this gain on spam email data, and section 4.3 repeats the comparison on the same data.

Extra: Why does difference between the trees matter so much? Suppose each tree's prediction has **variance** (G-2074) σ^2 , its spread around its average and any two trees are correlated by ρ (0: unrelated, 1: identical).

1. **In words:** the variance of the average of n trees has a part that more trees can remove and a part, set by the correlation, that they cannot.

2. **Formula:**

$$\text{variance of the average} = \rho \sigma^2 + \frac{1 - \rho}{n} \sigma^2$$

3. **Example:** with $\sigma^2 = 1$ and $n = 100$ trees. For $\rho = 0.6$:

$$0.6 + 0.4/100 = 0.604$$

For $\rho = 0.3$:

$$0.3 + 0.7/100 = 0.307$$




Figure 3: Variance of the average of n trees with $\sigma^2 = 1$, for trees correlated by $\rho = 0.6$ (blue) and $\rho = 0.3$ (orange); the dashed lines mark the floor $\rho \sigma^2$

Figure 3 plots the formula. Watch both curves flatten: by 20 trees each is within 0.04 of its dashed floor, so more trees stop helping, and only a lower ρ moves the floor down. Adding trees only shrinks the second term. The first term stays, however many trees we add, unless the trees become less alike. Lowering ρ is exactly what node-level sampling does. (ESL §15.2, eq. 15.1)

4. Checking it in code

Key point: Every bagged tree splits on exactly 2 features; the forest's trees split on 2 to 5 features, most often 4.

We use a dataset of 100 observations and 5 features, `col1` to `col5`, made with `make_classification`, and give both classes `max_features=2`.

4.1 The bagged trees

Key point: The first bagged tree only ever splits on `col5` and `col1`, the two features it was given.

Python: Bagging with 2 features per tree, and the features each tree got.

```
from sklearn.ensemble import BaggingClassifier
```

```
bag = BaggingClassifier(max_features=2, random_state=42)
```

```
bag.fit(df.iloc[:, :5], df["target"])
```

```
bag.estimators_features_[0]          # array([4, 0])
```

`bag.estimators_[0]` is the first trained tree. `estimators_features_` (G-710) [0] lists the features it was given (the [bagging classifier Note](#), section 3.3): positions 4 and 0, that is `col5` and `col1`.

When we print the first tree, every split is on `feature_0` or `feature_1`. Careful: these are **not** `col1` and `col2`. The tree was trained on a 2-column table, so it numbers those two features 0 and 1 itself; `estimators_features_` maps them back to `col5` and `col1`.

Python: Printing a tree with the real feature names.

```
from sklearn.tree import export_text
```

```
cols = bag.estimators_features_[0]
```

```
print(export_text(bag.estimators_[0],  
                  feature_names=list(df.columns[cols])))
```

`export_text` (G-737) prints a tree as indented text, one line per split. `feature_names` replaces `feature_0`, `feature_1`, ... with names.

All 10 bagged trees use exactly 2 features: (`col1`, `col5`), (`col1`, `col3`), (`col3`, `col5`), and so on.

4.2 The forest's trees

Key point: The first forest tree splits on col3, col5, col1 and col2: four different features with `max_features=2`.

The first tree of `RandomForestClassifier(max_features=2)` splits first on col3, then on col5, col1 and col2. Every tree of a forest is trained on all the features, so its feature numbers are the real ones.




Figure 4: How many different features each of 100 trees splits on, out of 5, with `max_features=2`

Figure 4 counts, for 100 trees of each kind, how many different features each tree splits on:

- **bagging:** all 100 trees use exactly 2 features;
- **random forest:** 5 trees use 2 features, 28 use 3, 50 use 4, and 17 use all 5.

The counts prove the point: bagging samples features at the tree level, a random forest at the node level.

4.3 Does it pay off?

Key point: On real spam data, the random forest scores 0.953, bagging 0.946 with all features and 0.922 with 7 features per tree.

We use the Spambase data (UCI, via OpenML): 4,601 emails, 57 features such as how often words like “free” or “\$” appear, and a target of spam or not. ESL §15.2 compares bagging and random forests on the same data. Each model has 200 trees, scored with 5-fold **cross-validation** (G-510) repeated 3 times:

Model	Features	Accuracy
Bagging	all 57 per tree	0.946
Bagging	7 per tree (tree level)	0.922
Random forest	7 per node ($\sqrt{57} \approx 7$, the default)	0.953

The random forest makes about 13% fewer mistakes than bagging (4.7% against 5.4% of emails wrong), the same ranking ESL Figure 15.1 shows on this data. Tree-level sampling is the worst of the three: each tree is stuck with 7 random features for its whole life and often misses the most useful words.

Extra: A bagging ensemble of trees that each sample features at every node, `BaggingClassifier(DecisionTreeClassifier(max_features="sqrt"))`, scores 0.953 too. So with node-level sampling, bagged trees **are** a random forest, which is how scikit-learn defines one (scikit-learn User Guide §1.11).

5. Summary

	Bagging	Random forest
Base model	any algorithm (estimator)	decision trees only
Feature sampling	once per tree (tree level)	at every node (node level)
Features one tree can use, with <code>max_features=2</code> of 5	exactly 2	up to all 5
Randomness	less	more
On the spam data (accuracy)	0.946 (all features)	0.953

- A bagging ensemble of decision trees is still not a random forest.
- Bagging decides each tree's features before it is grown; a random forest re-draws them before every split.
- More randomness makes the trees less alike, which makes the forest better (spam data: 0.953 against 0.946).
- In a bagged tree, `feature_0`, `feature_1`, ... are positions within that tree's own features; `estimators_features_` maps them back.

6. Sources

Built from

- CampusX, “Bagging Vs Random Forest | What is the difference between Bagging and Random Forest | Very Important”, YouTube, <https://www.youtube.com/watch?v=l93jRojZMqU>
- StatQuest with Josh Starmer, “StatQuest: Random Forests Part 1 - Building, Using and Evaluating”, YouTube, https://www.youtube.com/watch?v=J4Wdy0Wc_xQ (a random subset of features at every split, section 3.2)

Other references

- Hastie, T., Tibshirani, R. and Friedman, J. (2009). *The Elements of Statistical Learning*, 2nd ed. Springer. §15.2 (eq. 15.1, variance of an average of correlated trees; Figure 15.1, bagging against random forest on the spam data).
- Hopkins, M., Reeber, E., Forman, G. and Suermondt, J. (1999). Spambase [dataset]. UCI Machine Learning Repository. archive.ics.uci.edu/dataset/94/spambase
- scikit-learn developers. User Guide, §1.11 “Ensembles”, Random forests. scikit-learn.org/stable/modules/ensemble.html

7. Key terms

Term	Meaning
Tree-level feature sampling (G-2014)	Drawing one random set of features per tree, before the tree is grown; every split of that tree uses only those features (bagging)

Term	Meaning
Node-level feature sampling (G-1325)	Drawing a new random set of features before every split (random forest)
<code>export_text</code>	scikit-learn function that prints a trained tree as indented text
Correlation between base models	How alike two base models' predictions are; the less alike, the more an ensemble cuts variance