


OOB Score: Out-of-Bag Evaluation

Contents

1. Overview	2
2. Out-of-bag observations	3
3. How the OOB score is computed	4
3.1 The steps	4
3.2 The formula	4
4. The OOB score in scikit-learn	5
5. Rebuilding the OOB score by hand	6
6. How many trees the OOB score needs	8
7. When the OOB score is available	9
8. Summary	10
9. Sources	10
10. Key terms	10

Where this fits: Pipeline map steps 8 (Model), 9 (Evaluate). See the [Course map](#).



- **Builds on:** Overfitting ([Note ML-007](#)); Hyperparameter tuning ([Note ML-009](#)); Grid and random search ([Note ML-028](#)); Decision trees ([Note ML-091](#)); Feature importance ([Note ML-093](#)); Bagging ([Note ML-099](#)).
- **Leads to:** Balanced random forest ([Note ML-127](#)).
- **Compare with:** Bagging ([Note ML-101](#)); Cross-validation ([Note ML-106](#)); Dropout ([Note DL-024](#)).

1. Overview

Key point: Each tree in a forest never sees about 37% of the training observations. Predicting every observation with only the trees that never saw it, and scoring those predictions, gives a free estimate of how the forest will do on new data.

Out-of-bag evaluation: each row is judged only by trees that never saw it

	tree 1	tree 2	tree 3	tree 4
row 1 y = 1	x2	OOB		
row 2 y = 0	seen	seen		
row 3 y = 1	OOB	x2		
row 4 y = 0	seen	seen		
row 5 y = 1	x2	OOB		
row 6 y = 0	OOB	x2		

Out-of-bag evaluation: each row is judged only by trees that never saw it

	tree 1	tree 2	tree 3	tree 4
row 1 y = 1	x2	OOB	seen	seen
row 2 y = 0	seen	seen	OOB	x2
row 3 y = 1	OOB	x2	seen	seen
row 4 y = 0	seen	seen	x2	OOB
row 5 y = 1	x2	OOB	seen	OOB
row 6 y = 0	OOB	x2	seen	x2

blue: in the tree's bootstrap sample orange: out-of-bag (OOB) for that tree

Out-of-bag evaluation: each row is judged only by trees that never saw it

	tree 1	tree 2	tree 3	tree 4	OOB vote
row 1 y = 1	x2	says 1	seen	seen	1 right
row 2 y = 0	seen	seen	says 0	x2	0 right
row 3 y = 1	says 0	x2	seen	seen	0 wrong
row 4 y = 0	seen	seen	x2	OOB	
row 5 y = 1	x2	OOB	seen	OOB	
row 6 y = 0	OOB	x2	seen	x2	

blue: in the tree's bootstrap sample orange: out-of-bag (OOB) for that tree

Out-of-bag evaluation: each row is judged only by trees that never saw it

	tree 1	tree 2	tree 3	tree 4	OOB vote
row 1 y = 1	x2	says 1	seen	seen	1 right
row 2 y = 0	seen	seen	says 0	x2	0 right
row 3 y = 1	says 0	x2	seen	seen	0 wrong
row 4 y = 0	seen	seen	x2	says 0	0 right
row 5 y = 1	x2	says 1	seen	says 1	1 right
row 6 y = 0	says 0	x2	seen	x2	0 right

OOB score = 5 right / 6 rows = 0.83

Figure 1: Out-of-bag evaluation with 6 observations and 4 trees: each observation is predicted only by the trees whose bootstrap sample missed it

Out-of-bag (OOB) evaluation (G-1411) is a way to test a bagging model (many models, each trained on a random resample of the data, whose answers are combined; see [the core idea of bagging](#)), such as a [random forest](#) (bagging of decision trees), using only its training data. Each **observation** is one record (one row of the data table): its **features** (input variables, one column each) and its **target** (the output we predict). OOB evaluation is available in every bagging-based ensemble of scikit-learn (section 7). Figure 1 shows the whole process

on 6 observations and 4 trees.

The OOB score was introduced with [the out-of-bag score of the bagging classifier](#): `oob_score=True`, then `oob_score_`. This Note:

- explains exactly how the score is computed (section 3);
- rebuilds it by hand (section 5);
- shows when it can be trusted (sections 4, 6 and 7).

The Notebook (`ML-107-oob-score.ipynb`) runs every step.

2. Out-of-bag observations

Key point: Drawing with replacement means some observations are drawn several times for a tree and others not at all; the observations a tree missed are out-of-bag for that tree.

Each tree is trained on a **bootstrap sample** (G-319), drawn **with replacement** (`bootstrap=True`), which holds on average about 63.2% of the distinct observations ([drawing with replacement](#)). The other **36.8%**, the observations a tree never drew, are its **out-of-bag (OOB) observations** (G-1412): the tree has never seen them, so they can serve as test data for that tree.

Figure 2 draws one tree's bootstrap sample from 20 observations, one draw at a time. Each draw picks any of the 20 observations, including ones already drawn. Watch the bars: some observations are drawn two, three or four times, and 7 of the 20 (35%) are never drawn. Those 7 are the tree's out-of-bag observations (orange).




Figure 2: One tree's bootstrap sample from 20 observations: 20 draws with replacement. Bars count how often each observation was drawn; the 7 never drawn are out-of-bag (OOB)

Careful: being out-of-bag is a property of an observation **for one tree**, not for the whole forest. An observation missed by tree 1 is almost certainly seen by other trees. The chance that one observation is missed by all 100 trees is

$$0.368^{100} \approx 10^{-43},$$

so in practice no observation is hidden from the whole forest. What OOB evaluation uses is that every observation is hidden from **some** trees.

3. How the OOB score is computed

Key point: For each observation, ask only the trees that never saw it, combine their answers, and compare with the truth; the share of correct answers is the OOB score.

3.1 The steps

Key point: Four steps: record each tree’s sample, find each observation’s OOB trees, let them vote, score the votes.

Figure 1 runs these steps on 6 observations and 4 trees:

1. **Record each tree’s bootstrap sample.** Blue cells are observations a tree was trained on (“x2”: drawn twice); orange cells are its OOB observations. Tree 1 missed observations 3 and 6; tree 4 missed observations 4 and 5.
2. **For each observation, find the trees that never saw it.** Observation 5 is OOB for trees 2 and 4; observation 1 only for tree 2.
3. **Let only those trees vote.** Trees 2 and 4 both say 1 for observation 5, so observation 5’s **OOB prediction** (G-1392) is 1.
4. **Compare with the truth.** Observation 3 is really 1, but its only OOB tree, tree 1, says 0: wrong. The other five observations are right.

3.2 The formula

Key point: OOB accuracy is the number of observations whose OOB prediction is right, divided by the number of observations that have an OOB prediction.

1. **In words:** the share of training observations predicted correctly by the trees that never saw them.
2. **Formula:**

$$\text{OOB score} = \frac{\text{correct}}{\text{predicted}}$$

Here “correct” is the number of rows whose OOB prediction is correct, and “predicted” is the number of rows that have an OOB prediction.

3. **Example:** in Figure 1, 5 of the 6 observations are right:

$$\text{OOB score} = \frac{5}{6} = 0.83$$

For a regressor, the OOB prediction of an observation is the **mean** of its OOB trees' numbers, and `oob_score_` is the R^2 of those predictions (**R^2 score**; see also [the out-of-bag score of the bagging regressor](#)).

Extra: scikit-learn does not count hard votes. Each OOB tree gives its class probabilities for the observation (the share of each class in the leaf the observation lands in), the probabilities are added up, and the class with the largest total wins. Adding probabilities in this way is **soft voting** (G-1828) ([soft voting](#)). The totals, divided by the number of OOB trees, are stored in `oob_decision_function_`. On our data, hard votes and soft votes give the same score.

4. The OOB score in scikit-learn

Key point: On the heart disease data, the OOB score is as good an estimate as 5-fold cross-validation (0.820 against 0.820, averaged over 50 splits), at no extra cost.

We use [the heart disease data](#): 303 patients, 242 for training and 61 for testing.

Python: A random forest with its OOB score.

```
rf = RandomForestClassifier(oob_score=True,  
                           random_state=42)
```

```
rf.fit(X_train, y_train)
```

```
rf.oob_score_           # 0.835
```

```
rf.score(X_test, y_test) # 0.836
```

`oob_score=True` must be set **before fit**: the OOB predictions are collected while the forest is trained. Without it, `oob_score_` does not exist.

On this split the OOB score, **0.835**, is close to the test accuracy, **0.836**, and a 5-fold **cross-validation** (G-510) on the training set gives 0.806. One split of 61 test patients is noisy, though: one patient moves the test accuracy by 0.016. So the Notebook repeats the comparison on 50 random splits, with 500 trees:

Estimate (mean over 50 splits)	Accuracy
OOB score	0.820
5-fold cross-validation	0.820
Test accuracy (61 patients)	0.831

On a typical split, the OOB score and the cross-validation score differ by only 0.011, while each one misses the 61-patient test accuracy by 0.038: the small test set is the noisy one. The OOB estimate is almost identical to cross-validation (ESL §15.3.1), but cross-validation needs 5 more forests to be trained. The OOB score needs none. Think of a study group in which each member is quizzed only on questions they never practised: every quiz is fair, and no separate exam is needed.

Figure 3 plots all 50 splits. Each point is one split: its OOB score across, and another estimate up. Watch the blue cross-validation points stay close to the dashed line, where the

two estimates are equal, while the grey test-set points spread far above and below it.




Figure 3: 50 random splits of the heart disease data, 500 trees each. Across: the OOB score of the split. Up: the 5-fold cross-validation score (blue) and the accuracy on the 61 test patients (grey). Dashed: equal to the OOB score

So the OOB observations act as a **validation set** (G-2067; data not used for training, kept to check the model) that comes for free: about 37% of the data, unseen by each tree, without holding out any observations from training.

A free validation score is also a way to tune a forest: train forests with different settings, for example different values of `max_features` ([how many features max_features means](#)), and keep the one with the best OOB score.

5. Rebuilding the OOB score by hand

Key point: `estimators_samples_` tells us which observations each tree saw; predicting each tree's missing observations and adding up the probabilities reproduces `oob_score_` exactly.

A fitted forest keeps each tree's bootstrap sample in `estimators_samples_` (the same attribute as in [which observations and features each tree got](#)). From it, the Notebook builds a table with one row per tree and one column per training observation, marking which observations each tree missed. Figure 4 draws that table.




Figure 4: The out-of-bag table of the forest above: 100 trees (rows) by 242 training observations (columns), orange where a tree missed an observation. Bottom: the number of trees that missed each observation

Reading Figure 4 row by row and column by column gives two facts:

- each tree's share of OOB observations is between 32% and 41%, **36.7%** on average, matching the theory's 36.8%;
- every training observation is OOB for at least 24 of the 100 trees (36.7 on average, at most 47), so every observation gets an OOB prediction.

Python: The OOB predictions, step by step.

```
in_bag = np.zeros((n_trees, n_rows), dtype=bool)
for t, rows in enumerate(rf.estimators_samples_):
    in_bag[t, rows] = True
oob = ~in_bag                                # True where tree t missed row i

proba = np.zeros((n_rows, 2))
for t, tree in enumerate(rf.estimators_):
    proba[oob[t]] += tree.predict_proba(X_arr[oob[t]])
(proba.argmax(axis=1) == y_arr).mean()      # 0.835
~ flips True and False. proba[oob[t]] += ... adds tree t's probabilities only to the
observations it missed. argmax(axis=1) picks, for each observation, the class with the
larger total.
```

The result, **0.835**, is exactly `rf.oob_score_`, and the probabilities match `oob_decision_function_`.

6. How many trees the OOB score needs

Key point: With very few trees, some observations have no OOB tree and the estimate is poor; from about 20 trees on, the OOB score tracks the test accuracy.




Figure 5: OOB score and test accuracy on the heart disease data as the number of trees grows

Each observation is OOB for only about 37% of the trees. With 5 trees, the chance that an observation is seen by **every** tree is the product of five chances of 0.632 (each tree sees it with chance 0.632):

$$0.632^5 \approx 0.10$$

So about 10% of the observations get no OOB prediction at all: 27 of the 242 here. scikit-learn warns: “Some inputs do not have OOB scores. This probably means too few trees were used”, and still counts those observations, with empty probabilities, in the score.

Figure 5 shows the effect on the heart disease data:

- with **5 trees**, the OOB score is 0.711, far below the test accuracy of 0.787: 27 observations have no OOB prediction, and the rest are judged by only a few trees each;
- from about **20 trees** on, the OOB score settles between 0.80 and 0.84, close to the test accuracy (0.84 to 0.87 on 61 noisy test observations).

Here the OOB score sits a little below the test accuracy at every forest size from 20 trees on (for example 0.818 against 0.836 with 500 trees). The gap is a known effect: on two-class data the OOB score tends to underrate a forest, most of all with balanced classes and few observations, as here (Janitza and Hornung, 2018). The reason is that a bootstrap sample

that misses observation i holds slightly more observations of the *other* class, so the trees that judge observation i lean a little towards the wrong class.

7. When the OOB score is available

Key point: Only with `bootstrap=True`; it works for classifiers, regressors and bagging ensembles alike.

- **The OOB score needs bootstrapping.** With `bootstrap=False` (pasting, see [pasting](#)) there are no OOB observations, and scikit-learn raises `ValueError: Out of bag estimation only available if bootstrap=True`.
- **The OOB score works for regression.** On a 5,000-observation sample of the California housing data (`fetch_california_housing` in scikit-learn: districts of California, 8 features, the median house value as target), a `RandomForestRegressor` gives an OOB R^2 of 0.764 against a test R^2 of 0.741. Its per-observation OOB predictions are in `oob_prediction_`.
- **The OOB score works for any bagging ensemble:** `BaggingClassifier` and `BaggingRegressor` take the same `oob_score` ([the bagging classifier's OOB score](#) and [the bagging regressor's](#)).

8. Summary

	Value on the heart data (mean of 50 splits)	What it needs
OOB score	0.820	<code>oob_score=True</code> , <code>bootstrap=True</code> , no extra training
5-fold cross-validation	0.820	5 more forests
Test accuracy (61 patients)	0.831	a held-out test set

- A tree’s OOB observations are the training observations its bootstrap sample missed: about 37%.
- OOB is per tree: practically every observation is seen by some trees and missed by others.
- Each observation is predicted only by its OOB trees; the share of correct predictions is the OOB score.
- With enough trees (here about 20 or more), the OOB score is a free estimate as good as cross-validation.

9. Sources

Built from

- CampusX, “OOB Score | Out of Bag Evaluation in Random Forest | Machine Learning”, YouTube, <https://www.youtube.com/watch?v=tdDhyFoSG94>
- StatQuest with Josh Starmer, “StatQuest: Random Forests Part 1 - Building, Using and Evaluating”, YouTube, https://www.youtube.com/watch?v=J4Wdy0Wc_xQ (the per-observation OOB vote of section 3, and tuning with the OOB score)

Other references

- Hastie, T., Tibshirani, R. and Friedman, J. (2009). *The Elements of Statistical Learning*, 2nd ed. Springer. §15.3.1 (out-of-bag samples).
- Janitza, S. and Hornung, R. (2018). On the overestimation of random forest’s out-of-bag error. *PLoS ONE* 13(8): e0201904.

10. Key terms

Term	Meaning
Out-of-bag (OOB) evaluation (G-1411)	Testing a bagging model by predicting each training observation with only the base models that never saw it
OOB prediction	An observation’s prediction from only the trees whose bootstrap sample missed it

Term	Meaning
<code>oob_score_</code>	The accuracy (classifier) or R^2 (regressor) of the OOB predictions
<code>oob_decision_function_</code>	Each training observation's class probabilities from its OOB trees
<code>oob_prediction_</code>	Each training observation's OOB prediction, for a regressor
Validation set	Data held back from training to check and tune a model before the final test