

K-Means from Scratch in Python

Contents

1. Overview	2
2. Prerequisites	3
3. The shape of the class	3
4. Step 2: random starting centroids	4
5. Step 3: assigning clusters	5
5.1 One distance formula for any number of features	6
5.2 The nearest centroid	8
6. Step 4: moving the centroids	9
7. Step 5: stopping	10
8. Testing the class	11
9. Bad random starts	13
10. Summary	15
11. Sources	15
12. Key terms	15

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Unsupervised learning ([Note ML-003](#)); Feature scaling ([Note ML-006](#)); Elbow method and WCSS ([Note ML-122](#)); Expectation maximization (EM) ([Note MA-074](#)).
- **Leads to:** Hierarchical clustering ([Note ML-125](#)); DBSCAN ([Note ML-126](#)).
- **Compare with:** Hierarchical clustering ([Note ML-125](#)); DBSCAN ([Note ML-126](#)); Gaussian mixture model (GMM) ([Note MA-073](#)).

1. Overview

Key point: We write k-means (G-996) ourselves as a class with one public method, `fit_predict`, and two helpers: `assign_clusters` (step 3) and `move_centroids` (step 4).

Coding an algorithm from scratch is the surest test that we understand it. This Note turns [the five steps of k-means](#) (pick k starting centroids, assign every point to the nearest, move each centroid to the mean of its points, repeat until nothing moves) into about 40 lines of Python. Figure 1 shows how the steps map onto the methods of the class.




Figure 1: The KMeans class: which method does which step of k-means

The class is in `kmeans.py` in this Note's folder; running `python kmeans.py` checks that it finds the same clusters as scikit-learn. The Notebook ([ML-124-kmeans-from-scratch.ipynb](#)) runs every experiment below.

2. Prerequisites

- [The five steps of k-means](#) and [WCSS](#) (G-2102; the total squared distance from each point to its cluster's centroid).
- KMeans in scikit-learn ([training k-means](#)), the student data and [boolean indexing](#).

3. The shape of the class

Key point: The constructor stores `k` and a limit on the number of rounds; `fit_predict` runs the loop and returns one cluster number per observation.

We copy scikit-learn's interface, so our class is used exactly like the real one:

Python: Using the class.
`from kmeans import KMeans`

```
km = KMeans(n_clusters=4, max_iter=100)
y_means = km.fit_predict(X)    # one cluster number per row
```

from kmeans import KMeans loads the class from our own file `kmeans.py`, not from `scikit-learn`.

The **constructor** (G-457), the method that runs when the object is created, takes two settings:

- **n_clusters**: k , the number of clusters (step 1). Default 2.
- **max_iter**: the largest number of assign-and-move rounds. Default 100. The loop stops earlier if the centroids (G-367) stop moving.

Python: The constructor.

```
class KMeans:
    def __init__(self, n_clusters=2, max_iter=100,
                 random_state=None):
        self.n_clusters = n_clusters
        self.max_iter = max_iter
        self.random_state = random_state
        self.centroids = None    # set by fit_predict
```

`random_state` is our addition: it fixes the random start so that results can be repeated.

4. Step 2: random starting centroids

Key point: Pick k different row numbers at random and take those rows as the first centroids.

`X` is a 2-D NumPy array: with 100 points in 2 columns, it holds 100 rows of 2 numbers each. Each row is an **observation** (one record of the data) and each column a **feature** (an input variable). To start, we pick k different row numbers between 0 and 99 at random, then take those rows. Figure 2 shows one such start on the 200 students: the four chosen rows are real students. They fall in pairs inside two groups, and two groups start with no centroid at all; the loop of Figure 6 still recovers all four.




Figure 2: Step 2 on the students: rows 47, 16, 35 and 159, picked at random, become the first four centroids

Python: Choosing the starting centroids.

```
rng = np.random.default_rng(self.random_state)
random_index = rng.choice(X.shape[0],
                           size=self.n_clusters,
                           replace=False)
self.centroids = X[random_index].astype(float)
```

`X.shape[0]` is the number of rows. `rng.choice(100, size=2, replace=False)` returns 2 different numbers from 0 to 99, for example [26, 17]. `X[[26, 17]]` then picks rows 26 and 17. `astype(float)` lets the centroids hold decimals once they move.

Extra: The same can be written with Python's `random` module: `random.sample(range(X.shape[0]), k)` also returns `k` different numbers. NumPy's generator is used here because it takes a seed per object, so two models never share random state.

5. Step 3: assigning clusters

Key point: For every row, compute its distance to every centroid and record the index of the smallest one.

5.1 One distance formula for any number of features

Key point: The squared Euclidean distance (G-715) is the dot product (G-634) of the difference vector with itself, so `np.sqrt(np.dot(a - b, a - b))` works for 2, 3 or 100 features.

The Euclidean distance (the straight-line distance between two points, see [the KNN imputer](#)) gains one term per feature, so code written term by term would change with the number of features. The sum of squared differences is the dot product (multiply matching entries and add, see [projecting one point](#)) of the difference vector with itself, which needs no change:

$$d(a, b) = \sqrt{(b - a) \cdot (b - a)}$$

For $a = (1, 2)$ and $b = (4, 5)$:

$$b - a = (3, 3)$$

$$(3, 3) \cdot (3, 3) = 18$$

$$d = \sqrt{18} \approx 4.24$$

With a third feature, $(1, 2, 3)$ and $(4, 5, 6)$, the same code gives:

$$b - a = (3, 3, 3)$$

$$(3, 3, 3) \cdot (3, 3, 3) = 27$$

$$d = \sqrt{27} \approx 5.196$$

In Figure 3, watch the two dashed sides: each feature adds one squared side, and the dot product adds them all at once.




Figure 3: The distance from $(1, 2)$ to $(4, 5)$: the difference vector $(3, 3)$, its dot product with itself, and its square root

Figure 4 plays the same calculation step by step: the difference vector, its dot product with itself, the square root. Watch the last part: a third feature adds one more term to the dot product, and the line of code stays the same.




Figure 4: The distance from $(1, 2)$ to $(4, 5)$ as a dot product, step by step; then a third feature is added, $(1, 2, 3)$ to $(4, 5, 6)$, and the same code gives 5.196

Python: The distance in NumPy.

```
np.sqrt(np.dot(b - a, b - a))    # 4.24 for (1, 2), (4, 5)
```

$b - a$ subtracts element by element. `np.dot(u, v)` multiplies matching elements and adds them up.

5.2 The nearest centroid

Key point: With 100 rows and 2 centroids we compute 200 distances; each row gets the index of its smaller one.

A nested loop visits every row and, inside, every centroid. With 100 rows and 2 centroids the inner line runs:

$$100 \times 2 = 200$$

times. For each row we keep the position of the smallest distance: 0 if the first centroid is nearer, 1 if the second.

Python: `assign_clusters`.

```
def assign_clusters(self, X):
    cluster_group = []
    for row in X:
        distances = [np.sqrt(np.dot(row - c, row - c))
                     for c in self.centroids]
        cluster_group.append(int(np.argmin(distances)))
    return np.array(cluster_group)
```

The square brackets build a list of k distances, one per centroid. `np.argmin` returns the position of the smallest value, the same as `distances.index(min(distances))`. The result has one entry per row, such as `[0, 0, 1, ...]`, and is returned as a NumPy array so it can be used for boolean indexing.

6. Step 4: moving the centroids

Key point: For each cluster k , select its rows with `X[cluster_group == k]` and take the mean of each column with `mean(axis=0)`.

The cluster numbers tell us which rows belong to which cluster. The new centroid of a cluster is the mean of each column over its rows (see [steps 3 and 4 on two features](#)).

Take five rows and their clusters:

Row	Values	Cluster
0	(1, 2)	0
1	(2, 3)	1
2	(3, 4)	0
3	(4, 5)	0
4	(5, 6)	1

`X[cluster_group == 0]` keeps rows 0, 2 and 3. Their column means are:

$$(1 + 3 + 4)/3 = 2.67$$

$$(2 + 4 + 5)/3 = 3.67$$

So the new centroid of cluster 0 is (2.67, 3.67). Cluster 1 gets (3.5, 4.5). Figure 5 draws both: each cross lands in the middle of its own rows, and the other cluster's rows (grey) play no part.




Figure 5: Moving the centroids for the five-row example: each new centroid is the column mean of its own rows

Python: `move_centroids`.

```
def move_centroids(self, X, cluster_group):
    new_centroids = []
    for k in range(self.n_clusters):
        members = X[cluster_group == k]
        new_centroids.append(members.mean(axis=0))
    return np.array(new_centroids)
```

`axis=0` averages down each column, giving one mean per column. Without it, `mean()` averages all six numbers into a single value (3.17), which is not a point at all.

Extra: A cluster can end up empty: no row is nearest to its centroid. The mean of zero rows is undefined, and a loop over the cluster numbers that actually occur (for example with `np.unique(cluster_group)`) would silently return fewer than `k` centroids. The file `kmeans.py` keeps an empty cluster's old centroid instead. scikit-learn moves it to the data point that is farthest from its own centroid (scikit-learn source, `_k_means_common.pyx`).

7. Step 5: stopping

Key point: Save the centroids before moving them; if the new ones equal the old ones, break out of the loop.

Python: The loop inside `fit_predict`.

```
for i in range(self.max_iter):
    cluster_group = self.assign_clusters(X)
```

```

old_centroids = self.centroids
self.centroids = self.move_centroids(X, cluster_group)
if np.allclose(old_centroids, self.centroids):
    break
return cluster_group

```

`break` leaves the loop at once. `np.allclose` is True when every pair of numbers is equal up to a tiny rounding error. `(old == new).all()` also works, but computed means can differ in the last decimal place.

Figure 6 runs the whole loop on the 200 students from the start of Figure 2. Watch the largest move shrink round by round; in round 4 the centroids do not move at all, `np.allclose` is True, and the loop breaks.




Figure 6: The assign-and-move loop on the students: colours change at step 3, crosses move at step 4, and the loop stops in round 4

The loop ends in one of two ways (Figure 1): the centroids stop moving, or `max_iter` rounds have run. Either way, `fit_predict` returns the cluster number of every row.

8. Testing the class

Key point: On two, three and four blobs, and on the 200 students, the class finds the right groups in 2 to 5 rounds.

We generate test data with `make_blobs` (see [k-means in three dimensions](#)): 100 points around 2, 3 or 4 centres, each with spread 1. Then we run our class with the matching `k`, and finally on the 200 students with `k = 4`.




Figure 7: The from-scratch class on 2, 3 and 4 blobs and on the students; crosses are the final centroids

Figure 7 shows all four results: every group found, every centroid in the middle of its group. The loops finished after 2, 5, 3 and 4 rounds. Plotting needs one scatter call per cluster, so the plotting code (not the class) must be extended each time k grows.

Python: The self-check in `kmeans.py`.

```
from sklearn.metrics import adjusted_rand_score
```

```
ours = KMeans(n_clusters=3, random_state=1).fit_predict(X)
```

```
theirs = SkKMeans(n_clusters=3, random_state=0).fit_predict(X)
```

```
assert adjusted_rand_score(ours, theirs) == 1.0
```

Cluster numbers are only names, so we cannot compare them directly. `adjusted_rand_score` is 1.0 when two labelings put exactly the same points together, whatever the numbers.

9. Bad random starts

Key point: If the starting centroids fall badly, k-means gets stuck in poor clusters, a **local optimum** (G-1111); more rounds do not help, but restarting from other random points does.

Sometimes the class returns wrong clusters. Blaming `max_iter` and raising it is tempting, but the round counts say otherwise. Across 30 random starts on the four blobs, k-means never needed more than 10 rounds; at most 10 of the 100 allowed rounds were ever used.

Figure 8 shows the real cause on the students. With one random start (left), two centroids end up in the bottom-left group. That group gets split in two, while two real groups, the top ones, have to share one centroid placed in the empty space between them. After 3 rounds nothing changes any more, so the algorithm stops, stuck. Raising `max_iter` to 500 gives exactly the same result.




Figure 8: The same students and $k = 4$ from two random starts: stuck in poor clusters (left) and the right clusters (right)

Figure 9 runs both starts round by round. Watch the status in each panel's title. The bad start stops in round 3 and the good start in round 4: in both, the centroids no longer move, so the loop breaks long before `max_iter`. From then on nothing changes in either panel, and the bad run stays at WCSS 2,280. Extra rounds are never used, so they cannot repair a bad start.




Figure 9: The bad start (left) and the good start (right) on the students, one assign-and-move round per frame, with the WCSS after each round. The bad run stops in round 3 at WCSS 2,280, the good run in round 4 at WCSS 682

The two runs can be compared by their WCSS (see [WCSS](#)): 2,280 for the bad start and 682 for the good one. Lower is better, so the fix is simple (ESL §14.3.6):

1. run k-means from several random starts;
2. keep the run with the lowest WCSS.

In the Notebook, the best of 10 starts is the good clustering (WCSS 682). scikit-learn's `n_init` (see [training k-means](#)) does exactly this job, and its `k-means++` start spreads the first centroids out.

Extra: The class also stores the WCSS of its result as `inertia_`, like scikit-learn:

```
self.inertia_ = ((X - self.centroids[cluster_group]) ** 2).sum()
self.centroids[cluster_group]
```

repeats, for every row, the centroid of its own cluster; subtracting, squaring and summing gives the WCSS. With it, the elbow method takes one loop over k . Taking the best of 10 starts for each k , the students give WCSS 29,958, 4,184, 2,363, 682, 514, ...: the elbow is at $k = 4$, as with scikit-learn.

10. Summary

Method	Step	What it does
<code>__init__</code>	1	Stores <code>k</code> (<code>n_clusters</code>) and the round limit (<code>max_iter</code>)
<code>fit_predict</code>	2	Picks <code>k</code> random rows as the first centroids, then runs the loop
<code>assign_clusters</code>	3	Distance from every row to every centroid; keeps the index of the nearest
<code>move_centroids</code>	4	Each centroid becomes the column means of its rows
check in <code>fit_predict</code>	5	Stops when the centroids no longer move

- The Euclidean distance for any number of features: `np.sqrt(np.dot(a - b, a - b))`.
- `X[cluster_group == k].mean(axis=0)` is the new centroid of cluster `k`.
- On the four blobs, k-means stopped within 10 rounds from all 30 starts; the wrong result on the students came from a bad random start, not from too few rounds.
- Fix bad starts by restarting several times and keeping the lowest WCSS.

11. Sources

Built from

- CampusX, “K-Means Clustering Algorithm From Scratch In Python | ML Algorithms From Scratch”, YouTube, <https://www.youtube.com/watch?v=MFraC1JObUo>
- StatQuest with Josh Starmer, “StatQuest: K-means clustering”, YouTube, <https://www.youtube.com/watch?v=4b5d3muPQmA> (a bad start is fixed by rerunning from new starts and keeping the lowest total variation, section 9)

Other references

- ESL: Hastie, Tibshirani and Friedman, *The Elements of Statistical Learning*, 2nd ed., Springer, 2009, §14.3.6 (k-means; restarting from several random starts and keeping the lowest objective).
- scikit-learn source, `_k_means_common.pyx`: the function `_relocate_empty_clusters_dense` in `sklearn/cluster/_k_means_common.pyx`, scikit-learn 1.9.

12. Key terms

Term	Meaning
Constructor (<code>__init__</code>)	The method that runs when an object is created and stores its settings
<code>max_iter</code> (G-113)	The largest number of assign-and-move rounds k-means may run
<code>np.argmin</code>	NumPy function returning the position of the smallest value
<code>mean(axis=0)</code>	The mean of each column of an array
Local optimum (k-means)	A clustering where k-means has stopped but a better one exists, caused by a bad start
Adjusted Rand score (G-175)	A number that is 1.0 when two labelings group the points identically, whatever the label numbers