


Imbalanced Data: Undersampling, Oversampling, SMOTE and Cost-Sensitive Learning

Contents

1. Overview	2
2. What imbalanced data looks like	3
2.1 Our dataset	3
3. Two problems with imbalanced data	4
3.1 Bias towards the majority class	4
3.2 Accuracy is not reliable	5
4. Why imbalanced data matters	6
5. Random undersampling	7
6. Random oversampling	8
7. SMOTE	10
7.1 Interpolation	10
7.2 The algorithm	10
7.3 Advantages and disadvantages	12
7.4 Resampling only the training data	14
8. Ensemble methods: the balanced random forest	15
9. Cost-sensitive learning	16
9.1 Class weights	16
9.2 A custom loss function	18
10. Comparing the techniques	20
10.1 Other techniques	22
11. Summary	23
12. Sources	23
13. Key terms	24

Where this fits: Pipeline map steps 3 (Understand data), 5 (Engineer features), 8 (Model). See the [Course map](#).



- **Builds on:** K-nearest neighbours ([Note ML-006](#)); Cross-validation ([Note ML-028](#)); Log loss (binary cross entropy) ([Note ML-072](#)); ROC curve and AUC ([Note ML-077](#)); Random forest ([Note ML-102](#)).
- **Leads to:** ANN for classification ([Note DL-011](#)).

1. Overview

Key point: When one class is rare, a model learns mostly the common class; we fix this by changing the data, the ensemble, or the cost of mistakes.




Figure 1: Five ways to handle imbalanced data, in three groups

Imbalanced data (G-921) was defined in the [accuracy Note](#), section 6: data in which one class is much rarer than another. This Note shows what goes wrong when we train on it and teaches the five techniques of Figure 1:

- **Change the data** (resampling): random undersampling, random oversampling, SMOTE.
- **Change the ensemble:** the balanced random forest.

- **Change the learning** (cost-sensitive learning): class weights and a custom loss function.

The Notebook (`notebook.ipynb`) runs every technique on the same dataset.

2. What imbalanced data looks like

Key point: One class (the majority) has far more observations than another (the minority); this happens with two classes or with many.

Take a college where we know each student's CGPA and IQ and whether they were placed. Each student is one **observation** (G-1374; one record, one row of the data table); CGPA and IQ are **features** (G-772; input variables, one column each); placed or not is the **target** (G-1949), the output we predict. At a top college, most students are placed: out of 500 final-year students, perhaps 450 placed and 50 not. Placed is the **majority class** (G-1145), the class with many observations; not placed is the **minority class** (G-1229), the rare one.

The same can happen with three or more classes. Suppose 10 more students did not sit for placements at all, a third class. The data now has one large class and two small ones, and it is still imbalanced.

2.1 Our dataset

Key point: 400 observations, two features, about 7% in class 0; the training set has 299 observations of class 1 and only 21 of class 0.

To see the effects, we use a synthetic dataset made with scikit-learn's `make_classification`: 400 observations, two features `x1` and `x2`, and two classes. Class 1 is the majority and class 0 the minority. After an 80/20 split, the training set has 320 observations (299 of class 1, 21 of class 0) and the test set 80 observations (73 and 7) (Figure 2).



Figure 2: The Note’s dataset: 299 against 21 observations in the training set, 73 against 7 in the test set

Python: Making an imbalanced dataset.

```
from sklearn.datasets import make_classification
```

```
X, y = make_classification(
    n_samples=400, n_features=2, n_informative=2,
    n_redundant=0, n_clusters_per_class=2,
    weights=[0.07], class_sep=0.8, flip_y=0,
    random_state=59)
```

`weights=[0.07]` puts about 7% of the rows in class 0. `n_clusters_per_class=2` gives each class two blobs, so no straight line separates them well.

3. Two problems with imbalanced data

Key point: The model becomes biased towards the majority class, and accuracy hides it.

3.1 Bias towards the majority class

Key point: The algorithm sees mostly majority observations, so it learns them well and the minority badly.

A model learns from the observations it sees. With 299 observations of class 1 and 21 of class 0, almost all of the training signal is about class 1. The model learns to get class 1 right and pays little attention to class 0.




Figure 3: Left: the training observations per class. Right: the line logistic regression learns on this data

Figure 3 shows the result for logistic regression. The decision boundary sits low, below almost all the red points. The model keeps every blue point on the correct side and gives up on most of the red ones: it is **biased towards the majority class**.

3.2 Accuracy is not reliable

Key point: The model in Figure 3 has 92.5% accuracy but finds only 1 of the 7 minority observations in the test set.

The classification report of the model in Figure 3, on the 80 test observations:

Class	Precision	Recall	F1	Observations
0 (minority)	1.00	0.14	0.25	7
1 (majority)	0.92	1.00	0.96	73
Accuracy			0.925	80

The accuracy, 92.5%, looks excellent. The recall of class 0 is 0.14: of the 7 minority observations, the model found only 1. If class 0 is the class we care about, this model is nearly useless. Figure 4 shows where the 92.5% comes from: the 73 majority observations, all right.




Figure 4: Confusion matrices on the 80 test observations. Left: logistic regression finds 1 of the 7 minority observations (red row). Right: a model that always answers class 1 finds none, and still scores 91.25%

A “dumb” model that always answers class 1 already scores $73 / 80 = 91.25\%$ here, the same trap as in the [accuracy Note](#), section 6. So on imbalanced data we judge models by the minority class’s precision, recall and F1 ([precision Note](#)) and by ROC AUC ([ROC Note](#)).

Python: The classification report.

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

clf = LogisticRegression().fit(X_train, y_train)
print(classification_report(y_test, clf.predict(X_test)))
```

4. Why imbalanced data matters

Key point: Many of the problems classical ML is used for are imbalanced, and in most of them the rare class is the one we need to find.

Domain	Problem	The rare class
Finance	fraud detection	fraudulent transactions
Finance	credit risk assessment	customers who will not repay a loan
Healthcare	rare disease prediction	patients who have the disease (say 1%)
Manufacturing	predicting machine failure	machines about to break down
Consumer internet	churn prediction	users who will leave the platform next month

Domain	Problem	The rare class
Environmental science	earthquake or volcano prediction	days with an eruption or quake

In every row of the table, predicting the minority class correctly is the whole point. A model that is biased towards the majority, like the one in Figure 3, fails exactly where it matters.

5. Random undersampling

Key point: Randomly drop majority observations until both classes have the same number of observations.

The simplest fix is to make the data balanced. **Random undersampling** (G-1619) keeps every minority observation and draws, at random and without replacement, the same number of majority observations. The other majority observations are thrown away.

With 900 majority and 100 minority observations, we keep 100 random majority observations and all 100 minority observations: 200 observations instead of 1,000. On our training set, 320 observations (299 + 21) become **42** observations (21 + 21).

Python: Random undersampling with the imbalanced-learn library.

```
import numpy as np
from imblearn.under_sampling import RandomUnderSampler
```

```
rus = RandomUnderSampler(random_state=42)
X_under, y_under = rus.fit_resample(X_train, y_train)
X_under.shape           # (42, 2)
np.bincount(y_under)    # [21, 21]
```

imbalanced-learn (G-922; imported as `imblearn`) collects the techniques of this Note behind one method, `fit_resample`, which returns the new inputs and labels. `np.bincount` counts how many labels are 0, 1, and so on.




Figure 5: The training data after random undersampling, random oversampling and SMOTE. Solid line: logistic regression trained on the resampled data; dashed line: trained on the original data

Figure 5, left, shows the 42 observations and the new boundary. Logistic regression now treats both classes as equally important, and the line moves up through the middle of the data. The minority recall on the test set rises from 0.14 to **1.00** (all 7 found), while the minority precision falls from 1.00 to 0.33: the model now also flags many majority observations as class 0.

Extra: A straight line cannot separate this data well, because each class has two blobs. Undersampling fixes the bias, not the shape of the boundary; an algorithm that draws curved boundaries, such as a random forest, may do better.

Advantages:

- Undersampling removes the bias towards the majority class.
- The data gets smaller, so training is faster. With 10 lakh majority and 1 lakh minority observations, we train on 2 lakh observations instead of 11 lakh.

Disadvantages:

- We throw data away: here 278 of 320 observations. Those observations may hold information the model needs.
- The sample can be unrepresentative. A random sample is usually fine, but by bad luck it can keep one kind of majority observation and drop another.

6. Random oversampling

Key point: Randomly copy minority observations until both classes have the same number of observations.

Random oversampling (G-1613) goes the other way: keep every observation and draw

minority observations at random *with replacement* until the minority is as large as the majority. With 900 and 100 observations, we draw 900 minority observations from the 100, so each one appears 9 times on average.

On our training set, 320 observations become **598** ($299 + 299$) (Figure 6).




Figure 6: The training set’s class counts before and after each resampler. Undersampling cuts the majority to 21; oversampling copies the minority up to 299; SMOTE adds 278 new minority points

Each of the 21 minority observations now appears between 9 and 21 times. In a scatter plot the copies sit exactly on top of each other, which is why Figure 5, middle, draws each repeated observation as one dot whose size grows with the number of copies.

Python: Random oversampling.

```
from imblearn.over_sampling import RandomOverSampler

ros = RandomOverSampler(random_state=42)
X_over, y_over = ros.fit_resample(X_train, y_train)
np.bincount(y_over)      # [299, 299]
```

The boundary moves up much as with undersampling (Figure 5, middle). The minority recall rises to **0.86** (6 of 7) and the precision falls to 0.32.

Advantage: random oversampling removes the bias towards the majority class without throwing any data away.

Disadvantages:

- The data grows: here from 320 to 598 observations. A 1 GB dataset can become nearly 2 GB.

- Copies are not new information. An observation repeated 21 times looks very important to the algorithm, which can **overfit** (G-1429) to it, especially decision trees: the tree carves out tiny regions around the repeated points, with more splits and more leaves (Chawla et al. 2002, §4.1).

7. SMOTE

Key point: Instead of copying minority observations, create new ones on the line between a minority observation and one of its minority neighbours.

SMOTE (G-1825; Synthetic Minority Over-sampling Technique) is the most popular technique for imbalanced data. SMOTE is an oversampling technique: it grows the minority class until the classes are equal. The difference is that it makes new observations instead of copies, so the model does not see the same point again and again.

7.1 Interpolation

Key point: A new point is placed a random fraction of the way from a minority point to its neighbour.

SMOTE creates observations by **interpolation** (G-964): placing a new point on the straight segment between two existing points.

The new point, step by step:

1. **In words:** start at the minority point x ; move a random fraction λ (between 0 and 1) of the way towards its neighbour n .
2. **Formula:**

$$x_{\text{new}} = x + \lambda(n - x), \quad 0 \leq \lambda \leq 1$$

With $\lambda = 0$ the new point is x itself; with $\lambda = 1$ it is the neighbour. The same formula is often written $x - \lambda(x - n)$.

3. **Example:** $x = (1, 1)$, $n = (2, 2)$ and $\lambda = 0.5$:

$$x_{\text{new}} = (1, 1) + 0.5 \times ((2, 2) - (1, 1)) = (1, 1) + (0.5, 0.5) = (1.5, 1.5)$$

The new point lies exactly halfway between the two.

7.2 The algorithm

Key point: Find each minority point's 5 nearest minority neighbours, then repeat: random point, random neighbour, random fraction, new point.




Figure 7: SMOTE step by step: a random minority point and its 5 nearest minority neighbours, one neighbour chosen, a new point between them, and the result after a few repetitions

Figure 7 shows the steps:

1. Take only the minority observations, and find the k nearest minority neighbours of each one, usually $k = 5$ (a KNN search, as in the [KNN Note](#)).
2. Pick a minority point at random.
3. Pick one of its k neighbours at random.
4. Draw a random λ between 0 and 1 and create the new point $x + \lambda(n - x)$.
5. Repeat steps 2 to 4 until the minority class is as large as the majority.

Python: SMOTE written by hand in a few lines of NumPy, to see every step.

```
from sklearn.neighbors import NearestNeighbors

def smote(X, y, k=5, seed=42):
    rng = np.random.default_rng(seed)
    mino = X[y == 0]
    # step 1: k nearest minority neighbours
    # (k + 1, since each point is its own nearest)
    nn = NearestNeighbors(n_neighbors=k + 1).fit(mino)
    nbrs = nn.kneighbors(mino, return_distance=False)
    nbrs = nbrs[:, 1:]
    n_new = (y == 1).sum() - len(mino)
    i = rng.integers(len(mino), size=n_new) # step 2
    j = nbrs[i, rng.integers(k, size=n_new)] # step 3
    lam = rng.random((n_new, 1))           # step 4
```

```

    new = mino[i] + lam * (mino[j] - mino[i])
    return (np.vstack([X, new]),
            np.r_[y, np.zeros(n_new)])
rng.integers(21, size=278) draws 278 random positions at once, so all new points
are made in one go instead of in a loop.

```

Python: SMOTE with imbalanced-learn, the version we use in practice.

```
from imblearn.over_sampling import SMOTE
```

```

sm = SMOTE(k_neighbors=5, random_state=42)
X_smote, y_smote = sm.fit_resample(X_train, y_train)
new = X_smote[len(X_train):]    # the 278 synthetic rows

```

imbalanced-learn builds each new point exactly as our function does (imbalanced-learn source, `_smote/base.py`): a random minority point, one of its `k_neighbors` nearest minority neighbours, and a random factor between 0 and 1. `fit_resample` returns the original observations first and the synthetic observations after them.

On our data, SMOTE creates 278 new observations: 320 observations become 598 (299 + 299). Figure 5, right, shows them in orange, lying on segments between red points. The boundary moves up again; the minority recall rises to **0.71** (5 of 7) and the precision falls to 0.29.

7.3 Advantages and disadvantages

Key point: SMOTE reduces bias without copies, but it fails on categorical features, is slow on big data, depends on `k`, spreads outliers and may invent unrealistic observations.

Advantages:

- SMOTE reduces the bias towards the majority class.
- There are no duplicates, so there is less risk of overfitting than with random oversampling (section 10 tests this on real data).

SMOTE is popular, but whether to use it is much debated, because of five disadvantages:

1. **SMOTE cannot handle categorical features.** Interpolating between category 0 and category 1 can give 0.7, which is not a category (Figure 8, left).
2. **SMOTE is slow on large data.** SMOTE needs a nearest-neighbour search, which measures many distances and becomes slow with many observations or many features.
3. **SMOTE depends on k .** With $k = 1$, every new point lies on the segment to the single nearest neighbour, so the new points have little variety. With a very large k , neighbours can be far away and new points appear all over the feature space. The middle is hard to find; $k = 5$ is the usual default, but the best value depends on the data. Figure 9 shows all three cases on our training data.
4. **SMOTE is sensitive to outliers.** A noisy minority point far from the others gets picked too, and the segments to its neighbours cross the majority region, so one noisy observation creates more noisy observations (Figure 8, right). Our own data shows this:

the red outlier at the bottom of Figure 5, right, pulls a trail of orange points down with it.

5. **The new observations may not be realistic.** Nothing guarantees that the invented points follow the true distribution of the population. They follow the straight lines between existing points, which real data may not.




Figure 8: Two SMOTE failures: a new point between two categories, and an outlier that spreads new points into the majority




Figure 9: The 278 new SMOTE points (orange) on our training data for $k = 1$, 5 and 20. With $k = 1$ the new points line up on a few short segments; with $k = 20$ almost every minority point is a neighbour of every other, and the new points spread over the area between them, into the majority.

Extra: imbalanced-learn has variants aimed at some of these problems (imbalanced-learn user guide, Over-sampling). **SMOTENC** handles a mix of numeric and categorical features (for a categorical feature it takes the most common category among the neighbours instead of interpolating). **BorderlineSMOTE** creates points only near the class boundary, and **ADASYN** creates more points where the minority is hardest to learn.

7.4 Resampling only the training data

Key point: Resample after the split, and inside each cross-validation fold; resampling first makes the scores look far better than they are.

We resample **only the training set**, after the train-test split. The test set must keep its real class balance, or the scores no longer describe real data. And SMOTE points made from test observations would leak test information into training (data leakage, the [toy project Note](#)).

The same holds inside cross-validation ([pipelines Note](#)). If we run SMOTE on the whole training set and then cross-validate, the validation folds are full of synthetic points built from the training folds' own minority observations, and they no longer have the real class balance (imbalanced-learn user guide, Common pitfalls).

imbalanced-learn's own **Pipeline** fixes this: it runs the resampler only while fitting, so each fold resamples its own training part and validates on untouched observations. On our training set, with 5-fold cross-validation and the F1 score of class 0:

Setup	Cross-validated F1 (class 0)
No resampling	0.200
SMOTE inside an imblearn Pipeline (correct)	0.426
SMOTE on all observations, then cross-validation (leaky)	0.885

The leaky setup reports more than twice the honest score, and almost all of the gap comes from the synthetic validation points: scoring only the real observations of the same leaky folds gives 0.398, close to the honest 0.426 (Notebook).

Python: SMOTE inside cross-validation.

```
from imblearn.pipeline import Pipeline
from sklearn.metrics import f1_score, make_scorer
from sklearn.model_selection import (StratifiedKFold,
                                     cross_val_score)

pipe = Pipeline([("smote", SMOTE(random_state=42)),
                 ("model", LogisticRegression())])
cv = StratifiedKFold(n_splits=5, shuffle=True,
                    random_state=42)
f1_minority = make_scorer(f1_score, pos_label=0)
cross_val_score(pipe, X_train, y_train, cv=cv,
```

```
scoring=f1_minority).mean()    # 0.426
```

scikit-learn's own Pipeline cannot hold a resampler, because it only knows fit/transform steps. StratifiedKfold keeps the class ratio the same in every fold, so each fold gets some of the 21 minority observations. `make_scorer(f1_score, pos_label=0)` scores the F1 of class 0 instead of class 1.

8. Ensemble methods: the balanced random forest

Key point: A balanced random forest trains every tree on a balanced sample: the minority observations (or a bootstrap sample of them) plus the same number of random majority observations.

Ensemble methods ([ensemble learning Note](#)) such as bagging and random forests can be changed to handle imbalanced data. A random forest ([random forest Note](#)) trains many decision trees, each on a random sample of the observations, and combines their votes. A **balanced random forest** (G-256) changes only how each sample is drawn: every sample is balanced, even though the data is not.




Figure 10: A balanced random forest: every tree is trained on a balanced sample, then the trees vote

In Figure 10, the training data has 900 majority and 300 minority observations. Each sample takes 300 minority observations and 300 majority observations drawn at random, 600 observations in all, and one tree is trained on it. To predict, every tree votes and the majority vote wins: here 1, 0 and 1 give 1.

Undersampling (section 5) threw most majority observations away for good. Here each tree throws different observations away, so together the trees still see most of the majority class. On our data, 100 trees each draw 21 of the 299 majority observations, so a given observation is missed by every tree with probability $(1 - 1/299)^{2100} \approx 0.001$.

Python: A balanced random forest with imbalanced-learn.

```
from imblearn.ensemble import BalancedRandomForestClassifier
```

```
brf = BalancedRandomForestClassifier(
    n_estimators=100, random_state=42)
brf.fit(X_train, y_train)
y_pred = brf.predict(X_test)
```

The balanced forest is used like any scikit-learn classifier. By default (`sampling_strategy="all"`, `replacement=True`), every tree gets 21 observations of each class, drawn with replacement: a bootstrap sample of the minority and an equally small random sample of the majority.

On our data, the balanced forest finds 5 of the 7 minority test observations (recall **0.71**, precision 0.28). An ordinary random forest finds 3 (recall 0.43, precision 0.50).

Extra: imbalanced-learn has other balanced ensembles (imbalanced-learn API reference): `BalancedBaggingClassifier` (bagging with a resampler on every sample), `EasyEnsembleClassifier` (boosted models on balanced samples) and `RUSBoostClassifier` (boosting with random undersampling at every round).

9. Cost-sensitive learning

Key point: Leave the data alone and change the learning itself, so that mistakes on the minority class cost more.

Cost-sensitive learning (G-493) changes how the algorithm learns, so that it handles imbalanced data better. There are two ways: class weights, and a custom loss function.

9.1 Class weights

Key point: Give the minority class a bigger weight; each of its mistakes then adds more to the loss, and the model works harder to avoid them.

`class_weight` was introduced in the [logistic regression hyperparameters Note](#), section 5.1: it multiplies each observation's loss by a weight for its class. Here we set the weights by hand.

With 900 observations of class 1 and 100 observations of class 0, we might give class 0 a weight of 10 and class 1 a weight of 1. Every mistake on class 0 then costs 10 times as much. For an algorithm trained by gradient descent ([gradient descent Note](#)), those mistakes produce 10 times larger updates, so the parameters move more to correct them.

Python: Class weights in logistic regression.

```
clf = LogisticRegression(class_weight={0: 5, 1: 1})
clf.fit(X_train, y_train)
```

The dictionary maps each class label to its weight. `class_weight="balanced"` instead sets each weight to $(\text{observations}) / (2 \times \text{observations of that class})$: here $320 / 42 = 7.62$ for class 0 and $320 / 598 = 0.54$ for class 1.




Figure 11: The line logistic regression learns as the weight of class 0 grows from 1 to 50

Figure 11 shows the effect. With weight 1 (no weighting), the line sits below the red points; at 5, 25 and 50 it moves further up into the majority class. The bigger the weight, the more pressure on the model to classify the minority correctly:

Weight of class 0	Accuracy	Precision (0)	Recall (0)
1	0.925	1.00	0.14
5	0.900	0.44	0.57
25	0.788	0.29	1.00
50	0.700	0.23	1.00

The weight is a dial: raise it and recall rises while precision falls. We pick the value that gives the best result for the problem, for example by cross-validation.

Most scikit-learn classifiers that minimise a loss accept `class_weight`: `LogisticRegression`, `SVC` ([SVM Note](#)), `DecisionTreeClassifier`, `RandomForestClassifier` ([random forest hyperparameters Note](#)). KNN and Naive Bayes do not, because they do not learn by minimising a loss.

Extra: Among scikit-learn’s boosting models, `HistGradientBoostingClassifier` accepts `class_weight`, but `GradientBoostingClassifier` and `AdaBoostClassifier` do not. For those, pass weights per observation through `fit(X, y, sample_weight=...)`, which many scikit-learn models accept (`KNeighborsClassifier` does not).

9.2 A custom loss function

Key point: Some libraries (gradient boosting, XGBoost, LightGBM) let us write our own loss, so we can set exactly how much each kind of mistake costs.

The losses we have used so far are standard ones: mean squared error, log loss ([log loss Note](#)). Some algorithms let us replace them with a **custom loss function** (G-523; also called a custom objective) that fits our problem. Suppose missing a minority observation (a false negative, if the minority is the positive class) is worse than a false alarm: we write a log loss that charges more for that mistake.

The weighted log loss, step by step:

1. **In words:** the usual log loss, but each observation’s loss is multiplied by the cost of its class: a for observations of class 1, b for observations of class 0.
2. **Formula:** with p_i the predicted probability of class 1,

$$L = -\frac{1}{n} \sum_{i=1}^n \left[a y_i \log p_i + b (1 - y_i) \log(1 - p_i) \right]$$

3. **Example:** a class 0 observation ($y = 0$) that the model gives $p = 0.9$, so only 0.1 for its true class. With $a = 1$, $b = 3.5$:

$$\text{loss} = -3.5 \times \log(1 - 0.9) = -3.5 \times (-2.303) = 8.06$$

instead of 2.303 with $b = 1$: the same mistake now costs 3.5 times as much.

To train with a custom loss, a gradient boosting library needs its derivatives with respect to the model’s raw output z (where $p = \sigma(z)$, the sigmoid). The symbol $\partial L_i / \partial z_i$ is the slope of observation i ’s loss when only z_i changes. Take the loss of one observation, $L_i = -[a y_i \log p_i + b (1 - y_i) \log(1 - p_i)]$, and use $\partial p_i / \partial z_i = p_i(1 - p_i)$ (the [sigmoid derivative Note](#)). One step per line:

$$\frac{\partial L_i}{\partial p_i} = -\frac{a y_i}{p_i} + \frac{b (1 - y_i)}{1 - p_i}$$

$$\frac{\partial L_i}{\partial z_i} = \frac{\partial L_i}{\partial p_i} \times p_i(1 - p_i) = -a y_i (1 - p_i) + b (1 - y_i) p_i$$

$$\frac{\partial^2 L_i}{\partial z_i^2} = [b (1 - y_i) + a y_i] \times p_i(1 - p_i)$$

The two results are

$$\frac{\partial L_i}{\partial z_i} = b(1 - y_i) p_i - a y_i (1 - p_i), \quad \frac{\partial^2 L_i}{\partial z_i^2} = p_i(1 - p_i)(a y_i + b(1 - y_i))$$

For the class 0 observation above ($y = 0, p = 0.9, a = 1, b = 3.5$):

$$\frac{\partial L_i}{\partial z_i} = 3.5 \times 1 \times 0.9 - 0 = 3.15$$

$$\frac{\partial^2 L_i}{\partial z_i^2} = 0.9 \times 0.1 \times (0 + 3.5) = 0.315$$

Python: The custom loss in XGBoost.

```
import xgboost as xgb

def weighted_logloss(a, b):
    def objective(z, dtrain):
        y = dtrain.get_label()
        p = 1 / (1 + np.exp(-z))          # sigmoid
        grad = b * (1 - y) * p - a * y * (1 - p)
        hess = p * (1 - p) * (a * y + b * (1 - y))
        return grad, hess
    return objective

dtrain = xgb.DMatrix(X_train, label=y_train)
booster = xgb.train(
    {"max_depth": 3, "eta": 0.1, "base_score": 0.5},
    dtrain, num_boost_round=100,
    obj=weighted_logloss(a=1, b=3.5))

z = booster.predict(xgb.DMatrix(X_test),
                    output_margin=True)
p = 1 / (1 + np.exp(-z))          # probability of class 1
y_pred = (p > 0.5).astype(int)
```

A `DMatrix` is XGBoost's own table format. `obj=` replaces XGBoost's built-in loss: at every round it calls our function with the current raw scores `z` and asks for the **gradient** (G-863; first derivative) and **Hessian** (G-887; second derivative) of each observation's loss. With a custom loss, XGBoost does not know how to turn raw scores into probabilities (XGBoost docs, Custom Objective), so we ask for the raw scores (`output_margin=True`) and apply the sigmoid ourselves.

With $b = 1$ (the ordinary log loss), XGBoost finds 2 of the 7 minority test observations. With $b = 3.5$ it finds 3 (recall **0.43**, precision 0.38). Section 10 repeats the comparison on real data, where the effect is clear.

Extra: For logistic regression, this weighted log loss is exactly what `class_weight={0: b, 1: a}` does: scikit-learn multiplies each observation's loss by the weight of its class (scikit-learn source, `_logistic.py`). A custom loss is needed when the cost is not one number per class, for example the focal loss (Lin et al. 2017), which down-weights observations the model already gets right, or in a library without a `class_weight` setting.

10. Comparing the techniques

Key point: On real data, every technique raises the minority recall far above the plain model's, at the cost of precision; which trade is best depends on the problem.

The toy test set has only 7 minority observations, so one observation found or missed moves the recall by $1/7 \approx 0.14$. Seven observations are enough for pictures, but too few to compare techniques. For the comparison we use real data, in conditions where the differences show:

- **Data:** the mammography dataset (Woods et al. 1993), the same data the SMOTE paper tests on (Chawla et al. 2002, §4.2). It has 11,183 observations and 6 features computed from mammogram images. The target says whether a region shows a calcification: only 260 observations (2.3%) do, so the minority class is class 1 here.
- **Testing:** 5-fold stratified cross-validation, repeated 3 times with different shuffles: 15 test folds of 52 minority observations each. Resampling happens inside each training fold (section 7.4). Each number below is the average over the 15 folds.

Technique	Precision (1)	Recall (1)	F1 (1)	ROC AUC
Logistic regression, as is	0.79	0.40	0.53	0.920
+ random undersampling	0.17	0.85	0.28	0.923
+ random oversampling	0.17	0.85	0.28	0.923
+ SMOTE	0.17	0.85	0.28	0.922
+ <code>class_weight="balanced"</code>	0.17	0.85	0.28	0.923
Random forest, as is	0.88	0.54	0.67	0.941
Balanced random forest	0.31	0.85	0.46	0.952
XGBoost, ordinary log loss	0.83	0.54	0.65	0.953
XGBoost, custom loss (minority mistakes cost 10)	0.56	0.77	0.65	0.953

Figure 12 draws the table’s first two columns. Three things stand out:




Figure 12: Minority precision against minority recall on the mammography data. Each arrow goes from a model as is to the same model with a fix: every fix moves right (more calcifications found) and down (more false alarms)

- **Recall rises sharply.** Logistic regression finds 40% of the calcifications as is and 85% after any of the four fixes. The forest goes from 54% to 85%, and XGBoost from 54% to 77%.
- **Precision falls.** The models now flag many more majority observations as class 1. A screening test works the same way: sending more patients for a second look catches more real cases, and also more healthy people. Whether that trade is worth it depends on the cost of each mistake.
- **ROC AUC barely moves.** ROC AUC does not depend on the threshold ([ROC Note](#)). Its stability says the models rank the observations almost as before; what changes is where they draw the line. For logistic regression, the four fixes land on the same point for this reason.

Extra: Copies against SMOTE for a decision tree. Section 6 warned that a tree overfits to copied observations. Chawla et al. (2002, §4.1 and §4.2) explain why: copies make the tree carve smaller and more specific regions around the minority points, while SMOTE’s new points make the regions larger and more general, so the tree generalises better. The same 15 folds show it:

Decision tree	Precision (1)	Recall (1)	F1 (1)	ROC AUC
as is	0.62	0.60	0.61	0.749
+ random oversampling (copies)	0.62	0.57	0.60	0.730
+ SMOTE	0.43	0.70	0.53	0.808

With copies the tree finds even fewer calcifications than before; with SMOTE it finds more, and its ROC AUC rises from 0.75 to 0.81.

10.1 Other techniques

Key point: imbalanced-learn offers many more resamplers and ensembles; its API reference lists them.

The techniques here are the most common. imbalanced-learn has many more, grouped the same way:

- **Undersampling:** TomekLinks, NearMiss, EditedNearestNeighbours, ClusterCentroids and others, which choose which majority observations to drop instead of dropping them at random.
- **Oversampling:** RandomOverSampler, SMOTE, SMOTENC, BorderlineSMOTE, SVMSMOTE, ADASYN.
- **Combined:** SMOTETomek, SMOTEENN (oversample, then clean noisy observations).
- **Ensembles:** BalancedRandomForestClassifier, BalancedBaggingClassifier, EasyEnsembleClassifier, RUSBoostClassifier.

11. Summary

Technique	What it does	Good	Bad
Random undersampling	drops random majority observations	less bias, faster training	loses data
Random oversampling	copies random minority observations	less bias, keeps all data	bigger data, overfitting
SMOTE	creates minority observations by interpolation	less bias, no copies	categorical features, slow, depends on k , outliers, unrealistic observations
Balanced random forest	balanced sample for every tree	uses most of the data	only for tree ensembles
Class weights	minority mistakes cost more	one setting, no new data	weight must be tuned
Custom loss	any cost we can write	full control	needs derivatives, only some libraries

- Imbalanced data biases a model towards the majority class, and accuracy hides it: judge by the minority’s precision, recall, F1 and ROC AUC.
- Many real problems (fraud, credit risk, rare disease, churn) are imbalanced, and the rare class is the one that matters.
- Resampling changes the data, balanced ensembles change each model’s sample, and cost-sensitive learning changes the loss.
- SMOTE creates $x + \lambda(n - x)$, a point between a minority observation and one of its k nearest minority neighbours.
- Resample only the training data, and inside each cross-validation fold.

12. Sources

Built from

- CampusX, “Imbalanced Data in Machine Learning | Undersampling | Oversampling | SMOTE”, YouTube, <https://www.youtube.com/watch?v=yh2AKoJCV3k>

Other references

- Chawla et al. 2002: N. V. Chawla, K. W. Bowyer, L. O. Hall and W. P. Kegelmeyer, *SMOTE: Synthetic Minority Over-sampling Technique*, Journal of Artificial Intelligence Research 16, 2002, 321–357, §4.1 and §4.2 (copies against SMOTE on the mammography data).
- imbalanced-learn user guide: imbalanced-learn 0.14 documentation, sections “Over-sampling” and “Common pitfalls and recommended practices” (data leakage).
- imbalanced-learn API reference: `imblearn.ensemble`, imbalanced-learn 0.14.
- imbalanced-learn source, `_smote/base.py`: `imblearn/over_sampling/_smote/base.py`, `new point = X[rows] + steps * diffs` with `steps` uniform in $[0, 1]$.

- XGBoost docs, Custom Objective: XGBoost documentation, “Custom Objective and Evaluation Metric” (“XGBoost doesn’t know its link function”).
- scikit-learn source, `_logistic.py`: `sklearn/linear_model/_logistic.py`, scikit-learn 1.9 (class weights multiply the sample weights).
- Woods et al. 1993: K. S. Woods, C. C. Doss, K. W. Bowyer, J. L. Solka, C. E. Priebe and W. P. Kegelmeyer, *Comparative evaluation of pattern recognition techniques for detection of microcalcifications in mammography*, International Journal of Pattern Recognition and Artificial Intelligence 7(6), 1993, 1417–1436. The copy used here comes from imbalanced-learn’s `fetch_datasets`.
- Lin et al. 2017: T.-Y. Lin, P. Goyal, R. Girshick, K. He and P. Dollár, *Focal Loss for Dense Object Detection*, ICCV 2017.

13. Key terms

Term	Meaning
Observation	One record, one row of the data table
Feature	An input variable, one column of the data table
Target	The output we predict, here the class
Majority class	The class with the most observations in imbalanced data
Minority class	The rare class in imbalanced data, usually the one we need to find
Resampling	Changing the number of observations per class to balance the data
Random undersampling	Dropping randomly chosen majority observations until the classes are equal
Random oversampling	Copying randomly chosen minority observations until the classes are equal
SMOTE	Synthetic Minority Over-sampling Technique: new minority observations by interpolation between minority neighbours
Interpolation	Placing a new point on the segment between two existing points
Synthetic data	Rows created by an algorithm rather than collected
Balanced random forest	A random forest in which every tree is trained on a balanced sample
Cost-sensitive learning	Changing the learning so that mistakes on some classes cost more
Custom loss function	A loss written by the user, passed to libraries such as XGBoost with its gradient and Hessian

Term	Meaning
imbalanced-learn	A Python library (<code>imblearn</code>) of resampling techniques and balanced ensembles, with a <code>fit_resample</code> method