Maps and Landmarks

Contents

1. Overview . . . 2
2. Why a robot needs a map . . . 3
   - 2.1 What a map is for . . . 3
3. Feature-based and grid maps . . . 3
4. Obstacles as polygons made of half-planes . . . 5
   - 4.1 Why store obstacles exactly . . . 5
   - 4.2 Half-plane: one side of a line . . . 5
   - 4.3 A convex polygon is the overlap of half-planes . . . 6
   - 4.4 Non-convex obstacles are unions of convex pieces . . . 8
5. Triangle meshes and bitmaps . . . 9
   - 5.1 Triangle meshes . . . 9
   - 5.2 Bitmaps and occupancy grids . . . 10
6. Landmarks and feature extraction . . . 13
   - 6.1 Why extract features . . . 13
   - 6.2 Finding the poles in a scan . . . 13
   - 6.3 Range, bearing and signature . . . 14
7. Summary . . . 17
8. Sources . . . 18
9. Key terms . . . 18

Where this fits:

- **Builds on:** Equation of a hyperplane; Convex sets and convex optimisation; Pose.
- **Leads to:** Beam model; Landmark measurement model; Graphs and state spaces; Obstacle region and free space.

1. Overview

Key point: A robot finds where it is by comparing what it senses with a **map** (G-2331). A map can list a few landmarks or mark every cell as free or occupied; obstacles can be stored exactly as polygons built from half-planes, as triangles, or as a grid of cells; and landmarks are pulled out of raw sensor data as range, bearing and signature.

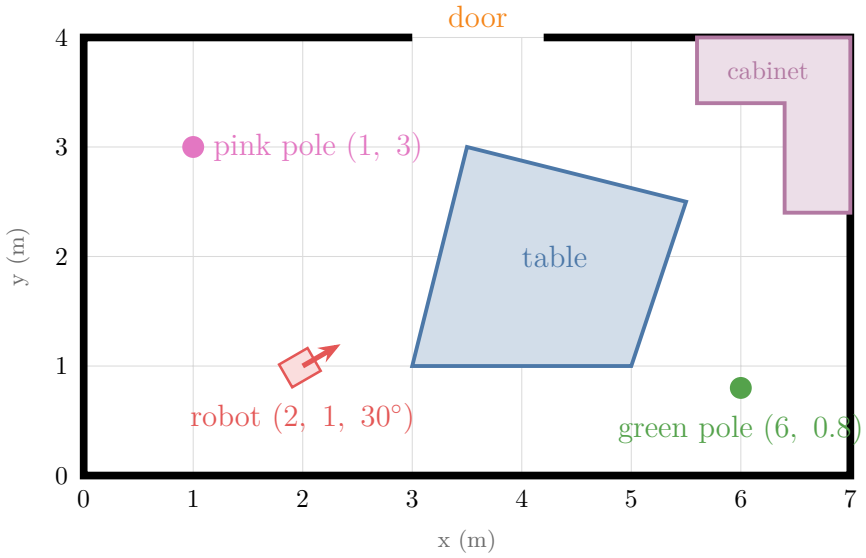


Figure 1: A room 7 m by 4 m with a door in the top wall, a four-sided table, an L-shaped cabinet in the top right corner, a pink pole at (1, 3) and a green pole at (6, 0.8); the robot stands at (2, 1) facing 30 degrees

So far the robot can say how a move widens the distribution of its pose: the **motion model** (G-2319) turns a drive from (2 m, 1 m, 30°) into a cloud of end **poses** (G-2280), positions with headings. The remaining problem is that a reading can shrink that cloud only if the robot knows where the sensed object stands. "The pink pole is 2.236 m away" means nothing until a map says the pink pole is at (1, 3). A robot compares each reading with a map: "door" fits only the places where the map has a door. This Note adds the map, for the same robot at (2 m, 1 m) facing 30°, now placed in the room of Figure 1 so that there is something to map:

- what a map is, and the main families of maps: feature-based and grid (Sections 2 and 3);
- how to store an obstacle exactly, as a polygon made of **half-planes** (G-2336), and test whether a point is inside it (Section 4);

- two other common models, **triangle meshes** (G-2338) and bitmaps, and what each costs (Section 5);
- how a robot pulls **landmarks** (G-2342) out of a laser scan and describes each by its range, bearing and signature (Section 6).

2. Why a robot needs a map

> **Key point:** A **map** (G-2331) is the robot's model of its surroundings: a list of the objects that matter, with where they are. Localization compares readings with it, and planning searches it for routes.

2.1 What a map is for

A reading such as "the pink pole is 2.236 m away, to the left" says where things are relative to the robot. To turn it into "where am I?", the robot needs to know where those things are in the room. A map gives exactly that: the robot compares what it sees with the map and keeps the poses where the two agree. Planning needs the map too: a route to the door must go round the table.

A **map** (G-2331) is a list of the objects in the environment that matter to the robot, with their positions and properties (Thrun et al. 2005 §6.2). We write it as a list of objects:

- m_i is object number i of the map, such as the pink pole at (1 m, 3 m);
- N is the number of objects, 2 when the map holds only the two poles (Section 3);
- m is the whole map, the list of all N objects.

$$m = \{m_1,\ m_2,\ \ldots,\ m_N\}$$

Fixed things, such as walls and furniture, belong in the map; things that move, such as people, belong in the robot's state (G-2314), the quantities it keeps track of, instead. What goes in the map is settled; how much of the room it records is not.

3. Feature-based and grid maps

> **Key point:** A feature-based map lists only landmarks and where they are; a grid **map** (G-2331) gives every cell of space a value, free or occupied. The first is small but says nothing about free space; the second is large but answers "can the robot be here?" at once.

What should the map of Figure 1's room hold: only the two poles, or every spot of floor? Maps come in two kinds (Thrun et al. 2005 §6.2). Figure 2 draws our room both ways.

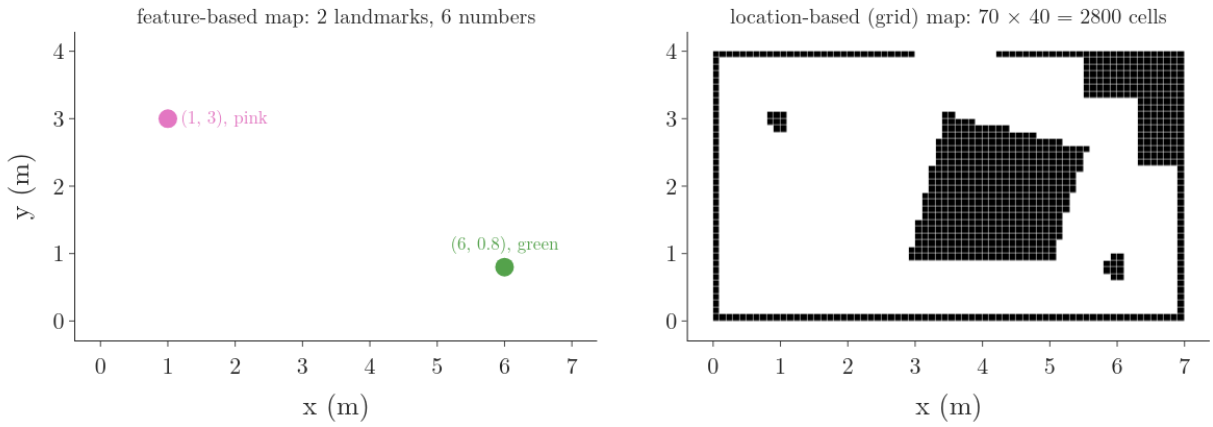


Figure 2: The room as a feature-based map, only the pink pole at (1, 3) and the green pole at (6, 0.8): 6 numbers; and as a grid map of 0.1 m cells, black where the walls, table, cabinet and poles are: 70 by 40 = 2800 cells

A **feature-based map** (G-2334) is a list of **landmarks** (G-2342), distinctive objects such as the two poles, each with its coordinates and a label that tells landmarks apart. Our room needs two entries, each a pole's x and y in metres and its colour. The map m, the list of the room's objects, is then:

$$m = \{(1,\ 3,\ \text{pink}),\ (6,\ 0.8,\ \text{green})\}$$

The map holds 6 numbers for the whole room (Brown CS148 tutorial). A feature-based map is cheap to store and to search, but it says nothing about the rest of the room: from it the robot cannot tell that the table is there.

A **location-based map** (G-2335), or grid map, splits space into small cells and stores a value for each: occupied or free. With cells of 0.1 m our room takes:

$$\frac{7}{0.1} \times \frac{4}{0.1} = 70 \times 40$$

$$= 2800 \text{ cells}$$

A grid map is called location-based, or volumetric, because it has an entry for every location, free space included. Those free-space entries let it answer "is this pose free?", the test that the map-consistent motion model (G-2329) uses to keep poses out of walls.

	Feature-based map	Grid map
Stores	landmarks: position and label	every cell: occupied or free
Our room	6 numbers	2800 cells
Free space	unknown	known
Matches	landmark readings	full scans
Weak point	needs landmarks to exist and be found	memory grows with area and fineness

Grid maps are the most common map in robotics today (Correll §4.1). Cells need not all be the same size: a **quadtree** (G-2533) splits a cell into four only where it contains an obstacle edge, so large free areas take few cells (Correll §4.1). Either kind of grid stores each obstacle only to its nearest cell.

4. Obstacles as polygons made of half-planes

> **Key point:** A line splits the floor into two half-planes, told apart by the sign of one formula. A convex obstacle is the overlap of a few half-planes, so a point is inside exactly when every formula is at most 0; any other obstacle is a **union** (G-2045) of convex pieces, the points in any one of them.

4.1 Why store obstacles exactly

The grid **map** (G-2331) of Section 3 says whether the robot can stand somewhere in the room, but it blurs every edge to the nearest cell. Collision checks in route planning often need the exact shape: does the robot's path clip the table's corner or not? The table in Figure 1 is a four-sided polygon, given by its corners in counter-clockwise order (LaValle §3.1.1):

$$(3,\ 1),\ (5,\ 1),\ (5.5,\ 2.5),\ (3.5,\ 3)$$

The list draws the table's outline, but it does not answer the question a planner asks repeatedly: is the point (4, 2) inside the table? **Half-planes** (G-2336) answer it with a few multiplications.

4.2 Half-plane: one side of a line

Every edge of the table lies on a line, which we can write in the general form of a line. For a point on the floor:

- x is the point's x coordinate, in m;
- y is its y coordinate, in m;
- a is the number that multiplies x, set by the line's direction (1.5 for the table's right edge below);
- b is the number that multiplies y (-0.5 for the right edge);

- c is the constant that shifts the line into place (-7 for the right edge);
- $f(x, y)$ is the line's function: 0 on the line, negative on one side and positive on the other.

$$f(x, y) = ax + by + c = 0$$

The vector (a, b) is the line's normal vector (G-1346), a direction square to the line, and points towards the positive side. All the points on one side of a line, together with the line, form a **half-plane** (G-2336) (LaValle §3.1.1):

- H is the half-plane on the negative side: the set of floor points (x, y) where f is 0 or less; in the braces below, the colon reads "such that".

$$H = \{(x, y) : f(x, y) \leq 0\}$$

Worked example: the right edge of the table. It runs from (5, 1) to (5.5, 2.5), so its direction is (0.5, 1.5). The corners run counter-clockwise, so the table always lies on the left of each edge as we walk along it. Turning the direction a quarter turn clockwise therefore gives a normal that points to the right of the edge, out of the table:

$$(a, b) = (1.5,\ -0.5)$$

The line passes through (5, 1), where f must be 0:

$$1.5 \times 5 - 0.5 \times 1 + c = 0$$

$$c = -7$$

Doubling every number gives whole numbers and the same line, with the same signs. We call this right-edge function f_2:

$$f_2(x, y) = 3x - y - 14$$

One edge gives one half-plane; the table has four edges.

4.3 A convex polygon is the overlap of half-planes

The table is a convex set (G-479): the segment between any two of its points stays inside it. A convex polygon is the overlap, the intersection (G-967), of one **half-plane** (G-2336) per edge, each taking the side that holds the polygon (LaValle §3.1.1). Our table's four edge functions of a floor point (x, y), in m, each scaled to whole numbers and signed so that the inside is negative:

- f_1 is the bottom edge's function, for the edge from (3, 1) to (5, 1);
- f_2 is the right edge's function, from (5, 1) to (5.5, 2.5) (Section 4.2);
- f_3 is the top edge's function, from (5.5, 2.5) to (3.5, 3);
- f_4 is the left edge's function, from (3.5, 3) to (3, 1).

$$f_1 = 1 - y$$

$$f_2 = 3x - y - 14$$

$$f_3 = 2x + 8y - 31$$

$$f_4 = 11 - 4x + y$$

Figure 3 adds the four half-planes one at a time; what is left inside all of them is the table.

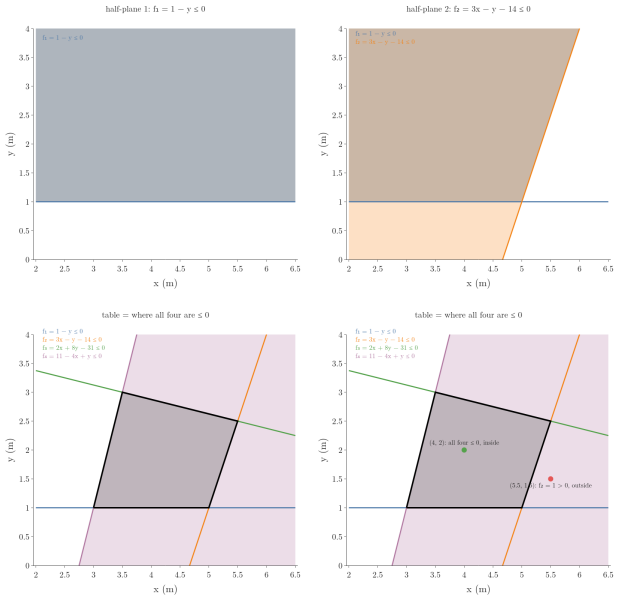


Figure 3: Animation: the half-planes f1 = 1 − y ≤ 0, f2 = 3x − y − 14 ≤ 0, f3 = 2x + 8y − 31 ≤ 0 and f4 = 11 − 4x + y ≤ 0 are shaded one at a time; the region inside all four is the table. The point (4, 2) gives all four values at most 0 and is inside; the point (5.5, 1.5) gives f2 = 1 and is outside

Inside test. A point is inside the table when all four values are at most 0. For the point (4, 2):

$$f_1 = 1 - 2 = -1$$

$$f_2 = 12 - 2 - 14 = -4$$

$$f_3 = 8 + 16 - 31 = -7$$

$$f_4 = 11 - 16 + 2 = -3$$

All four are negative: inside. For the point (5.5, 1.5):

$$f_2 = 16.5 - 1.5 - 14 = 1$$

One positive value is enough: outside, on the far side of the right edge. In logic, the table's test joins the four conditions with AND (LaValle §3.1.1):

$$f_1 \leq 0 \text{ AND } f_2 \leq 0$$

$$\text{AND } f_3 \leq 0 \text{ AND } f_4 \leq 0$$

4.4 Non-convex obstacles are unions of convex pieces

The cabinet is L-shaped, so it is not convex: the segment between a point in one arm and a point in the other leaves the cabinet (Figure 4). No set of **half-planes** (G-2336) taken together with AND can describe it, because their overlap is always convex. Instead we cut it into convex pieces, two rectangles, and take their union (G-2045) (LaValle §3.1.1). For a floor point (x, y), in m:

- R_1 is the cabinet's top arm, a rectangle along the top wall;
- R_2 is its right arm, a rectangle along the right wall;
- $\cup$ joins them into their union: the points in either piece.

$$R_1 : \ 5.6 \leq x \leq 7, \quad 3.4 \leq y \leq 4$$

$$R_2 : \ 6.4 \leq x \leq 7, \quad 2.4 \leq y \leq 4$$

$$\text{cabinet} = R_1 \cup R_2$$

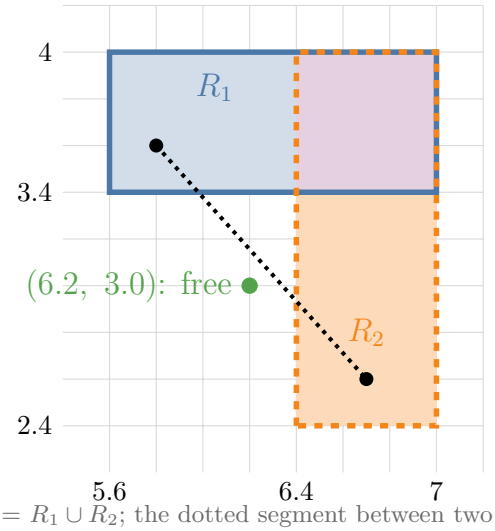


Figure 4: The L-shaped cabinet as the union of two rectangles R1 and R2, which overlap in the corner; the point (6.2, 3.0) in the notch of the L is in neither, so it is free; a dotted segment between a point of R1 and a point of R2 leaves the cabinet, so the cabinet is not convex

A point is in the cabinet if it is in R_1 OR in R_2. The point (6.2, 3.0), in the notch of the L, fails both:

- in R_1? No: $y = 3.0$ is below 3.4;
- in R_2? No: $x = 6.2$ is left of 6.4.

So it is free. Every obstacle made of straight edges can be written this way, as a union of **intersections** (G-967) of half-planes; the pieces may overlap, as R_1 and R_2 do in the corner (LaValle §3.1.1).

5. Triangle meshes and bitmaps

> **Key point:** In 3D, obstacles are often stored as many triangles, which graphics hardware draws fast but which do not say what is inside. A **bitmap** (G-2339) marks every cell that touches an obstacle; it is simple and fits sensor data, but finer cells cost many more cells.

5.1 Triangle meshes

Half-planes (G-2336), the regions on one side of a line, describe the table's outline on the floor plan, but the room's table and cabinet are really 3D, and the common 3D model is a set of triangles, each given by its three corner points in space. A set of triangles that together cover an object's surface is a **triangle mesh** (G-2338) (LaValle §3.1.3). Figure 5 builds our

table, 0.75 m high, from 8 corner points and 12 triangles: 2 for the top, 2 for the bottom and 2 for each of the 4 sides.

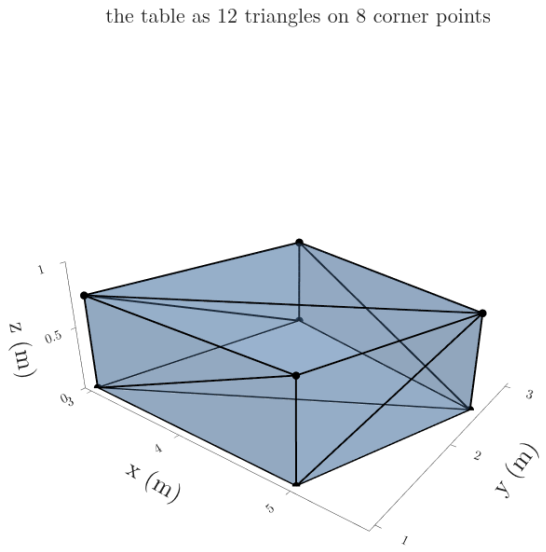


Figure 5: The table as a 3D triangle mesh: 8 corner points, at floor height and at 0.75 m, joined into 12 triangles covering the top, the bottom and the four sides

Triangles are popular because graphics hardware is built to draw them (LaValle §3.1.3). Neighbouring triangles share corners, so meshes often store each corner once and list triangles by corner numbers, as the table's 12 triangles share 8 corners. The drawback: a mesh is only a surface. Nothing in a loose set of triangles says which side is inside, and meshes made from scans or design files often have holes, so a planner cannot always tell whether a point is inside an object (LaValle §3.1.3).

5.2 Bitmaps and occupancy grids

The simplest model of all splits the room into square cells and marks each one: 1 if the cell contains any point of an obstacle, 0 if it does not. The result is a **bitmap** (G-2339) of the obstacles, the same thing as the black-and-white image of Figure 6 (LaValle §3.1.3).

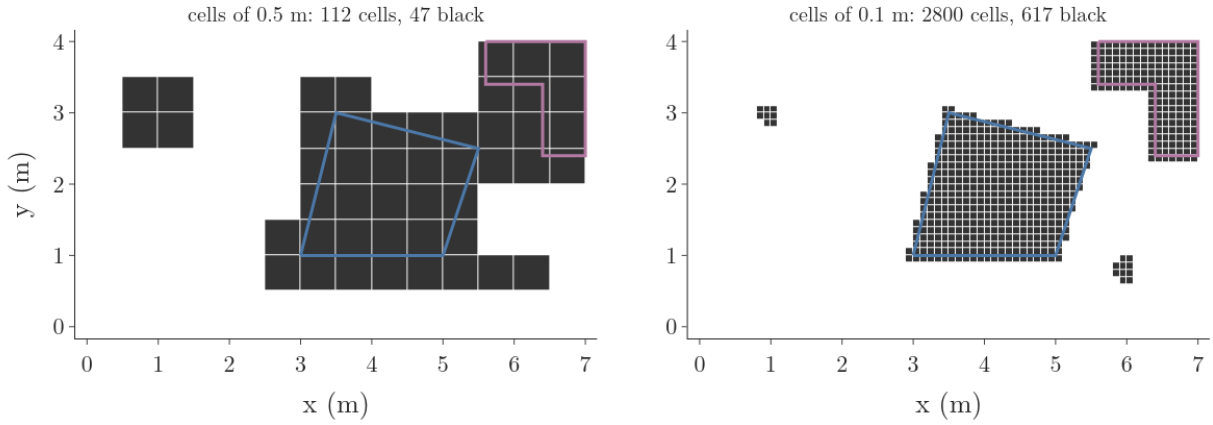


Figure 6: Bitmaps of the room's furniture: with 0.5 m cells, 112 cells of which 47 are black, a coarse blocky outline; with 0.1 m cells, 2800 cells of which 617 are black, close to the true table and cabinet, which are drawn as outlines

To read Figure 6: each small square is one cell, black for 1 (some obstacle inside) and white for 0; the coloured outlines are the exact polygons of Section 4. The cell size sets the trade-off. Black area overstates the obstacles, because a cell turns black when an obstacle covers any part of it. The table, cabinet and poles really cover 5.13 m^2:

Cell size	Cells	Black cells	Black area
0.5 m	112	47	11.75 m^2
0.1 m	2800	617	6.17 m^2

Cells five times smaller bring the black area close to the true 5.13 m^2 but need 25 times as many cells. Bitmaps arise naturally when a robot builds a **map** (G-2331), its model of the surroundings, with its own sensors (LaValle §3.1.3).

A bitmap cell holds only 0 or 1, so it cannot say how much of the cell an obstacle takes: with 0.5 m cells, a cell that the table barely clips is as black as a cell deep inside it. Storing a number between 0 and 1 in each cell instead of 0 or 1, such as the probability that an obstacle is in the cell, gives an **occupancy grid** (G-2340), the standard grid map of robotics (LaValle §3.1.3). No sensor has measured our room yet, so here we set each cell's number to the share of its area that the table, cabinet and poles cover. The share is the probability that a point picked at random in the cell, every spot equally likely, lies inside an obstacle:

- a cell fully inside the table or the cabinet gets 1;
- a free cell gets 0;
- a cell on an edge gets the share covered, such as 0.50 for the cell from 5.0 to 5.5 m along x and 1.5 to 2.0 m along y, which the table's right edge cuts in half.

Figure 7 puts the 0.5 m bitmap of Figure 6 beside the occupancy grid of the same cells, in the same greys: white is 0, black is 1, and each in-between cell carries its number. Look

along the edges of the table and the cabinet, and at the poles: where the bitmap marks a cell black, the occupancy grid shades it by how much of it is covered, so the cells round the poles drop to 0.03 and 0.06.

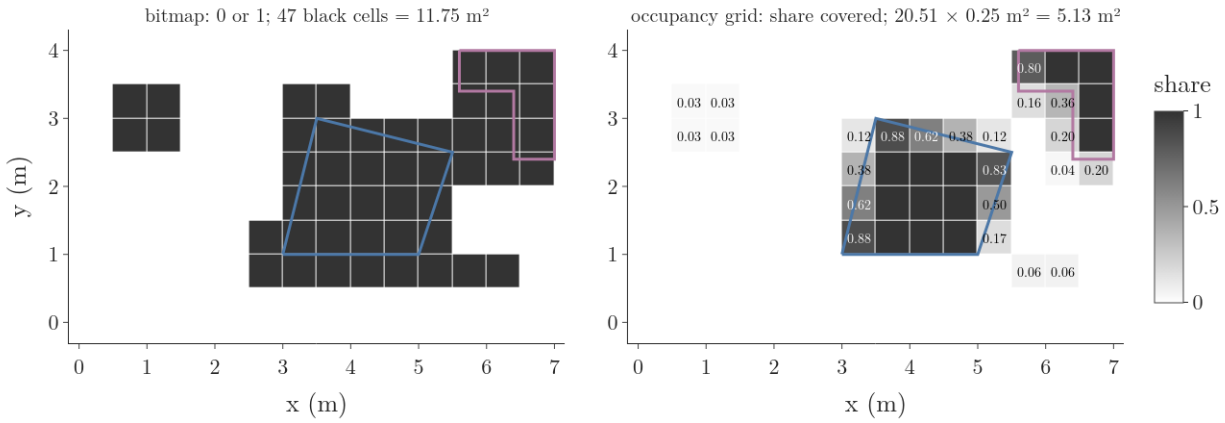


Figure 7: Two maps of the room's 0.5 m cells in the same greys, the table and cabinet drawn as outlines. Left, the bitmap: 47 black cells, 11.75 m^2. Right, the occupancy grid: each cell shaded by the share of it that the furniture and poles cover, 1 inside the table and cabinet, 0 in free space, numbers such as 0.12, 0.50 and 0.88 in the edge cells and 0.03 and 0.06 round the poles; the shares add up to 20.51, which times the 0.25 m^2 of one cell is 5.13 m^2

Over the 112 cells of Figure 7:

- 13 cells are fully covered and hold 1;
- 23 cells are partly covered and hold a number in between;
- 76 cells are free and hold 0;
- 11 of the bitmap's 47 black cells only touch an obstacle along their border, so none of their area is covered and their share is 0.

A cell's share times its area is the obstacle area inside that cell, so adding the shares gives back the true covered area:

- the shares of the 112 cells add up to 20.51;
- each cell is 0.5 m by 0.5 m, an area of 0.25 m^2.

$$\text{covered area} = 20.51 \times 0.25 \text{ m}^2$$

$$= 5.13 \text{ m}^2$$

The bitmap of the same cells gives 11.75 m^2. Each of these models describes the room; the robot still has to turn its own readings into something it can compare with them.

6. Landmarks and feature extraction

> **Key point:** Most of the 360 distances in a scan are readings of wall. **Feature extraction** (G-2341) finds the few objects in it that are easy to spot and tell apart, the landmarks, and reports each as a range, a bearing and a signature.

6.1 Why extract features

The other half of the comparison is the reading itself. A laser scanner on our robot at (2, 1) measures 360 distances, one per degree. Comparing all 360 with a **map** (G-2331), the robot's model of the room, is possible (the range-sensor Notes later in this chapter do it), but most of the numbers describe plain wall, which looks the same all along the room. A few distinctive objects carry most of the information about where the robot is. Pulling such objects out of raw sensor data, so that a scan of hundreds of numbers becomes a short list of objects, is **feature extraction** (G-2341) (Thrun et al. 2005 §6.6, as in the Freiburg sensor-model slides).

An object makes a good **landmark** (G-2342) when two things hold:

1. the robot can **detect** it reliably, such as a pole that stands out in a laser scan or a coloured marker in a camera image;
2. the robot can **tell it apart** from the other landmarks: tree 1 from tree 2, the pink pole from the green one.

Landmarks can be natural objects (trees, doors, road signs) or artificial ones put there for the robot (coloured markers, reflectors). Some are passive, only seen; others are active and send a signal, such as a radio beacon (Freiburg sensor-model slides). The two poles of our room are passive landmarks; the next step finds them in the scan.

6.2 Finding the poles in a scan

Our robot at (2, 1) facing 30° takes a simulated scan: 360 beams, one per degree, each with 1 cm of range noise. The beams hit the walls, the near side of the table, and two small clusters.

A pole shows up as a few neighbouring beams that are much shorter than the beams on both sides, which pass it and hit the wall behind. So the extraction rule, applied to the scan in order, is:

1. cut the scan wherever the range jumps by more than 0.3 m between neighbouring beams;
2. keep a piece that is narrower than 0.35 m and nearer than the beams on both sides of it;
3. report its **bearing** (G-2344), its direction measured from the robot's heading, as the average direction of its beams, and its **range** (G-2343), its distance, as its nearest beam's range plus the pole's radius of 0.1 m, so that the range is to the pole's centre.

Figure 8 applies the rule to the scan, one beam after another. Watch each reading land as a grey dot, the orange jump at each cut, and the test of each piece when the next cut closes it: only a 3-beam piece and a 5-beam piece pass, and they become the green and the pink pole.

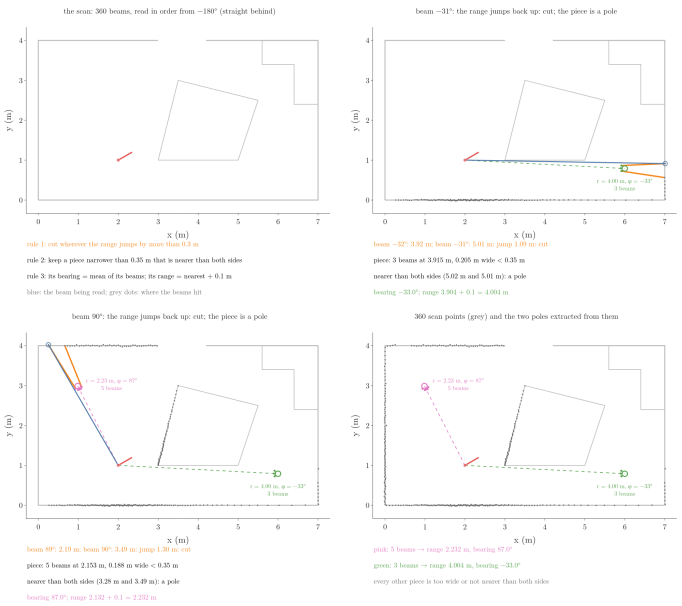


Figure 8: Animation: from the robot at (2, 1) facing 30 degrees, a blue cursor reads the 360 beams in order and each reading lands as a grey dot; at each cut, where the range jumps by more than 0.3 m, an orange line shows the jump, and the piece that the cut closes is tested: narrower than 0.35 m and nearer than the beams on both sides. Beams through the door read 8 m and leave the picture. Only two pieces pass: 3 beams, the green pole at range 4.004 m and bearing −33.0 degrees, and 5 beams, the pink pole at range 2.232 m and bearing 87.0 degrees. The last frame shows all 360 scan points in grey and the two extracted poles, each labelled with its range r and bearing ϕ

The rule finds exactly the two poles:

Pole	Beams	Extracted range	True range	Extracted bearing	True bearing
pink	5	2.232 m	2.236 m	87.0°	86.6°
green	3	4.004 m	4.005 m	−33.0°	−32.9°

The extracted **bearings** (G-2344) are within half a degree of the true ones, because the beams are one degree apart. The table's corners pass step 2 only if they are narrow and stand in front of something farther away, which they do not here; real extractors add more tests for such cases.

6.3 Range, bearing and signature

Each extracted **landmark** (G-2342) is reported as three numbers (Thrun et al. 2005 §6.6, as in the Brown CS148 tutorial), shown in Figure 9:

- its **range** (G-2343) r: the distance from the robot to the landmark, in m (2.236 m for the pink pole, worked out below);

- its **bearing** (G-2344) φ: the angle from the robot's heading to the landmark, positive to the left, negative to the right, in degrees (86.6° for the pink pole);
- its **signature** (G-2345) s: what kind of landmark it is, such as its colour (pink) or an ID number.

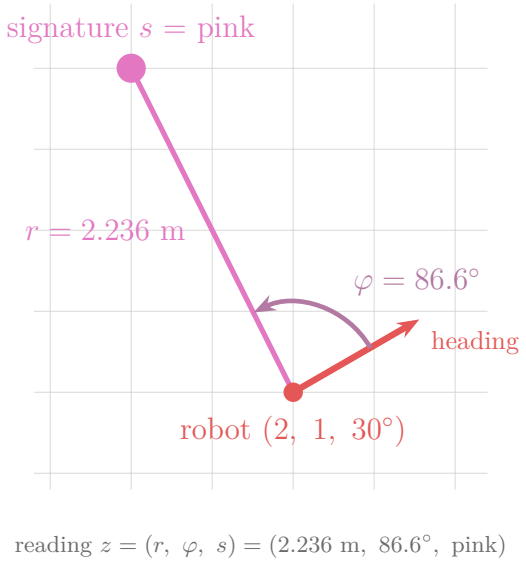


Figure 9: A landmark reading: from the robot at (2, 1) facing 30 degrees, the pink pole at (1, 3) is at range r = 2.236 m and bearing phi = 86.6 degrees from the heading; its signature is pink

For the pink pole, the true values follow from the **map** (G-2331). The pole is 1 m back in x and 2 m up in y from the robot:

$$r = \sqrt{1^2 + 2^2} = 2.236 \text{ m}$$

The pole's direction in the room comes from atan2 (G-2325), which turns the rise and the run into an angle over the full circle:

$$\text{atan2}(2,\ -1) = 116.6°$$

The bearing is measured from the heading, 30°:

$$\varphi = 116.6° - 30° = 86.6°$$

So the reading of the Overview, with its bearing and signature, is one landmark reading:

- z is the reading: its range, bearing and signature together.

$$z = (2.236 \text{ m},\ 86.6°,\ \text{pink})$$

Why each of the three numbers is needed:

- **Range and bearing** fix where the landmark is relative to the robot. Range alone puts it anywhere on a circle around the robot; bearing alone, anywhere along a ray from it.
- **The signature** says which landmark of the map the reading belongs to. Without it, two identical poles would leave the robot with two explanations of every reading, like the three identical doors of the hallway; with it, the reading "pink" can only be the pole at (1, 3).

In localization the map is known and the pose is not. Then a single landmark reading cannot fix the pose; with the range alone it only says the robot stands somewhere on a circle of radius r around the landmark, and two or three landmarks together fix it. The **landmark measurement model** (G-2401) builds that step.

7. Summary

Model	Stores	Inside test	Strength	Weakness
Feature-based map	landmarks: position, signature	none	tiny (6 numbers here)	no free space
Grid map	a value per cell	read the cell	free space known; fits scans	size grows with area and fineness
Polygons from half-planes	edge functions	every edge function f_i is at most 0 in some convex piece	exact straight edges	curved shapes need many edges
Triangle mesh	3D triangles	not defined by the mesh	fast to draw in 3D	inside and outside unclear
Bitmap / occupancy grid	0 or 1 (or a probability) per cell	read the cell	simple; built from sensors	overstates obstacles; finer costs more

- The reading "the pink pole is 2.236 m away" places the robot only once the **map** (G-2331) says the pink pole is at (1, 3): a map is the robot's list of the objects that matter, because readings say where things are relative to the robot and only a map turns them into a pose.
- A feature-based map stores 6 numbers for our room and a 0.1 m grid 2800 cells; the grid costs more but records where free space is, which the **map-consistent motion model** (G-2329) needs to keep sampled poses out of walls.
- A **half-plane** (G-2336) is one sign test: the function f of the line through an edge is at most 0; a convex obstacle is where all its edge tests pass, so the table's inside test is four multiplications and sums (-1, -4, -7, -3 for (4, 2)).
- Non-convex obstacles are **unions** (G-2045) of convex pieces, because the overlap of half-planes is always convex; the cabinet's notch point (6.2, 3.0) fails both rectangles and is free.
- **Triangle meshes** (G-2338) suit 3D graphics, because graphics hardware is built to draw triangles, but do not define an inside; bitmaps are simple, but 0.5 m cells inflate the obstacles from 5.13 m^2 to 11.75 m^2, because a cell turns black when any part of it is covered, and 0.1 m cells cost 25 times as many cells.
- An **occupancy grid** (G-2340) stores a number between 0 and 1 in each cell, such as the probability that an obstacle is there; with each 0.5 m cell holding the share of it that is covered, the shares add up to the true 5.13 m^2, because each share times the cell's area is the obstacle area inside it.
- **Feature extraction** (G-2341) turns 360 scan distances into two landmarks, because a few distinctive objects carry most of the information about the pose; each landmark is reported as range, bearing and signature, and the signature says which map landmark it is.

The robot can now hold its room as a map and turn a scan into readings such as (2.236 m, 86.6°, pink). Still missing: how likely is a reading if the robot stood at a guessed pose? The **beam model** (G-2377) answers that for each beam of a laser scan.

8. Sources

Built from

- Oleg Shipitko, "Lecture 5. Mapping", YouTube, 1:06–23:00, https://www.youtube.com/watch?v=2_xC (Shipitko, Mapping)
- Steven M. LaValle, "Motion Planning: Life in C-Space" (ICRA tutorial), YouTube, 3:35–8:16, https://www.youtube.com/watch?v=ERaZJru80bM (LaValle, Life in C-space)
- Cyrill Stachniss, "Observation Models", YouTube, 32:42–36:17, https://www.youtube.com/watch?v=Sf o (Stachniss, Observation models)
- Carlotta A. Berry, "Advanced Mobile Robotics: Lecture 4-1a", YouTube, 4:07–5:38, https://www.youtube.com/watch?v=zNanMwnBT5w (Berry, Lecture 4-1a)
- LaValle, S. M. (2006). *Planning Algorithms*. Cambridge University Press. §3.1.1 polygonal models (convex polygons as intersections of half-planes, non-convex polygons as unions, the logical predicate), §3.1.3 other models (3D triangles, bitmaps, occupancy grids). Free online: https://lavalle.pl/planning/node78.html (LaValle)
- Schwertfeger, J. (2007). "Tutorial on a Probabilistic Measurement Model based on Landmark Range and Bearing Information", Brown University CS148. https://cs.brown.edu/courses/cs148/tutorials/measurement_model_tutorial.pdf (Brown CS148 tutorial)

Other references

- Correll, N. et al. *Introduction to Autonomous Robots*, §4.1 "Map Representations". Free online: https://eng.libretexts.org/Bookshelves/Mechanical_Engineering/Introduction_to_Autonomou (Correll §4.1)
- Burgard, W. et al. (2023). "Probabilistic Sensor Models", *Introduction to Mobile Robotics* slides, University of Freiburg (features and landmarks). http://ais.informatik.uni-freiburg.de/teaching/ss23/robotics/slides/07-sensor-models.pdf (Freiburg sensor-model slides)
- Thrun, S., Burgard, W. and Fox, D. (2005). *Probabilistic Robotics*. MIT Press. §6.2 maps, §6.6 feature-based measurement models (range, bearing, signature). Cited only where the free sources above confirm it.

9. Key terms

Terms taught in this Note come first; linked terms are recaps, taught in the Note the link opens.

Term	Meaning
Map (G-2331)	A robot's list of the objects in its surroundings that matter, with their positions and properties, such as the pink pole at (1, 3); readings only give positions relative to the robot, and the map turns them into a pose.
Half-plane (G-2336)	All points on one side of a line, together with the line: the points where $ax + by + c \leq 0$; one sign test, used as a building block for polygon obstacles.
Triangle mesh (G-2338)	A set of triangles, each given by three corner points in space, that together cover an object's surface (a table as 12 triangles on 8 corners); fast to draw in 3D, but it does not say what is inside.
Landmark (G-2342)	An object the robot can detect reliably and tell apart from the others, such as a coloured pole or a door; its known position in the map lets readings of it fix the robot's pose.
Feature-based map (G-2334)	A map that stores only landmarks, each with its coordinates and signature, such as (1, 3, pink); very compact, but it says nothing about which space is free.
Location-based map (grid map) (G-2335)	A map that splits space into cells and stores a value for every cell, occupied or free, free space included (2800 cells for a 7 m by 4 m room at 0.1 m); larger, but it answers at once whether a pose is free.
Quadtree (G-2533)	A grid map whose cells are split into four smaller cells only where they contain an obstacle edge, so large free areas stay as a few big cells and the map needs far fewer cells than a uniform fine grid.
Bitmap (obstacle bitmap) (G-2339)	A grid of cells marked 1 if the cell contains any obstacle point and 0 otherwise, like a black-and-white image; simple and natural for sensor data, but it overstates obstacles, and finer cells cost many more cells.
Occupancy grid (G-2340)	A grid map that stores in each cell a number between 0 and 1, such as the probability that the cell is occupied, instead of a hard 0 or 1; the standard grid map of robotics, since sensor readings are uncertain.

Term	Meaning
Feature extraction (from sensor data) (G-2341)	Pulling a few distinctive objects out of raw sensor data, such as two poles out of a 360-beam laser scan; it shrinks hundreds of numbers to a short list that is easy to match with a map.
Bearing (G-2344)	The angle from the robot's heading to a sensed landmark, positive to the left, such as 86.6 degrees; together with the range it fixes where the landmark is relative to the robot.
Range (of a landmark reading) (G-2343)	The distance from the robot to a sensed landmark, such as 2.236 m; together with the bearing it fixes where the landmark is relative to the robot.
Signature (of a landmark) (G-2345)	A label that says what kind of landmark was sensed, such as its colour or ID number; it tells the robot which landmark in the map a reading belongs to.
Motion model $p(\mathbf{x}_t \mid u_t, \mathbf{x}_{t-1})$ (G-2319)	The probability distribution of a robot's pose after a move, given its pose before and the control (command or odometry); it describes how uncertain motion is, so a filter can score possible end poses or draw them at random.
Pose (G-2280)	Where a robot is and which way it faces: on a floor, its position (x, y) and heading θ, such as (2 m, 1 m, 30°); needed because a robot at one spot drives off differently depending on its heading.
State x_t (G-2314)	The quantities a robot keeps track of at time t because they decide what happens next, such as its pose, its speeds and moving things around it; it is what localization and filtering estimate.
Map-consistent motion model (G-2329)	A motion model that also uses the map, approximated for small steps by multiplying the plain motion model by 1 for free poses and 0 for poses inside obstacles; it stops the robot's estimate from passing through walls.
Union (A ∪ B) (G-2045)	The set of elements in A or B (or both), such as the points inside either of two convex pieces of an obstacle; for events, the event that A or B (or both) happens.
Normal vector (G-1346)	A vector perpendicular to a line, plane or hyperplane; for $w^{\mathsf{T}}x + w_0 = 0$ it is w.

Term	Meaning
Convex set (G-479)	A set that contains the whole segment between any two of its points.
Intersection (A ∩ B) (G-967)	The set of elements that are in both A and B, such as the points inside both of two half-planes; for events, the event that both A and B happen.
atan2 (G-2325)	A function atan2(rise, run) that returns the direction of a line over the full circle from its rise and run given separately; needed because the plain inverse tangent of the ratio cannot tell opposite directions apart.
Landmark measurement model (G-2401)	The density of a landmark reading (range, bearing) given a pose and the map: a Gaussian of the range error times a Gaussian of the bearing error, with the errors taken as independent; it scores how well a guessed pose explains what the robot saw.
Beam model (G-2377)	The measurement model of a range finder that mixes four densities for one reading (measurement noise, unexpected objects, failures, random readings) with weights that add up to 1, so no single odd reading can rule out the true pose.